



Computer Systems and Professional Practice

Professor Matthew Leeke
School of Computer Science
University of Birmingham

Topic 4 - Memory Systems

Session Outline

The memory hierarchy, including motivation and structure

Organising memory cells for addressing

Detecting and correcting errors in unreliable memory and network communications

The design and implementation of common memory components, including hard disks and optical discs

The Memory Hierarchy

The Choice of Memory Technology

To construct a memory system, we must think carefully about the properties of each storage approach

Only condition is that a memory system must facilitate data modification (writing) and data detection (reading)

Many factors influence the choice of memory technology

Frequency of access

Access time

Capacity required

Cost, e.g., cost per bit

The Designer's Dilemma for Memory Systems

The designer's dilemma

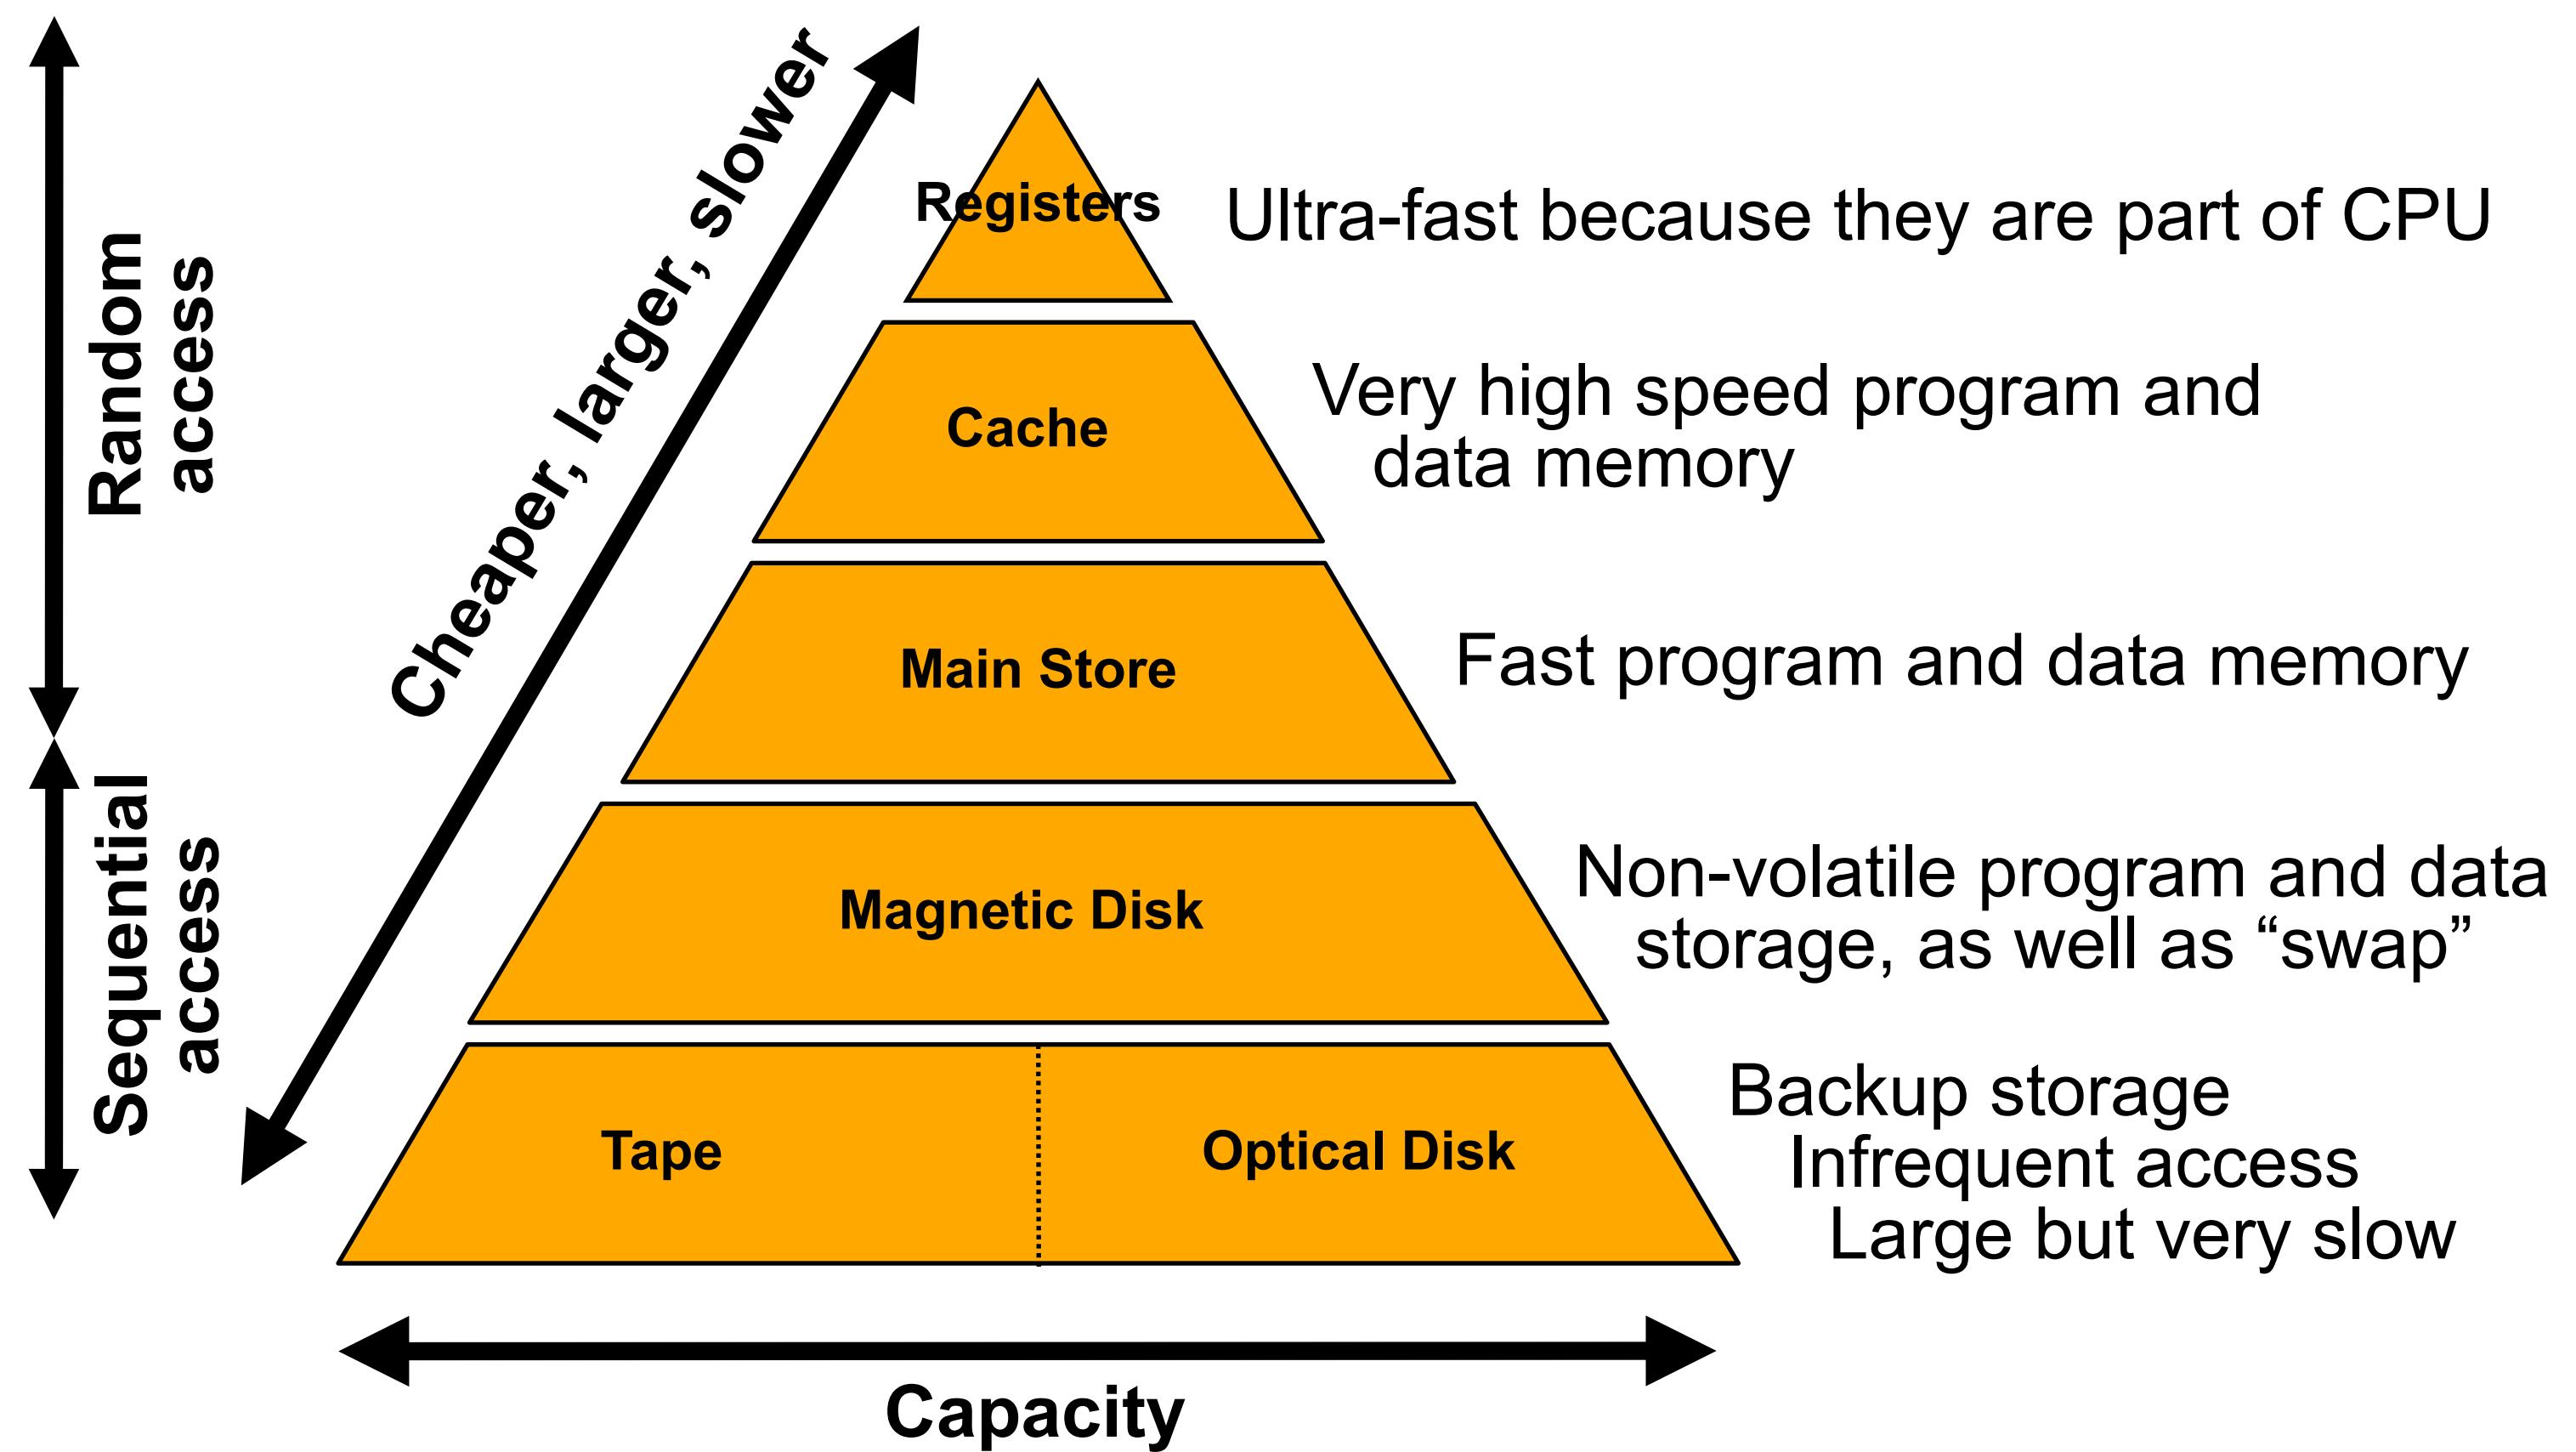
Desire low cost and high capacity - leads to choice of cheap per bit, high capacity and slow access time

Desire high performance - leads to choice of expensive per bit, low capacity and fast access time

To solve the problem we choose both by organising memory into a hierarchy

The Memory Hierarchy

A general view of the memory system that underpin modern computer systems



Why Have a Memory Hierarchy?

Economics!

Small sequential loops, subroutines, operations on arrays, tables and other structures mean that the locations in memory which will be accessed is somewhat predictable

Temporal locality - If a particular memory location is referenced, it is likely that the same location will be referenced again in the near future

Spatial locality - If a particular memory location is referenced, it is likely that nearby memory locations will be referenced in the near future

Reasonable assumption - It has been shown that 90% of memory accesses are within ± 2 kilobytes of previous PC position

Semiconductor Memory Types

Particularly interested in memory that provides random access

RAM is most common type of semiconductor memory, though the term is misleading since all types shown above allow random access

Memory Type	Category	Erasure	Write Mechanism	Volatility
Random Access Memory (RAM)	Read-write	Electrically at byte-level	Electrically written	Volatile
Read-only Memory (ROM)	Read-only	Not possible	Mask written	Non-volatile
Programmable ROM (PROM)		Not possible	Electrically written	
Erasable PROM (EPROM)	Read-mostly	UV-light at chip-level	Non-volatile	
Electrically Erasable PROM (EEPROM)		Electrically at byte-level		
Flash Memory		Electrically at block-level		

Cache Memory

What is a Cache?

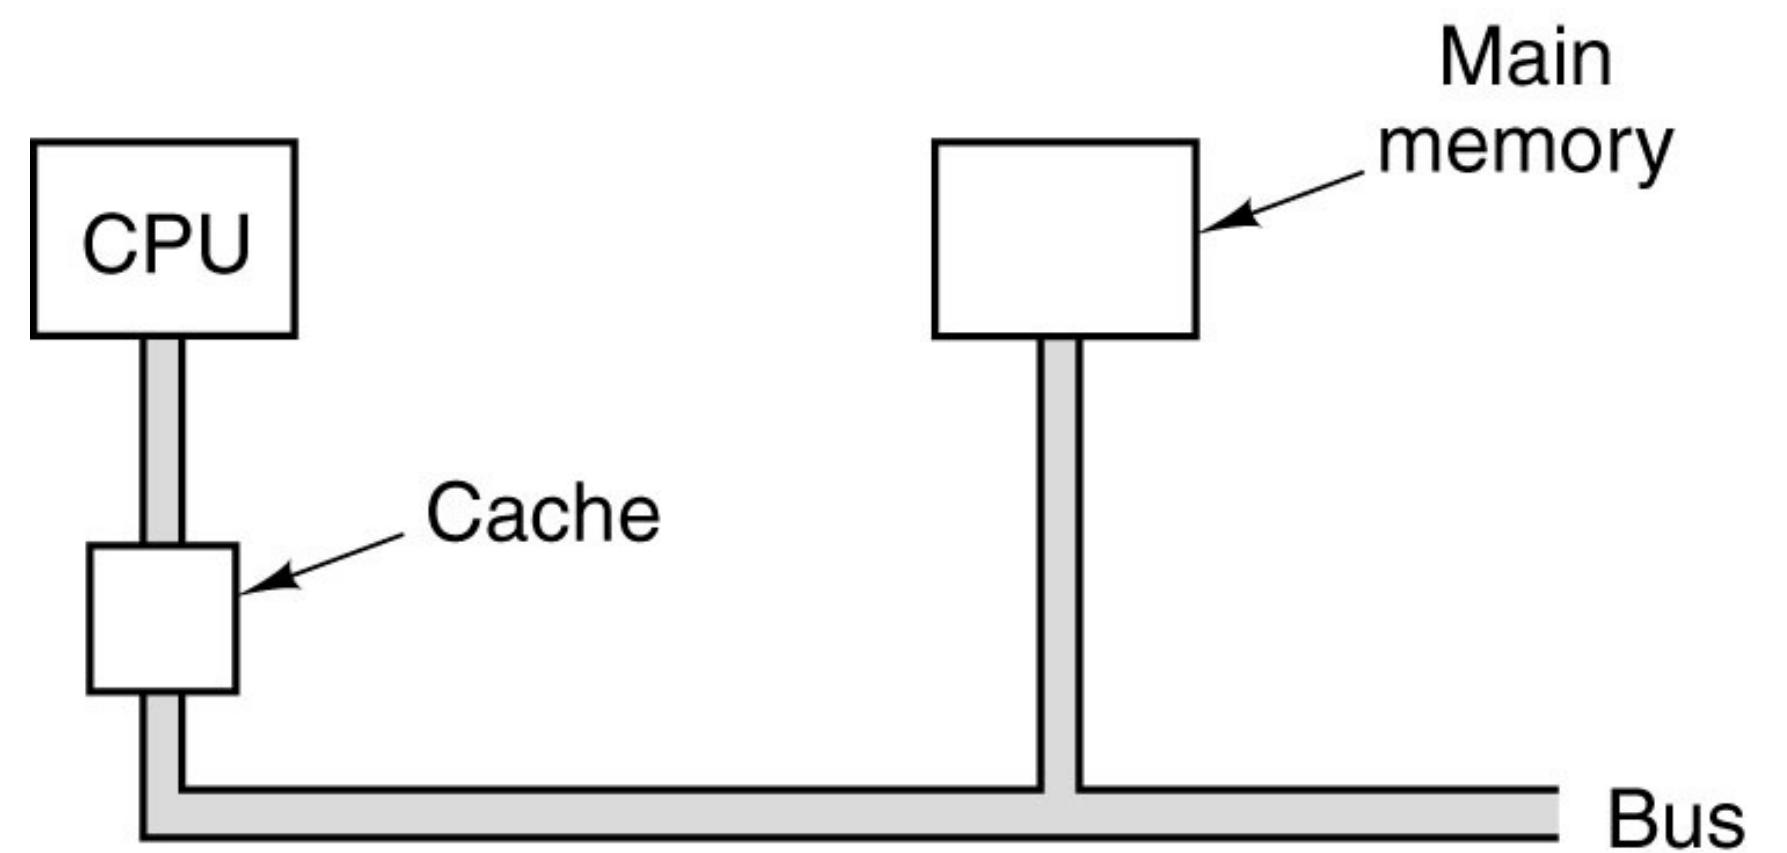
"90% of memory accesses within only 2 Kbytes"

So store those 2 Kbytes in a small, fast "cache" memory

If data required by CPU is in the cache (a "cache hit"),
big speed improvement

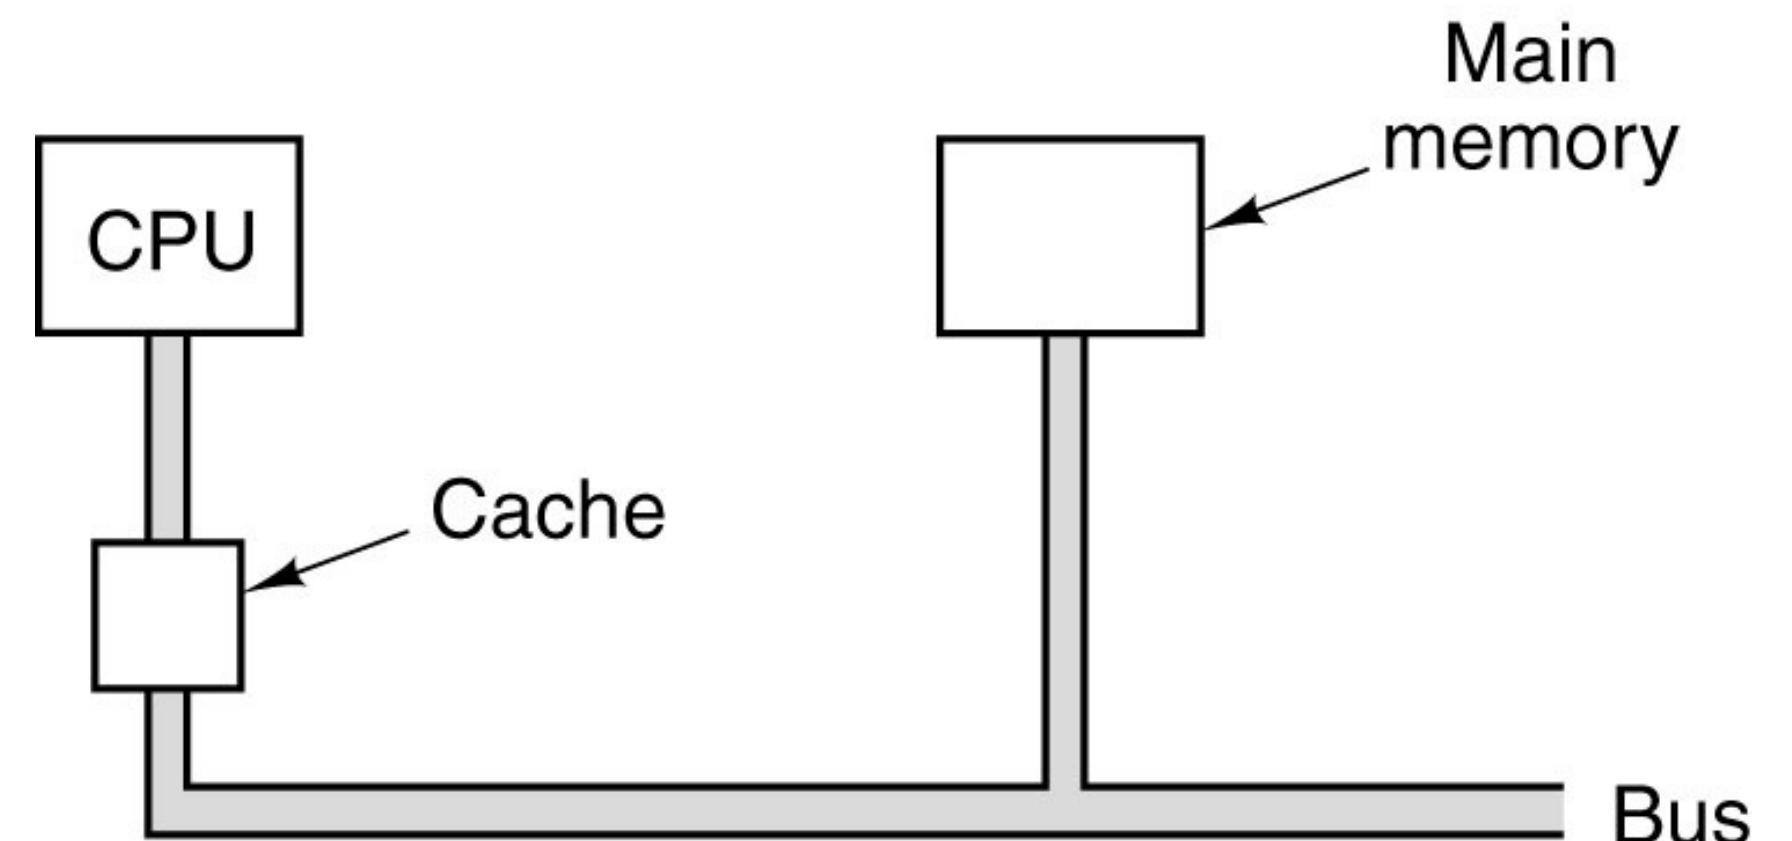
Cache is small to limit cost

Note that "cache" means "hiding place" - transparent
to programmer



What is a Cache?

execute typical instruction	1/1,000,000,000 sec = 1 nanosec
fetch from L1 cache memory	0.5 nanosec
branch misprediction	5 nanosec
fetch from L2 cache memory	7 nanosec
Mutex lock/unlock	25 nanosec
fetch from main memory	100 nanosec
send 2K bytes over 1Gbps network	20,000 nanosec
read 1MB sequentially from memory	250,000 nanosec
fetch from new disk location (seek)	8,000,000 nanosec
read 1MB sequentially from disk	20,000,000 nanosec
send packet US to Europe and back	150 milliseconds = 150,000,000 nanosec

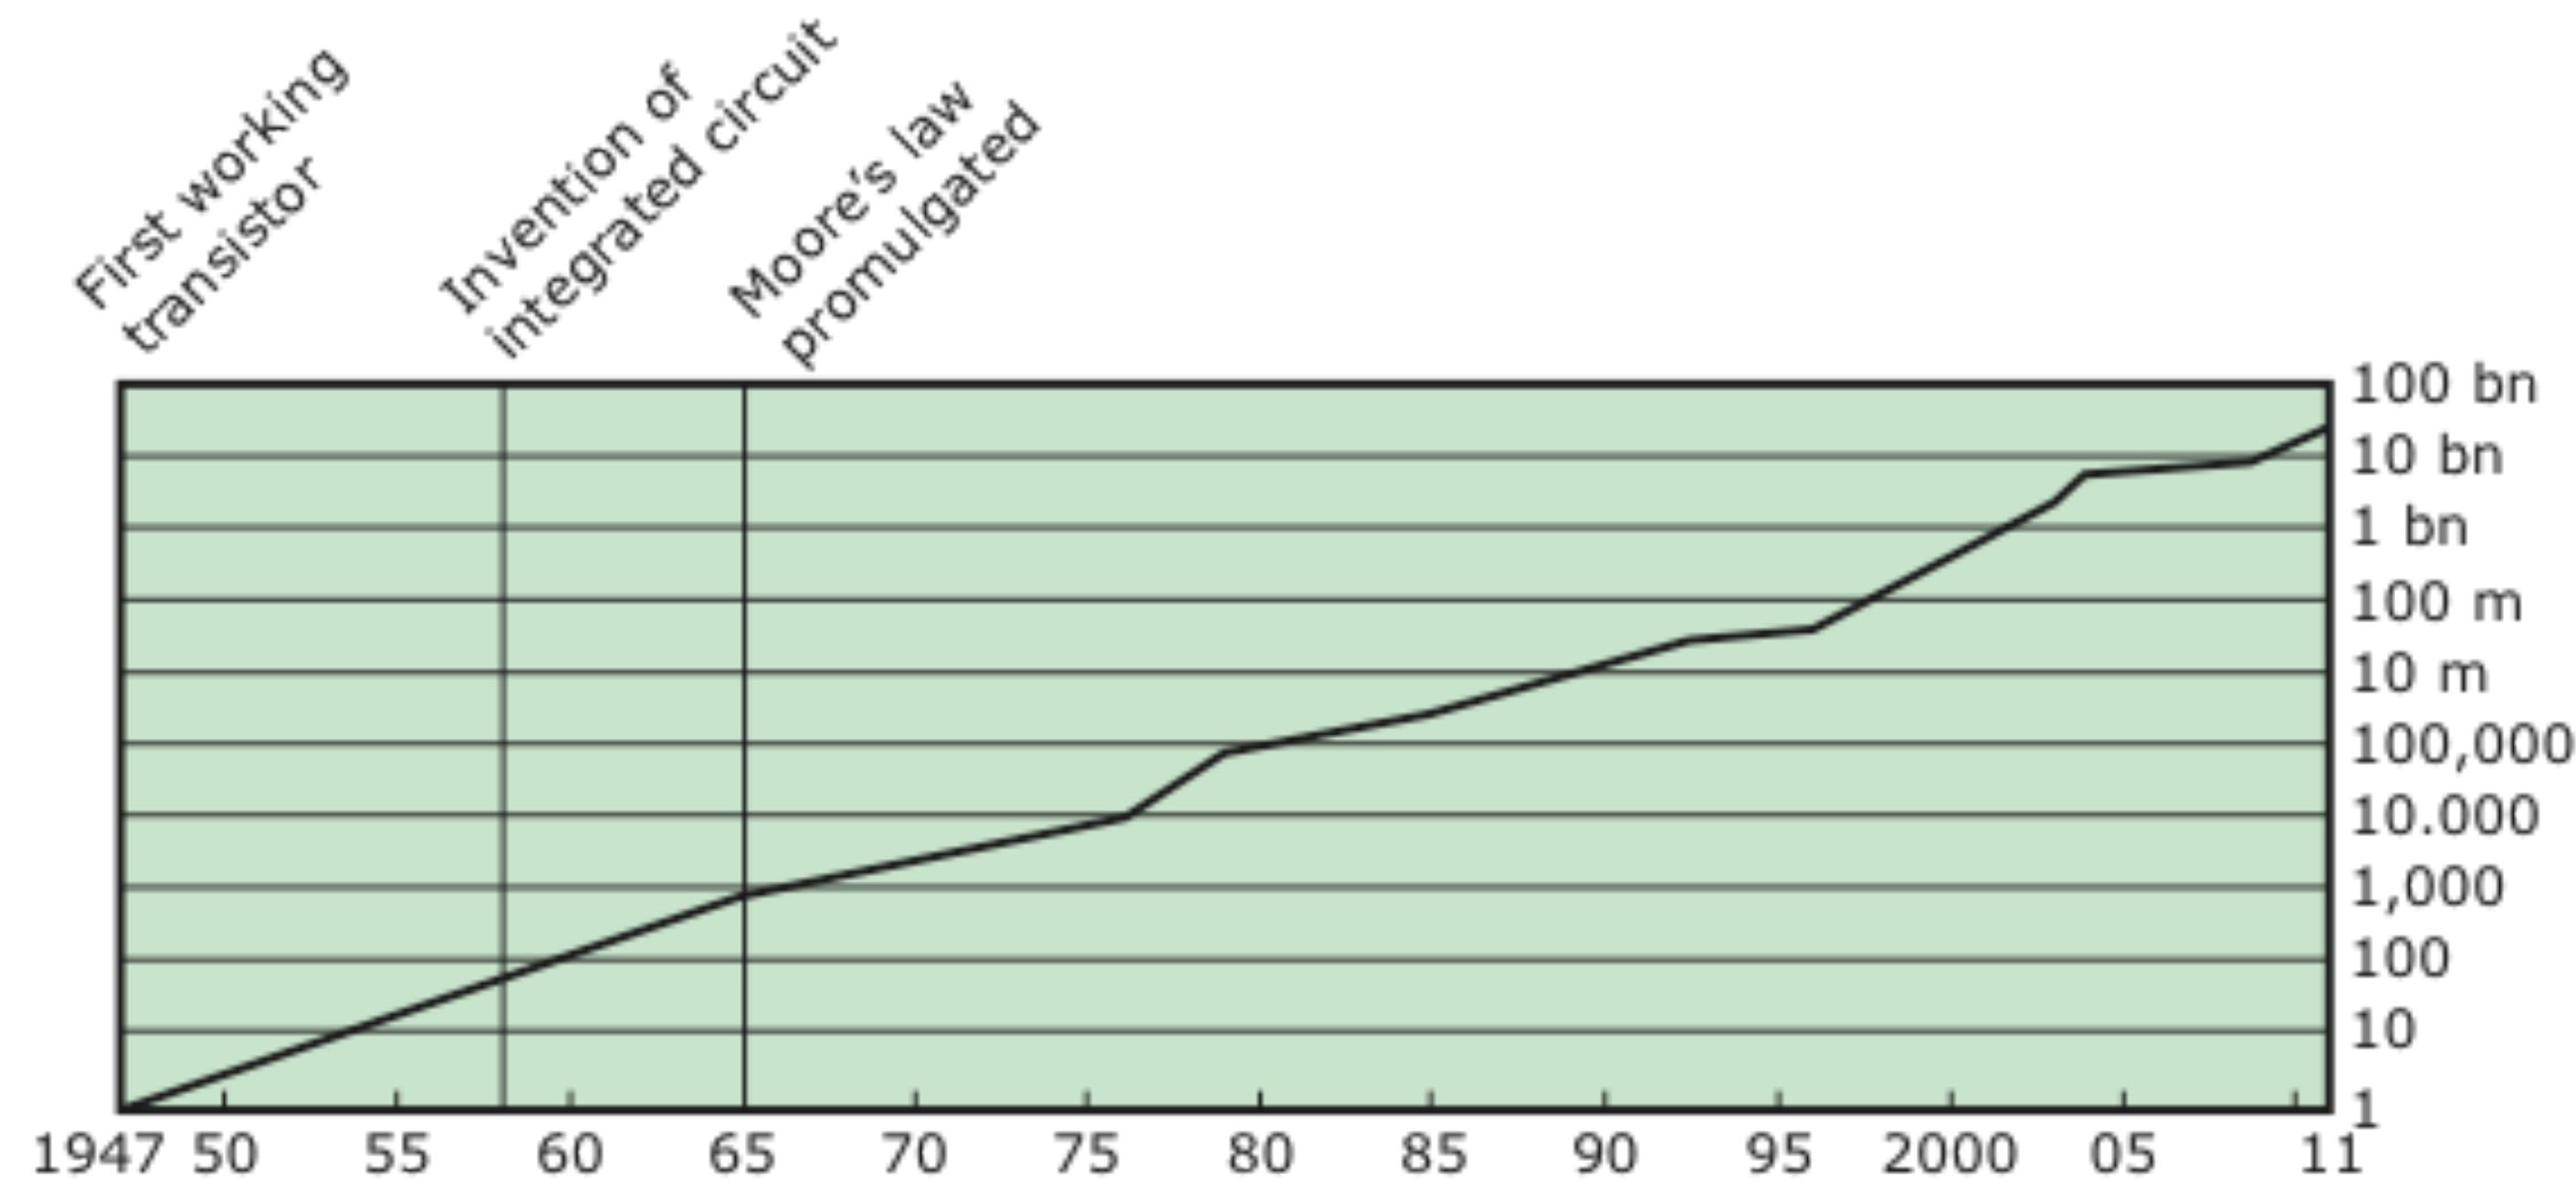


Moore's Law

"The complexity for minimum component costs has increased at a rate of roughly a factor of two per year. Certainly over the short term this rate can be expected to continue, if not to increase. Over the longer term, the rate of increase is a bit more uncertain, although there is no reason to believe it will not remain nearly constant for at least 10 years." - **Gordon Moore (co-founder of Intel), Electronics, Volume 38, No. 8, April 1965**

Now more commonly expressed as "The number of transistors on a memory chip doubles every 18 months"

Moore's Law



Reproduced from: W.Stallings, "Computer Organization and Architecture", 9th Edition, Pearson, 2013

Moore's Law

The cost of computer logic and memory circuitry has fallen at a dramatic rate

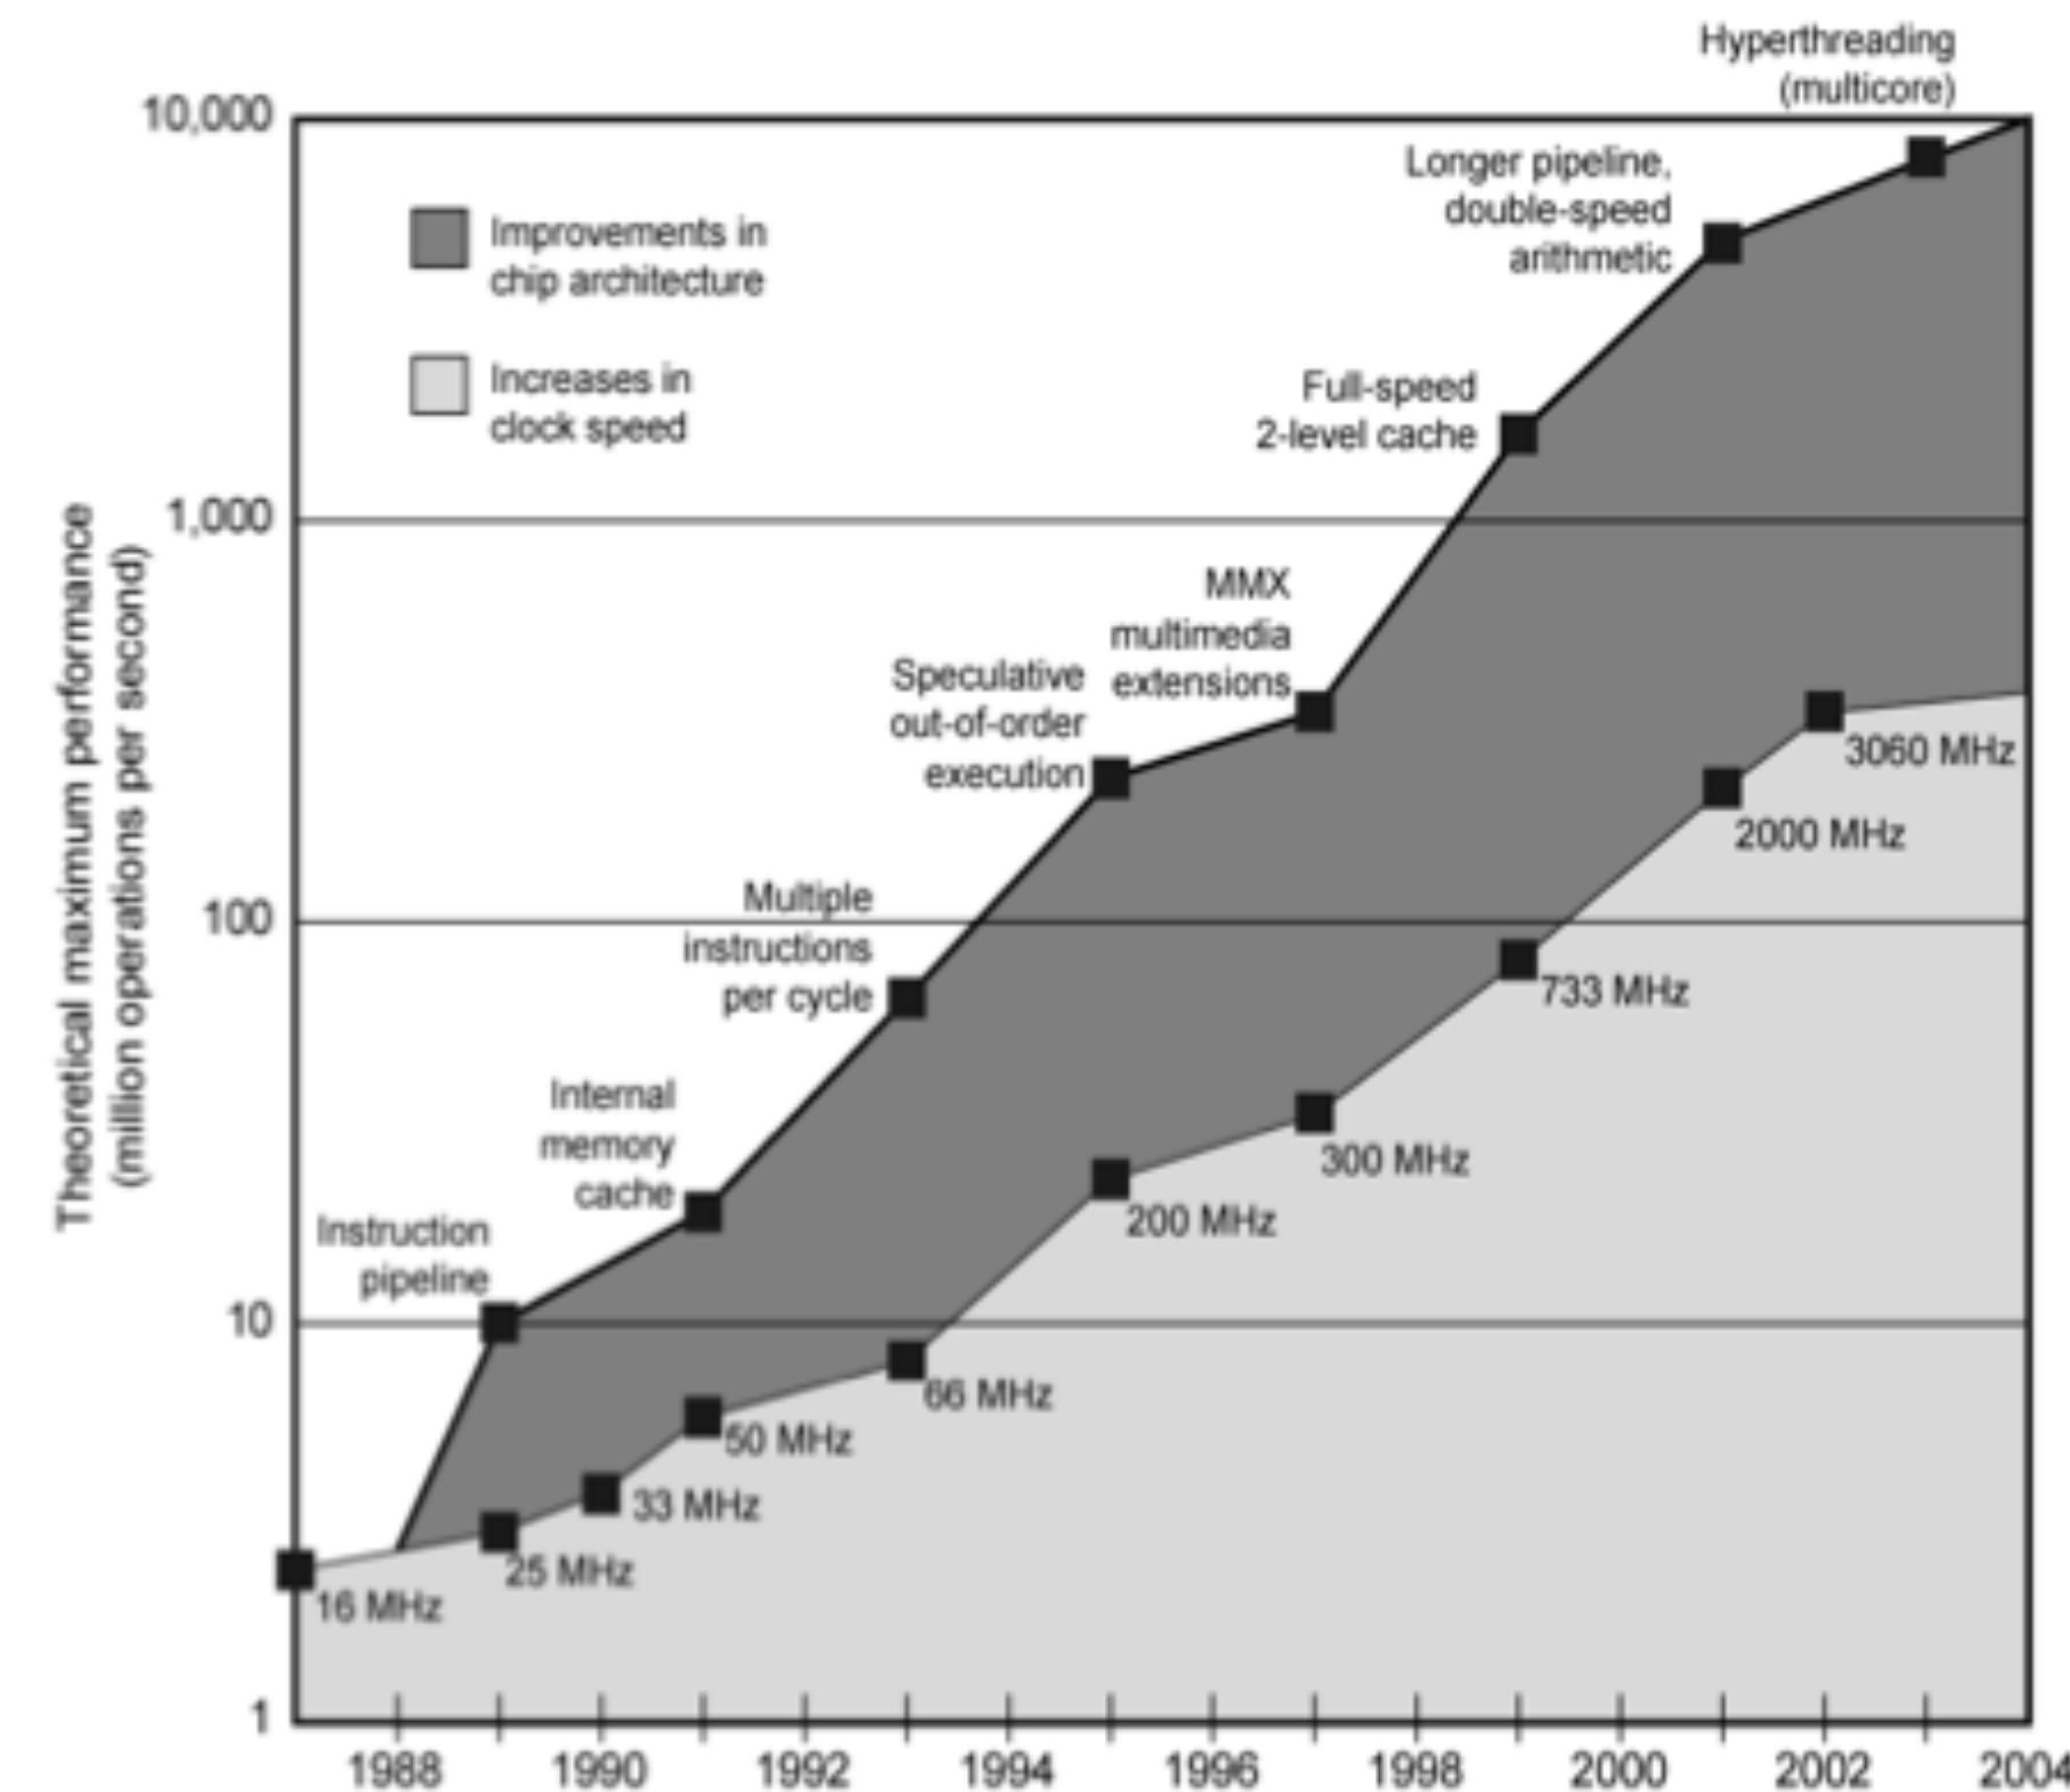
As logic and memory is placed closer together on more densely packed chips, the electrical path is shortened, increasing operating speed

Computers have become smaller, making them more convenient for use in a variety of environments, especially embedded systems

Reduction in power and cooling requirements

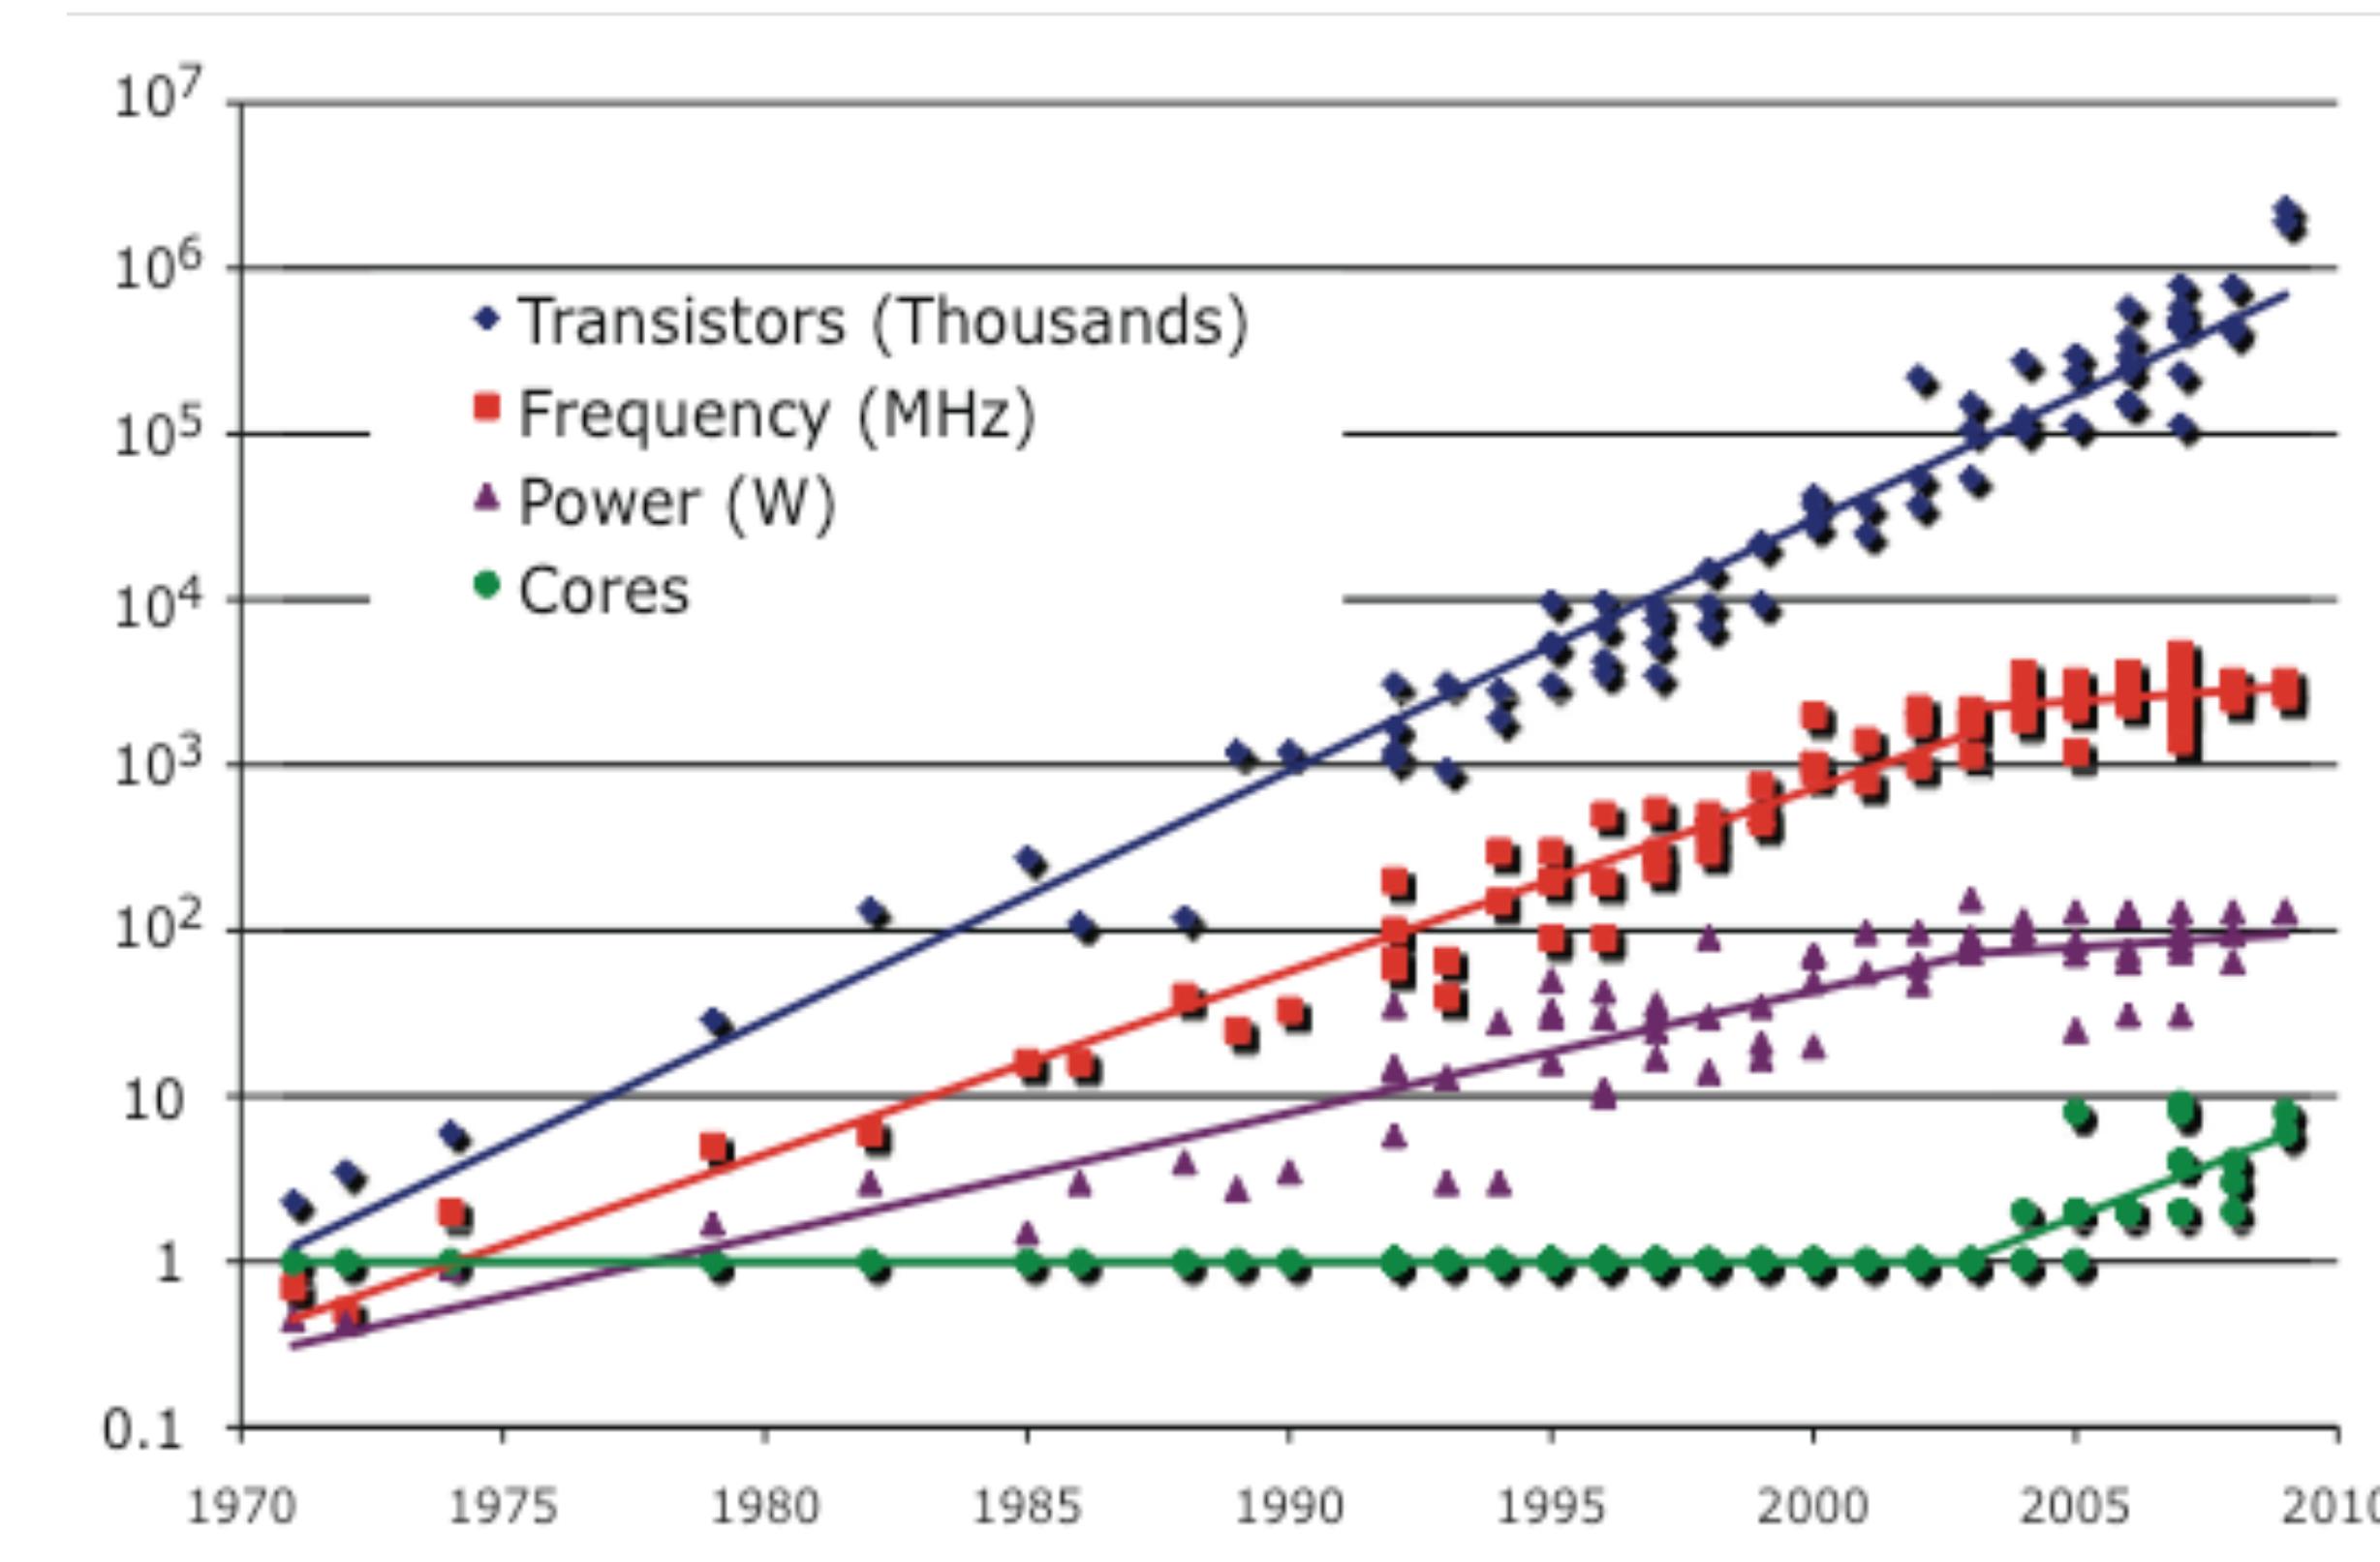
Interconnections on an IC are more reliable than solder connections, thus having fewer off-chip connections increases reliability

Intel Microprocessor Performance



Reproduced from: W.Stallings, "Computer Organization and Architecture", 8th Edition, Pearson, 2010

Current Outlook - Hardware

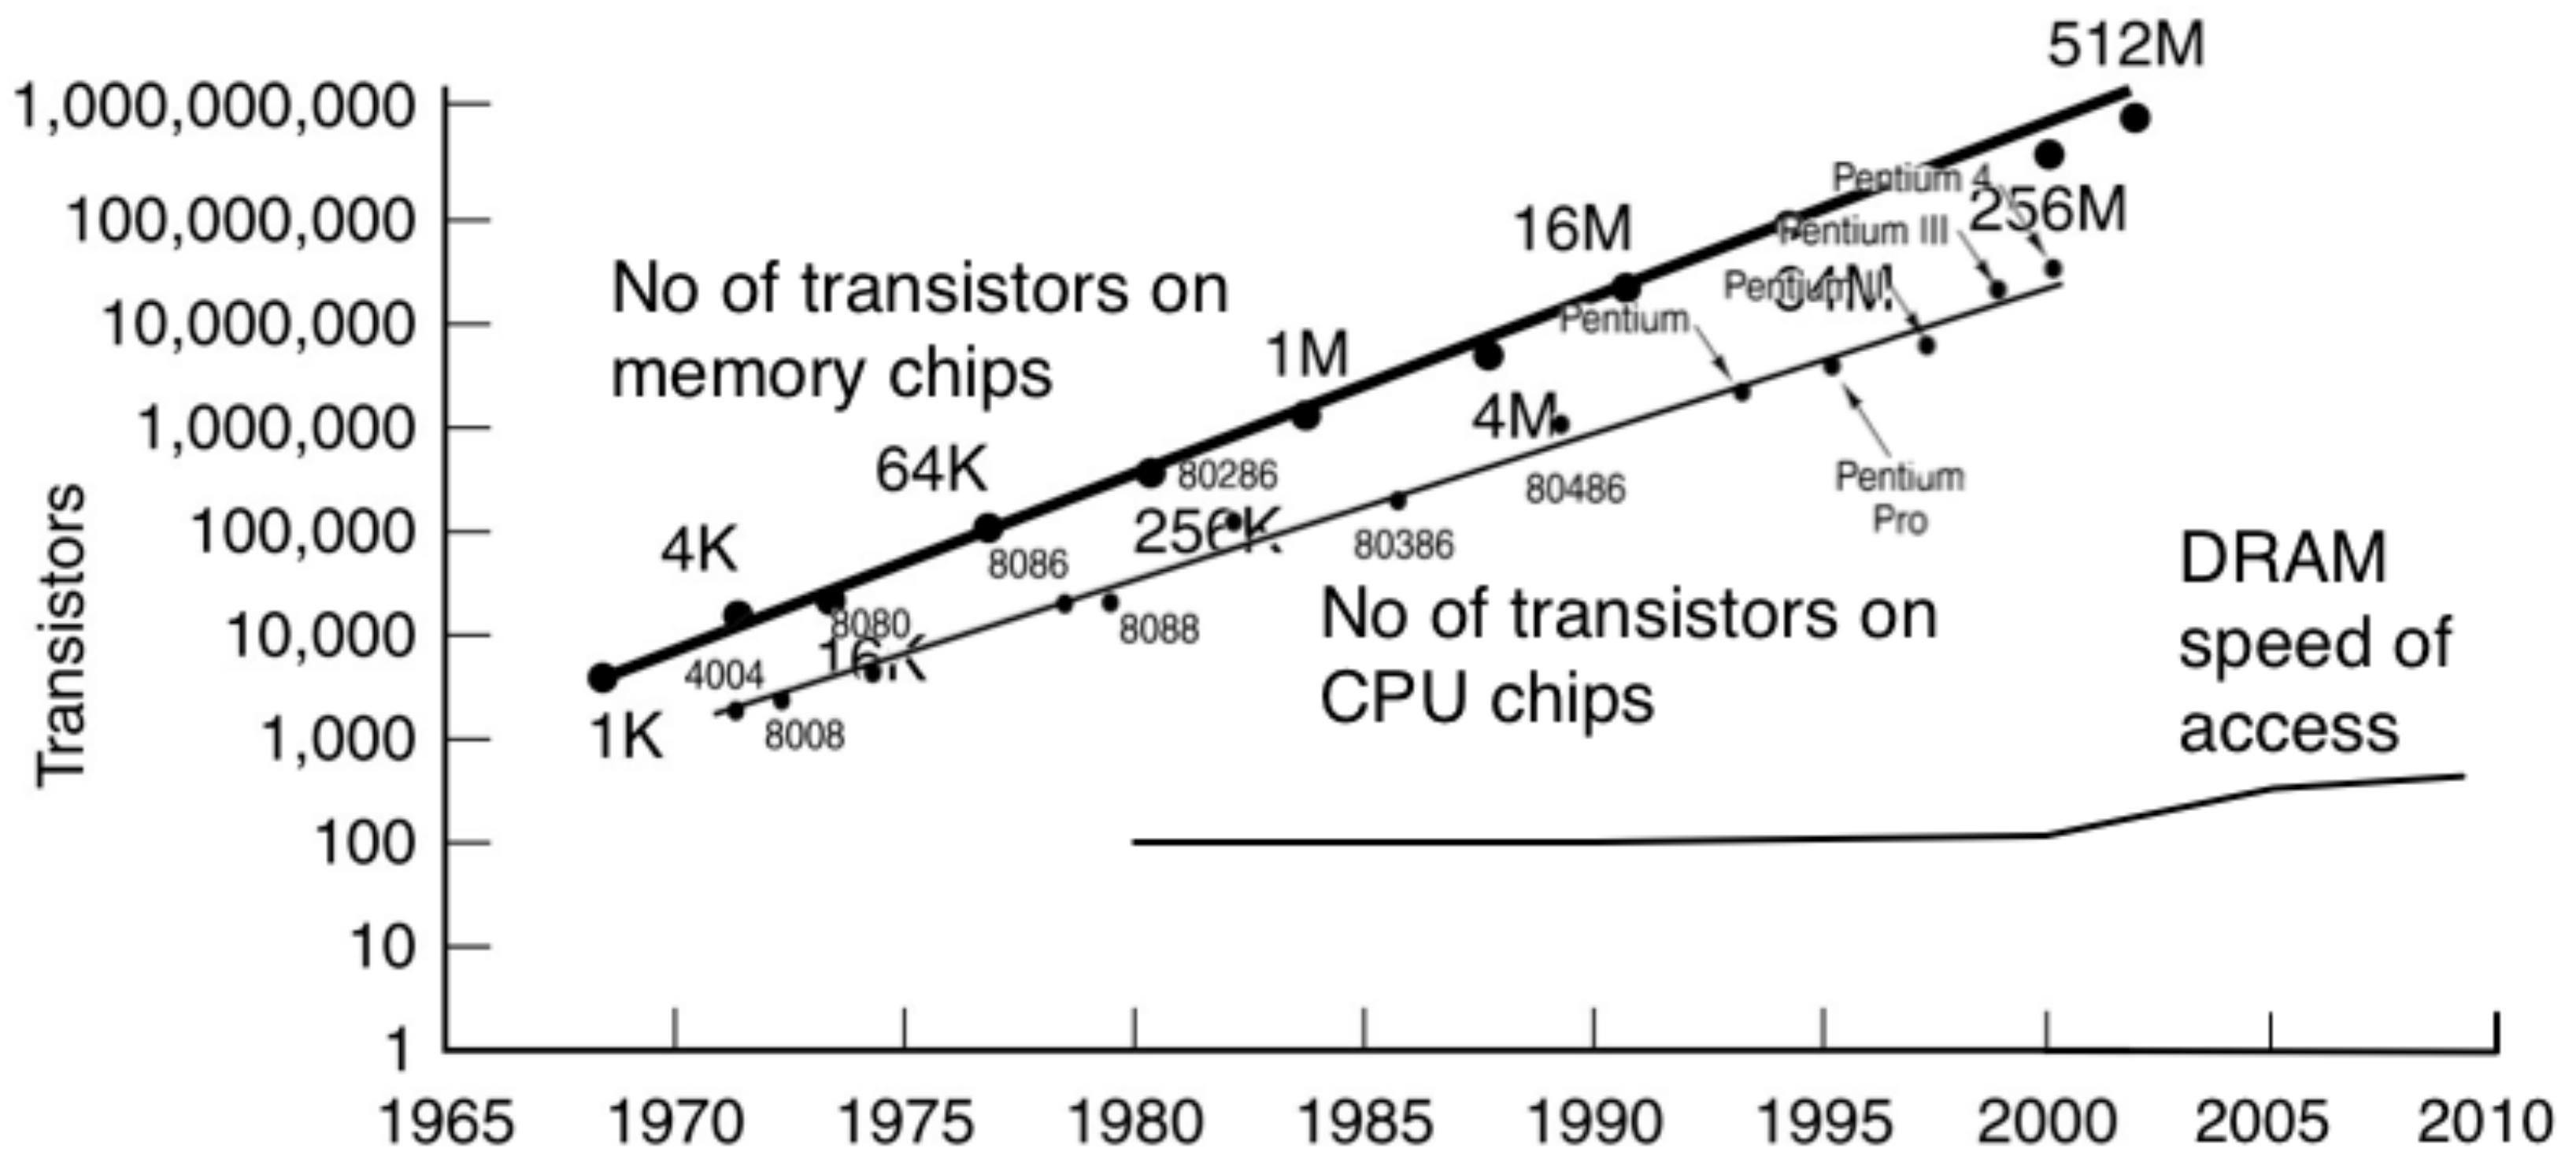


Reproduced from: W.Stallings, "Computer Organization and Architecture", 9th Edition, Pearson, 2013

Why Have a Cache?

Memory size is increasing according to Moore's Law

Memory access speed is improving much more slowly



Cache Concepts

Caching read-only data is relatively straightforward

We don't need to consider the possibility that items will change, hence copies across the memory hierarchy will remain consistent

Two general strategies can be adopted when caching writes, both of which can employ a write buffer to enhance performance

Write Through - Updates the item in cache and writes through to update lower levels of the memory hierarchy

Write Back - Only updates copy in cache, ensuring that blocks are copied back to memory when they are to be replaced

Types of Cache Miss

Measures of cache performance include hit rate (h) and miss rate ($1-h$)

$$h = \frac{\text{Number of times the words are in cache}}{\text{Total number of memory references}}$$

We can understand cache misses by sorting all misses into categories

Compulsory - Misses that would occur regardless of cache size, e.g., the first access to a block can not be in cache (regardless of cache size) so the block must be retrieved

Capacity - Misses occurring because a cache is not large enough to contain all blocks needed during the execution of a program

Conflict - Misses occurring as a result of the placement strategy for blocks not being fully associative, which means that a block may be discarded and retrieved

Coherency - Misses occurring due to cache flushes in multiprocessor systems

Measuring Cache Performance

Performance measures based solely on hit / miss rates don't factor in the cost of a cache miss - the real performance issue!

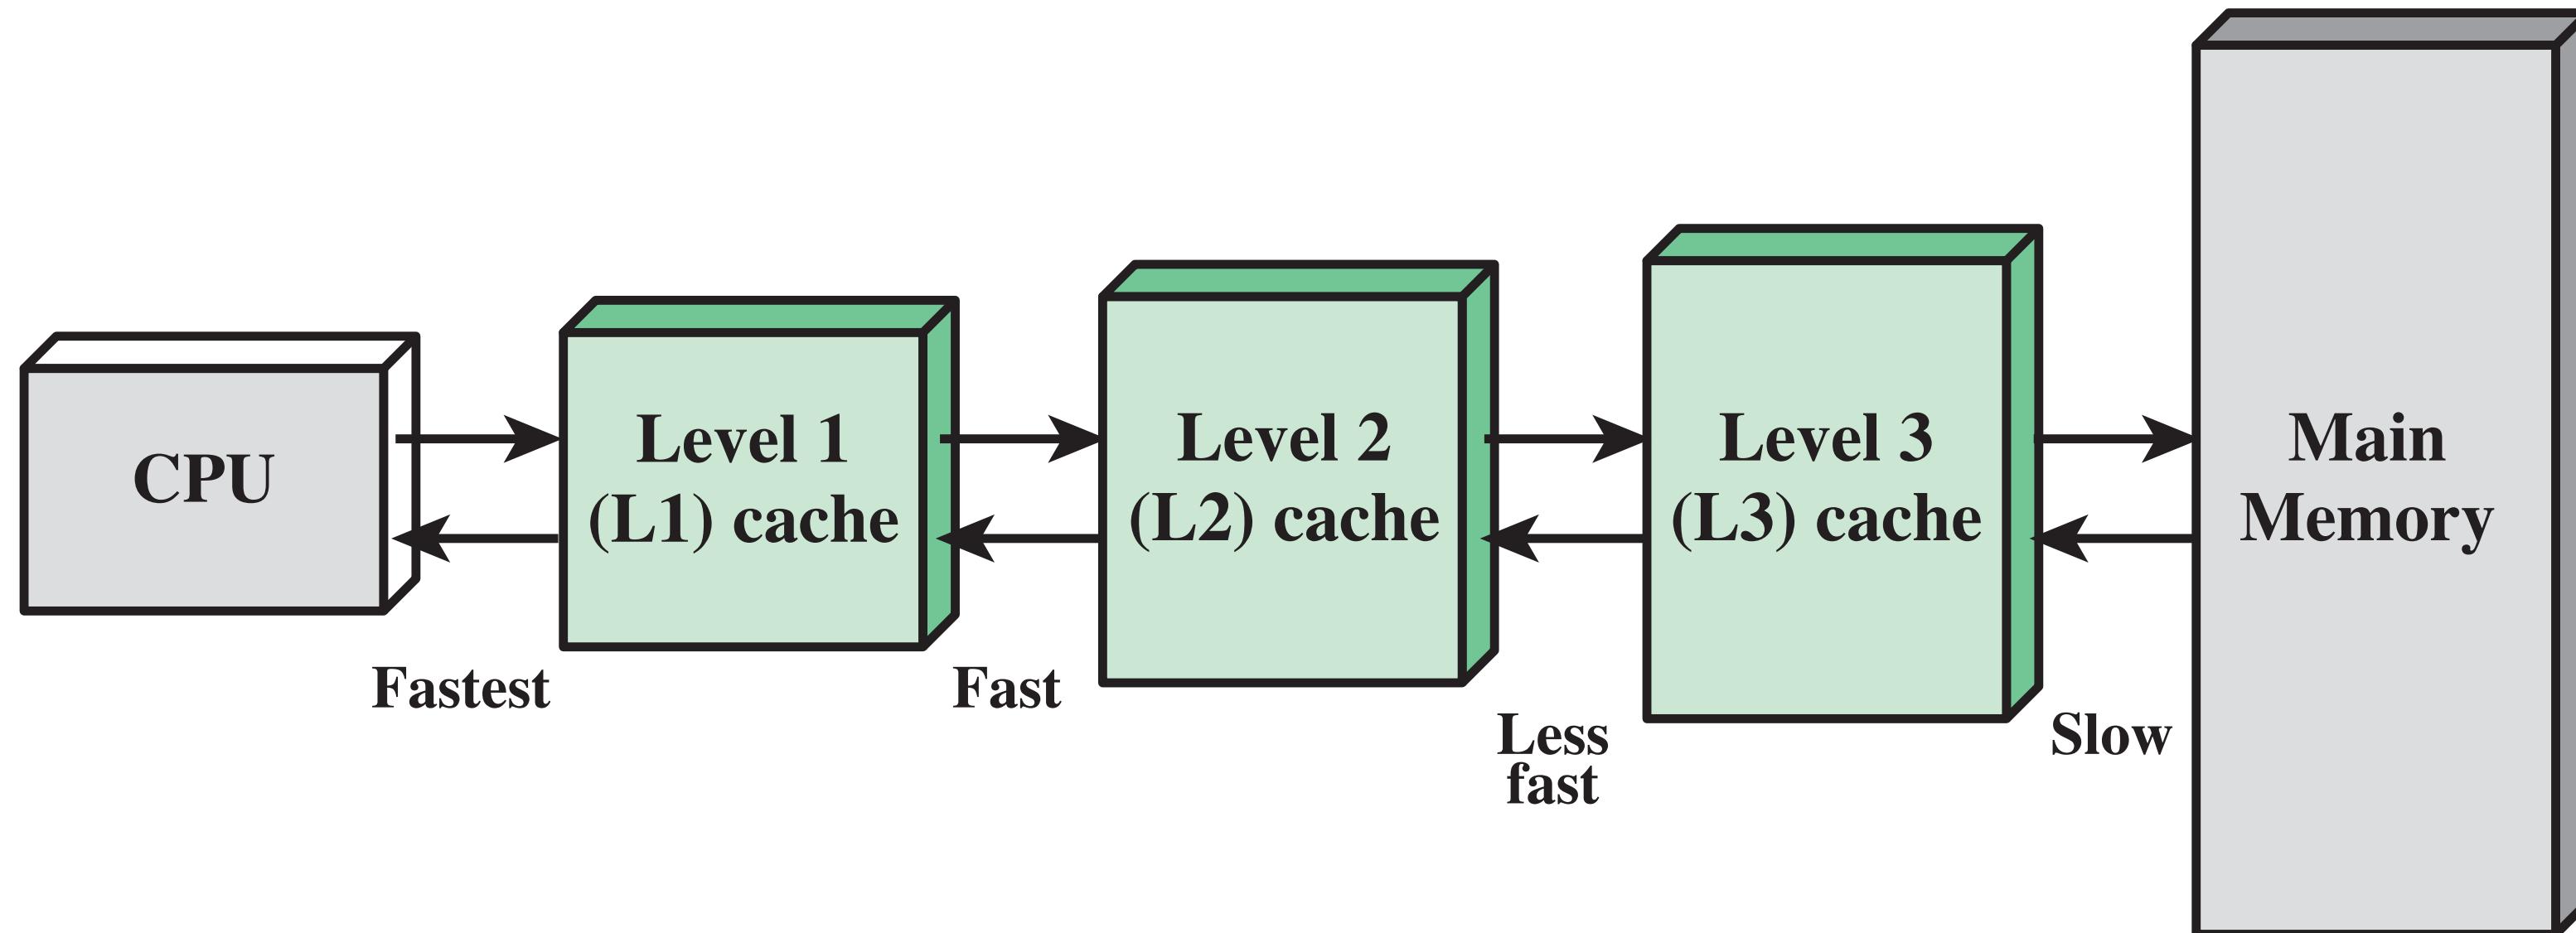
Average memory access time is an alternative measure that accounts for the cost of a cache miss

$$\text{Average memory access time} = \text{Hit time} + (\text{Miss rate} \times \text{Miss Penalty})$$

This is where hit time is the time to hit in the cache and the miss penalty is the time taken to replace the block from memory

Even average access time is still an indirect measure of performance - execution time can often be a better measure in real-world situations

Multilevel Cache



Reproduced from: W.Stallings, “Computer Organization and Architecture”, 9th Edition, Pearson, 2013

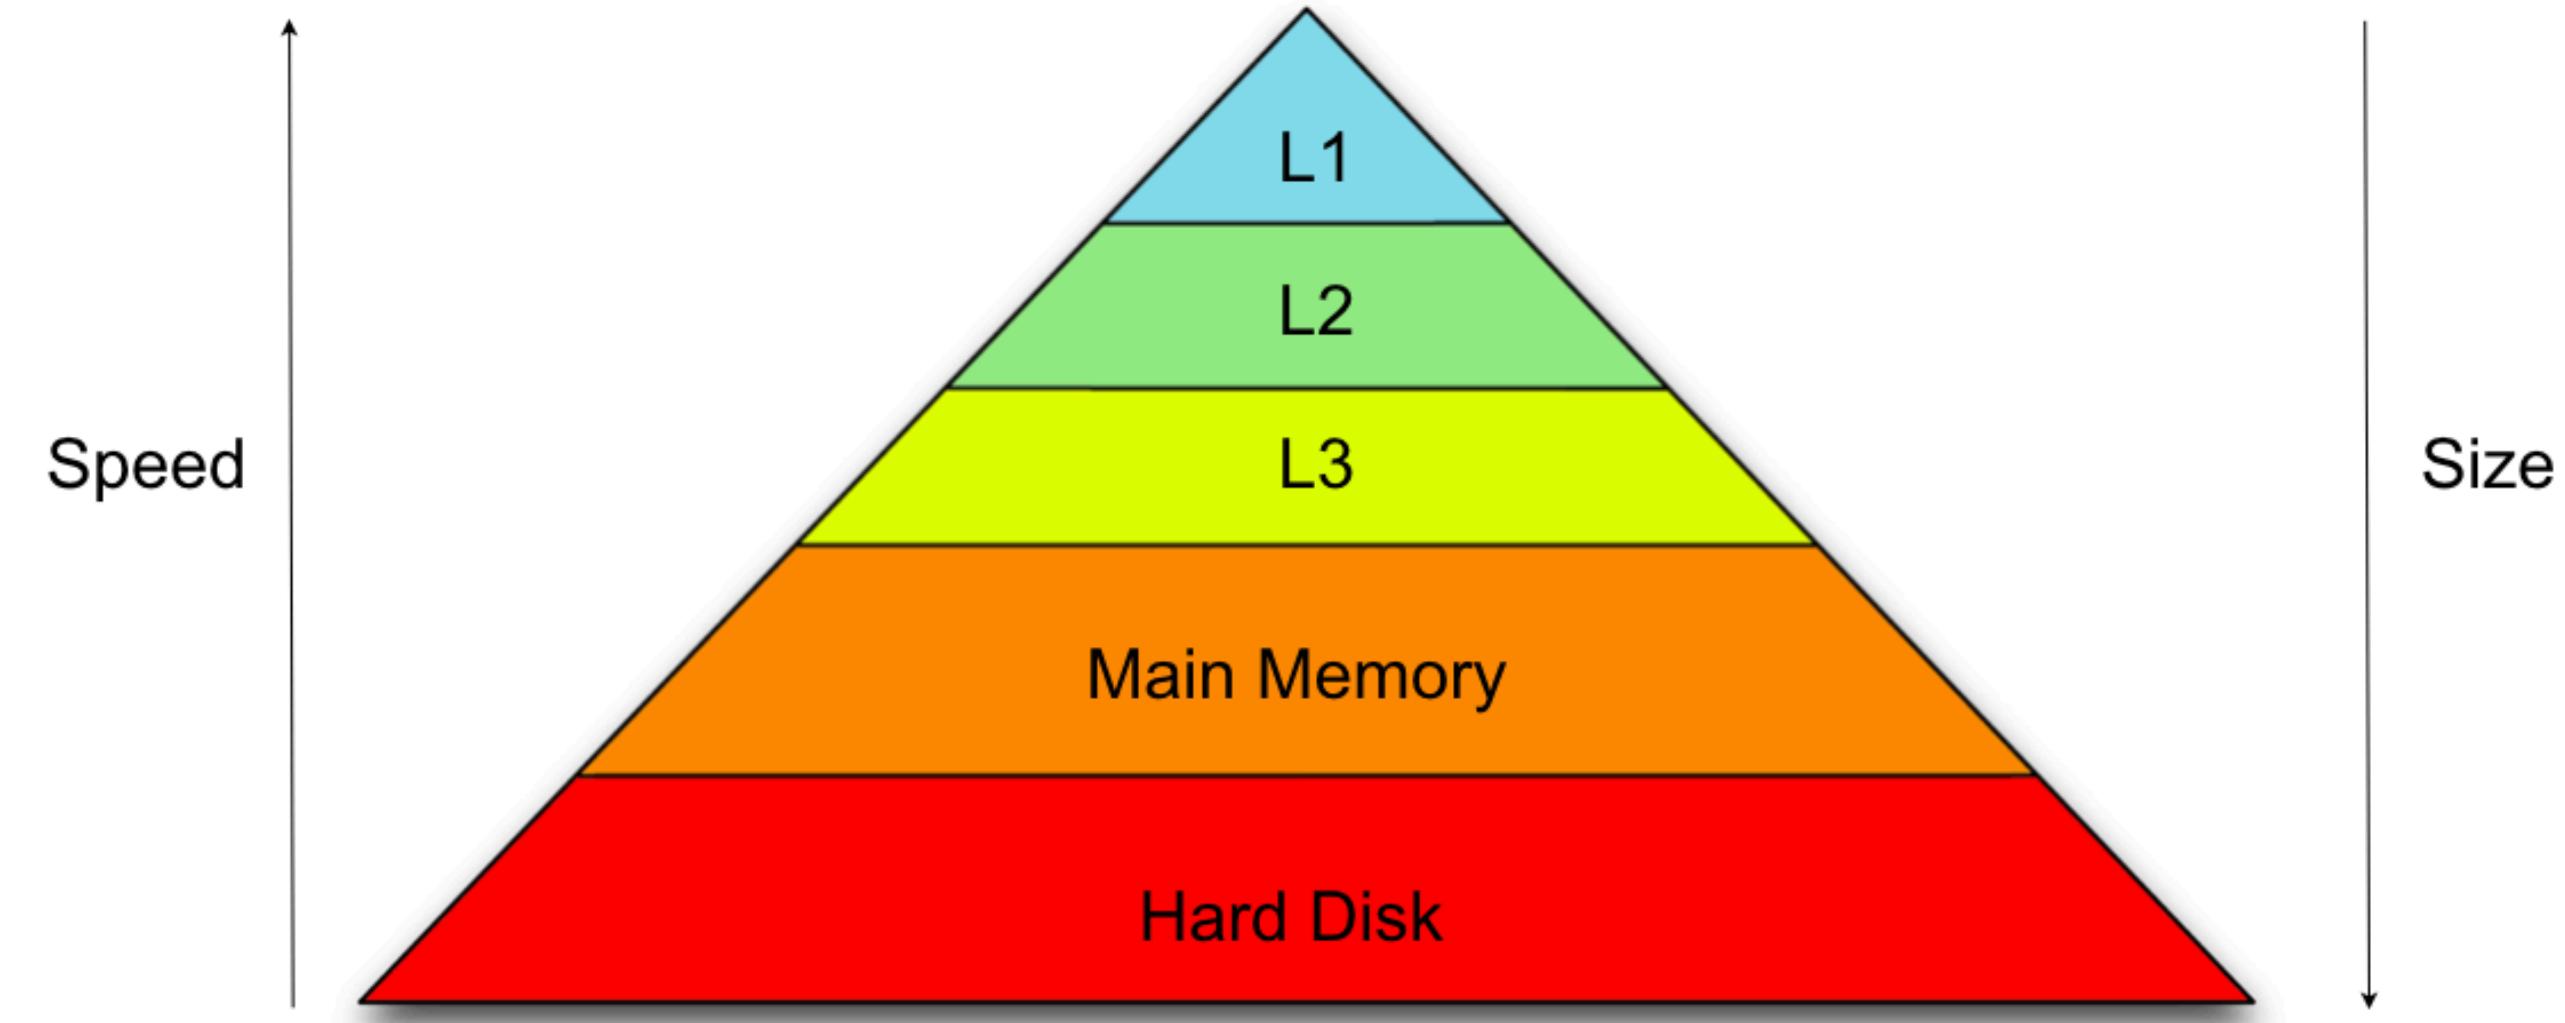
Relating the Memory Hierarchy to Cache

Data storage is a balance of size and speed:

Memory closer to the CPU (Cache - L1, L2, L3) is typically faster but has a far smaller capacity.

Main memory is further and slower but has a far greater capacity

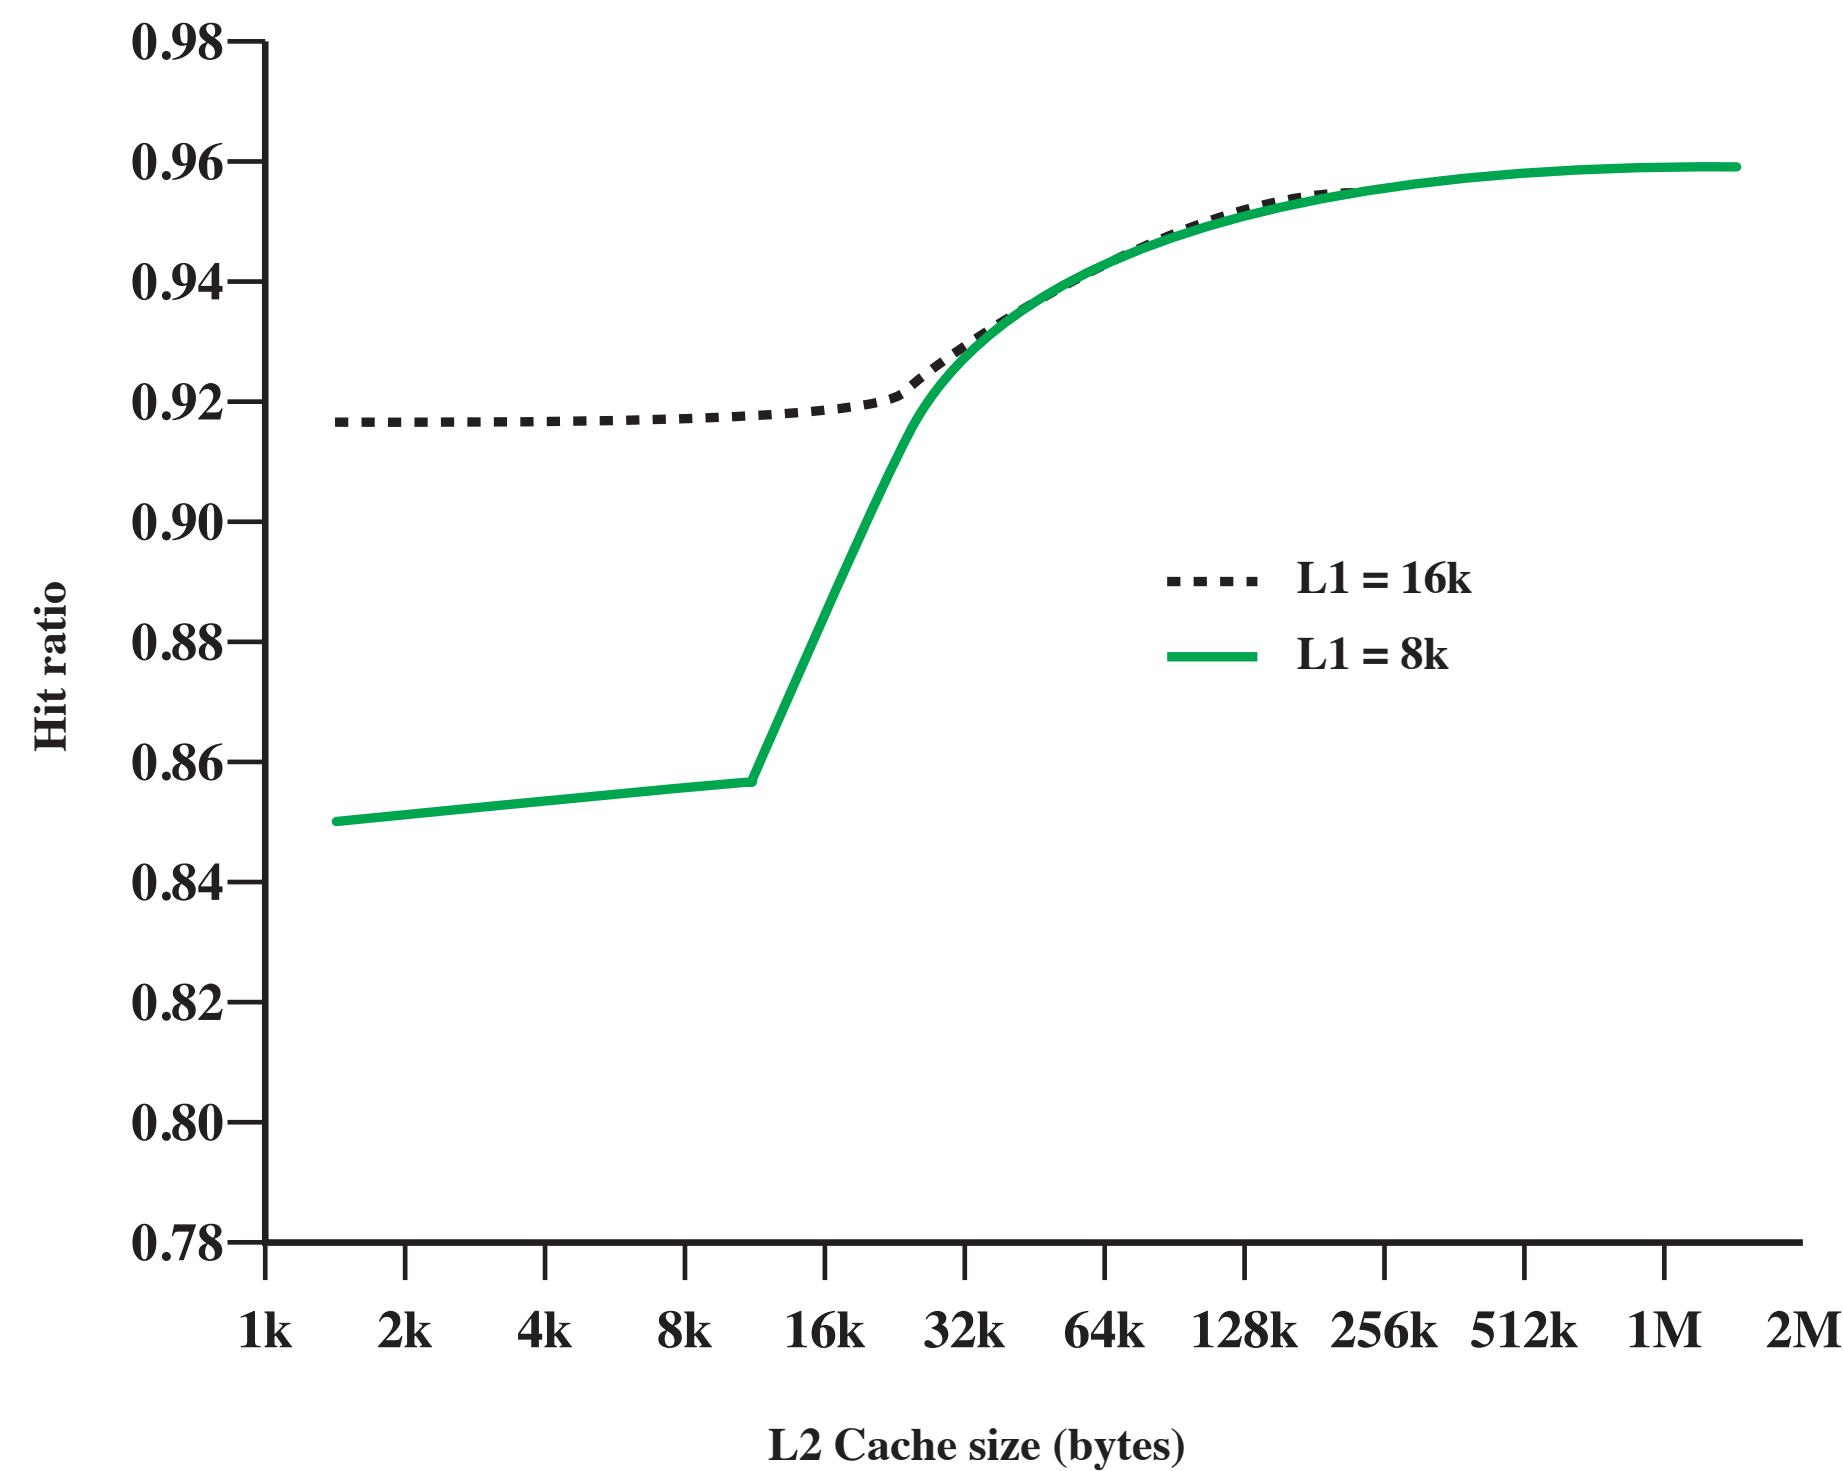
Hard disks have the greatest capacity, and do not need continuous power, but have the slowest access times.



Optimisation: We want to use fastest memory possible for high performance, but problems may large datasets.

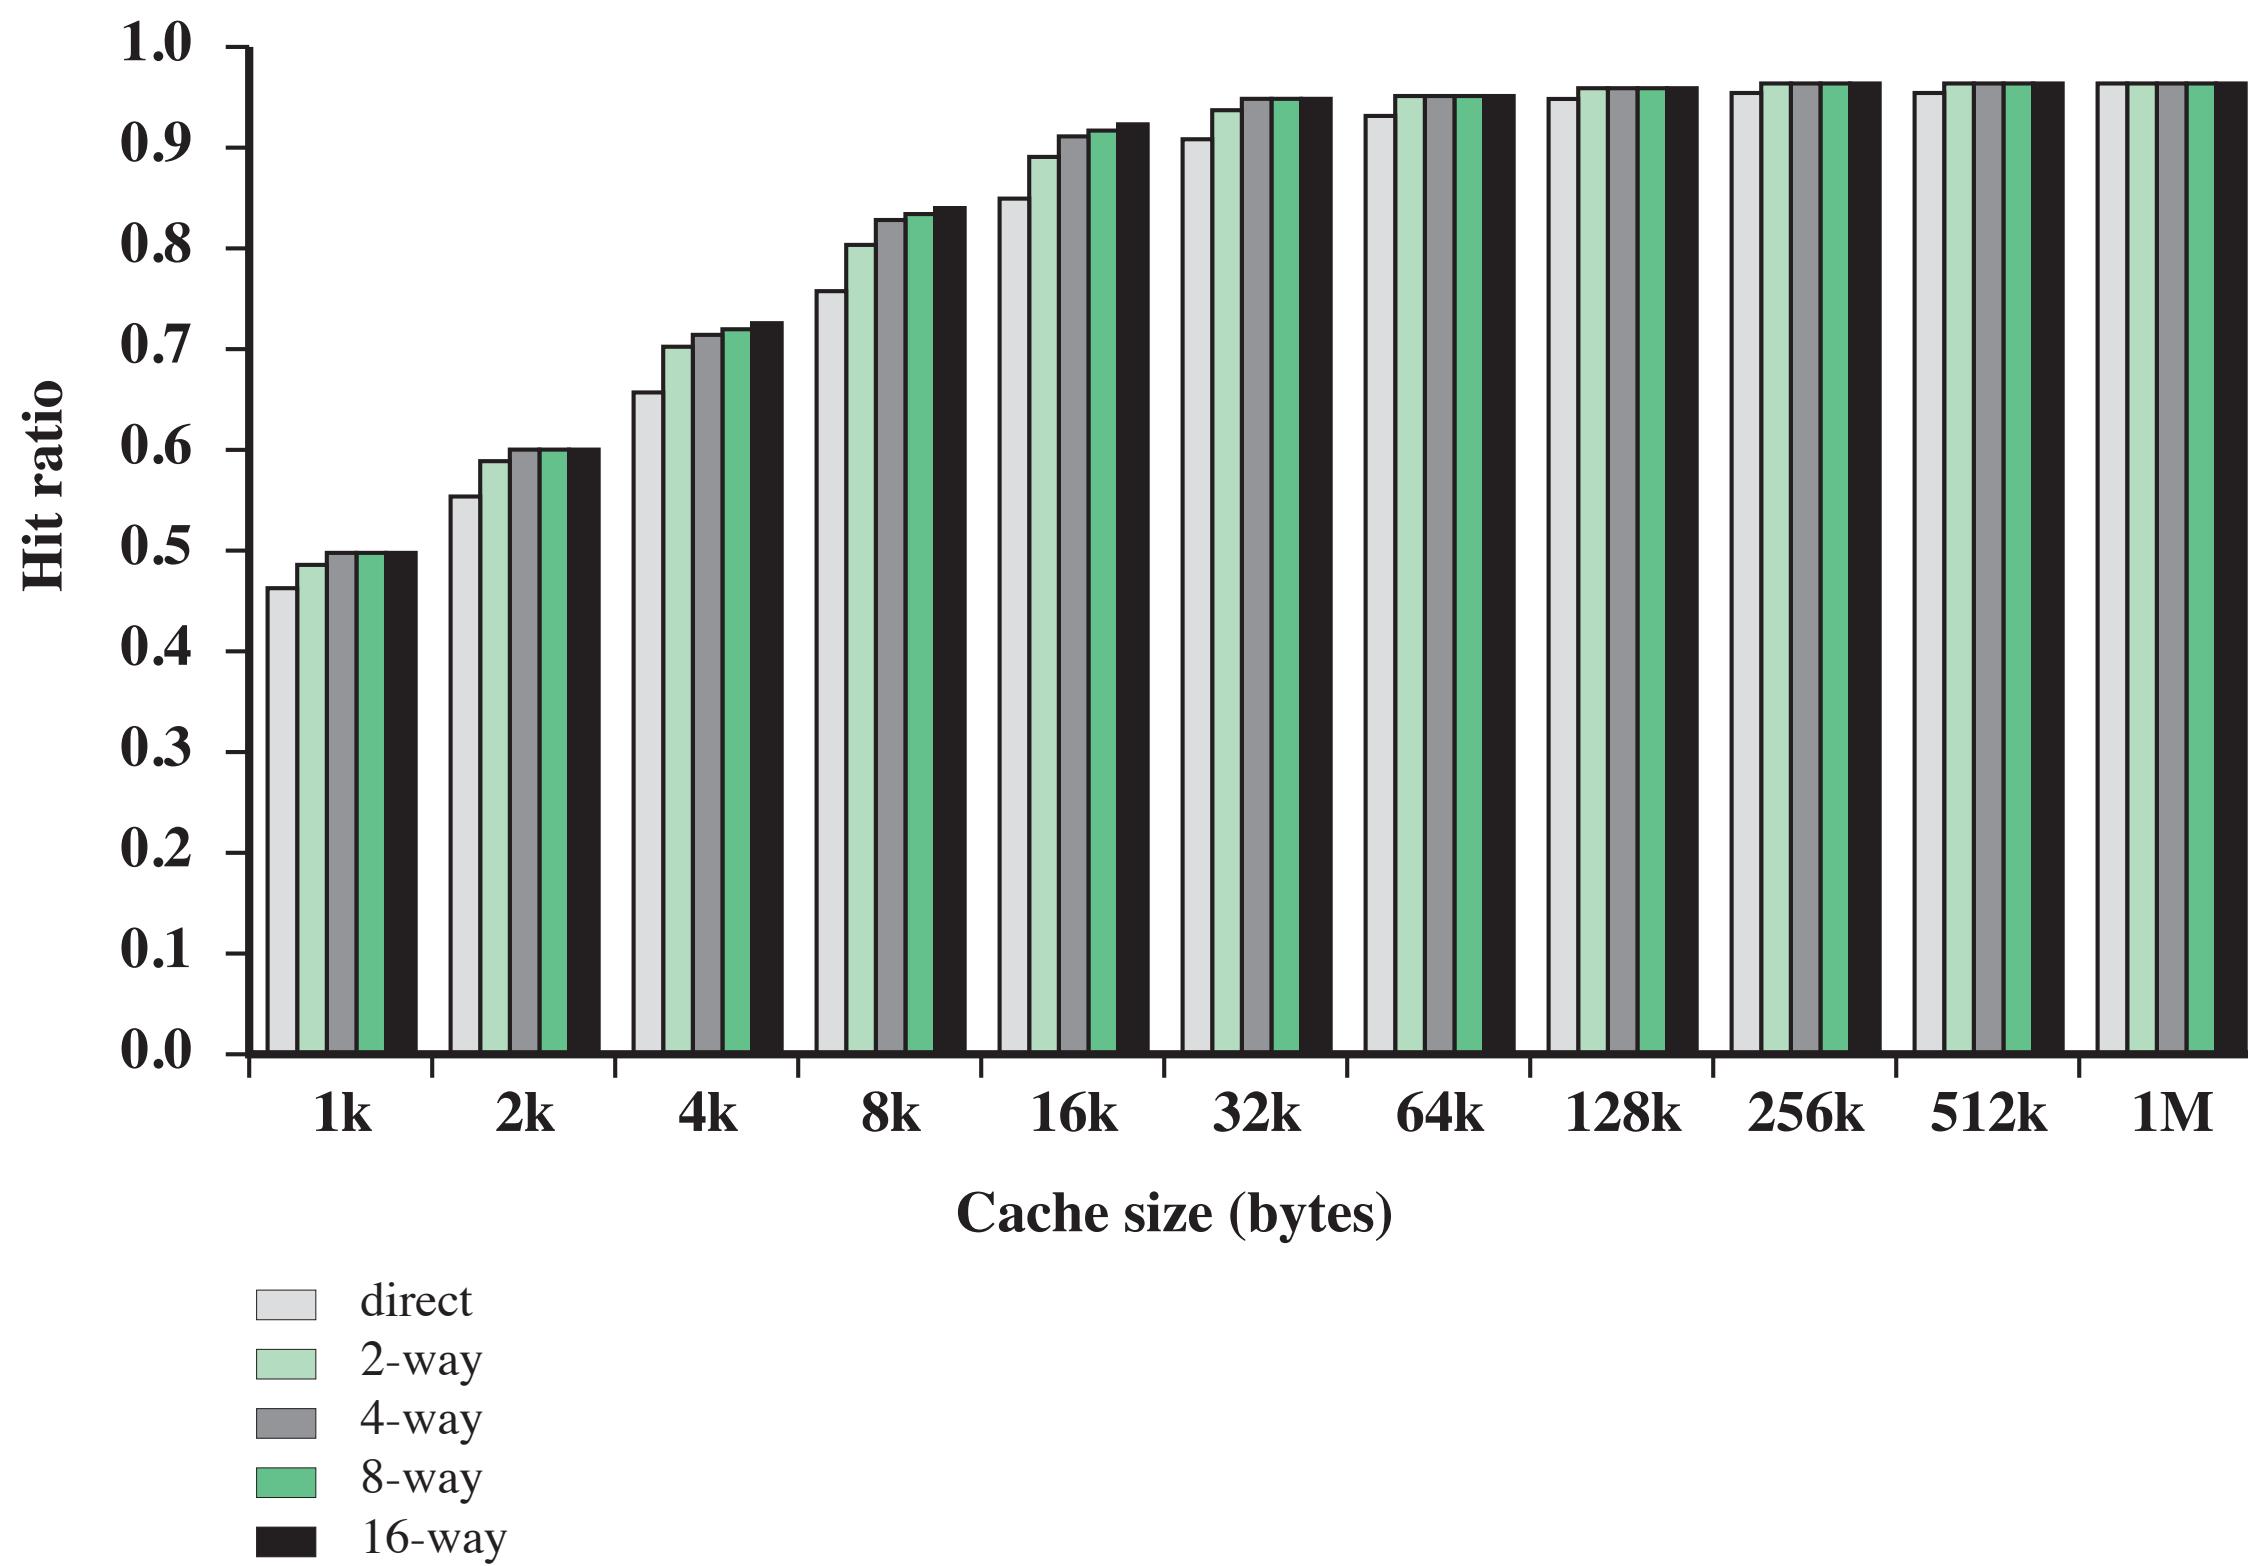
Solution: Reuse data already in cache as much as possible and prefetch data from main memory into cache before it is needed, masking load times

L1 Cache Size and Complexity



Reproduced from: W.Stallings, "Computer Organization and Architecture", 9th Edition, Pearson, 2013

L1 Cache Size and Complexity



Reproduced from: W.Stallings, "Computer Organization and Architecture", 9th Edition, Pearson, 2013

Memory Cell Organisation

Semiconductor Memory

Most common form of main store

RAM (Random Access Memory)

Two main technologies

Static RAM (SRAM)

Dynamic RAM (DRAM)

SRAM use a flip-flop as storage element for each bit

DRAM for each bit, use the presence or absence of charge in a capacitor to denote a 1 or 0

Capacitor charge leaks away over time - requires periodic refreshing (more complex) but is cheaper than SRAM so is more commonly used

Static RAM (SRAM)

SRAM stores data using configurations of flip-flops and logic gates

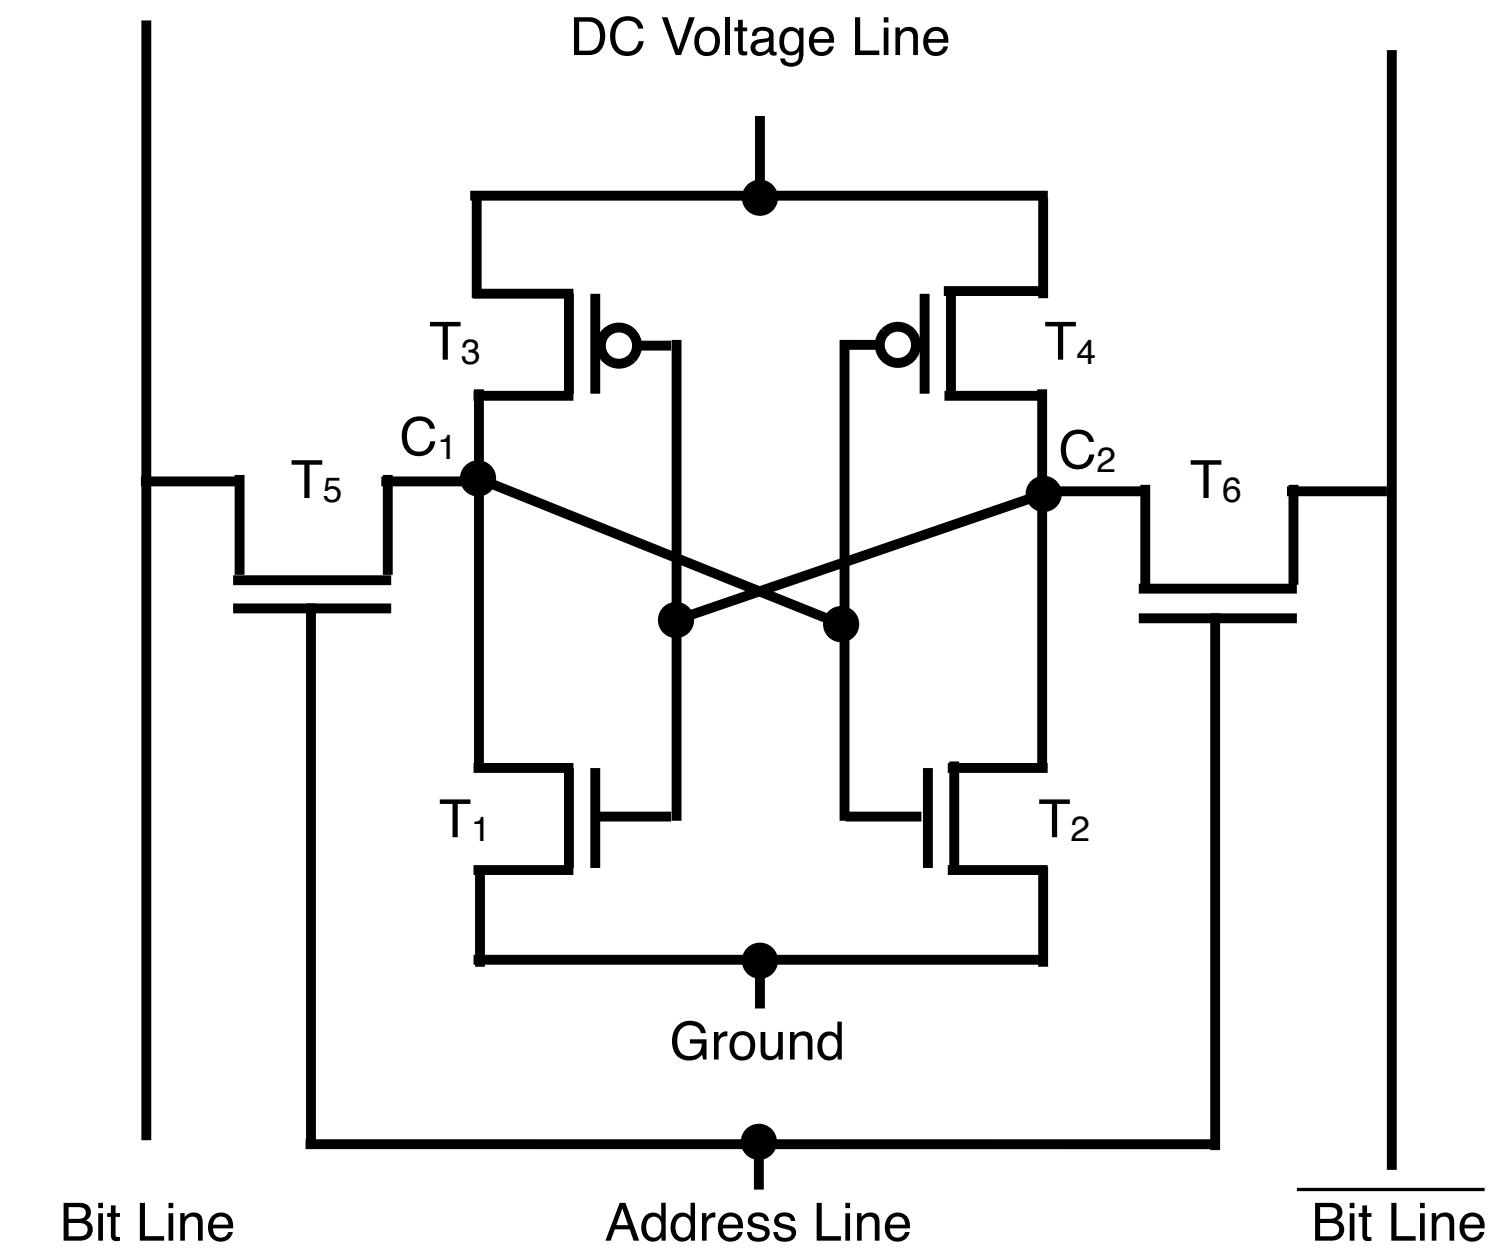
Memory cells hold data as long as power is supplied

Typical SRAM cell uses four cross-coupled transistors to establish a stable logic state

When storing logical 1, C1 is high and C2 is low (T1 and T4 off, T2 and T3 on)

When storing logical 0, C1 is low and C2 is high (T1 and T4 on, T2 and T3 off)

Address line controls two transistors (T5 and T6), where both transistors being on permits read and write



Write operation by applying bit value to Bit Line, which sets T1, T2, T3 and T4 to the appropriate bit value

Read operation by reading Bit Line

Dynamic RAM (DRAM)

DRAM memory cells store data as charge on capacitors

Capacitors naturally discharge, hence DRAM cells require periodic charge to maintain data storage

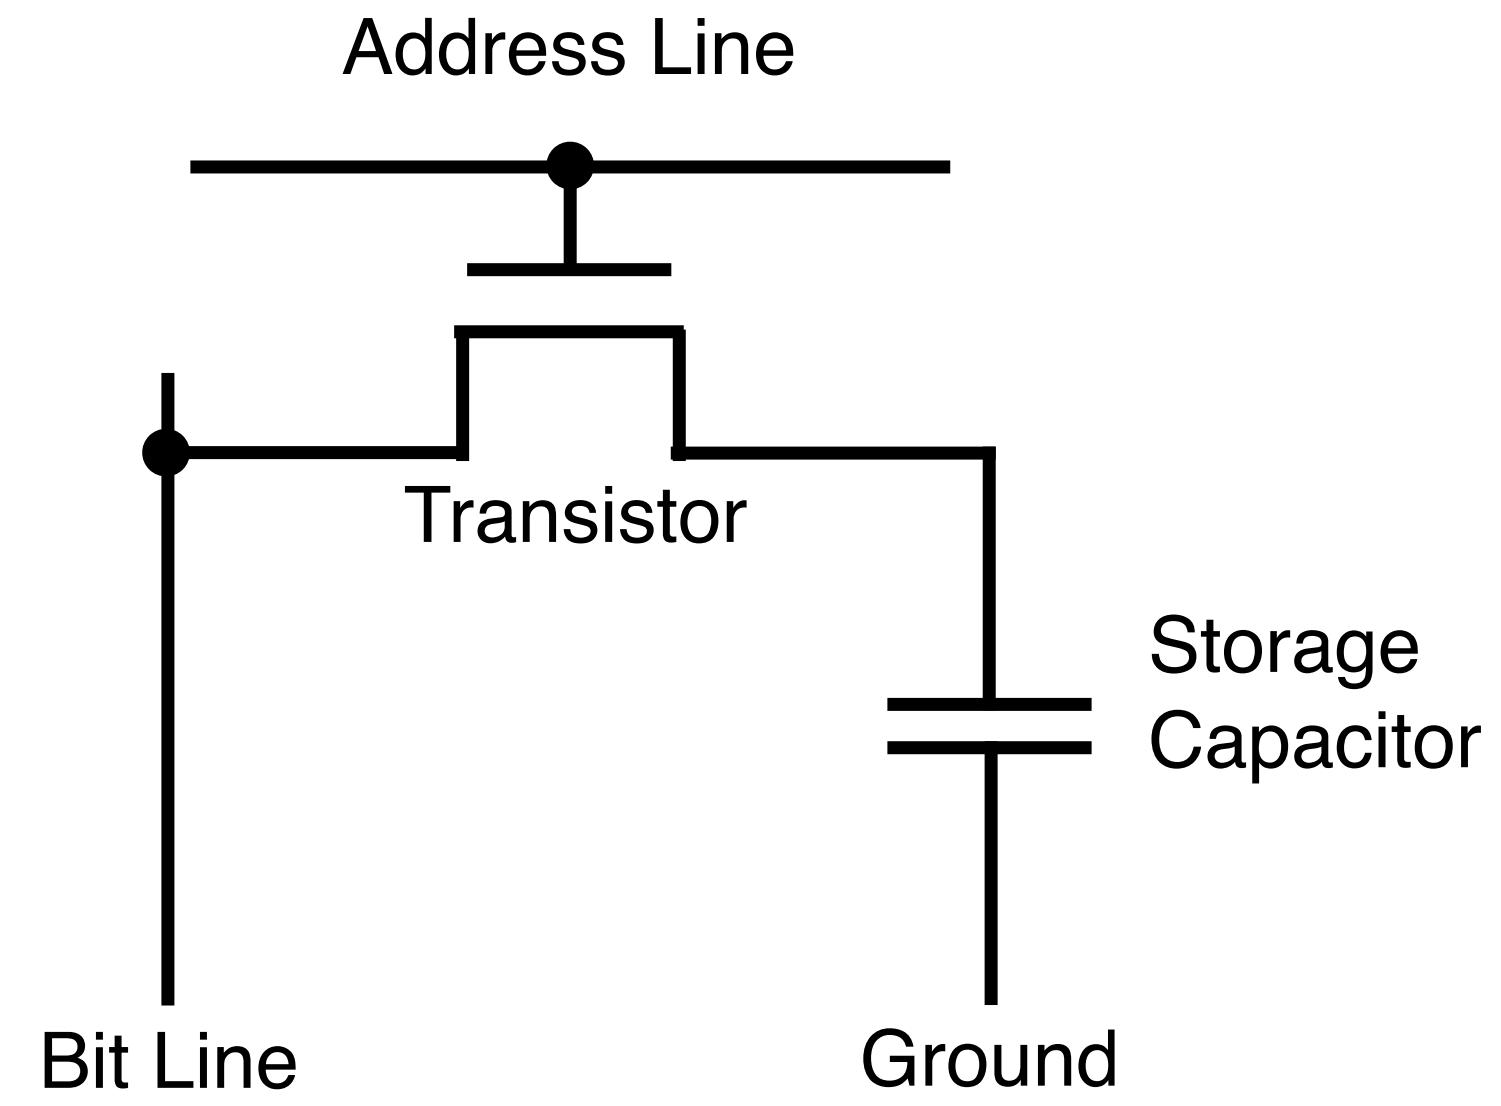
The term 'dynamic' refers to the natural leakage of charge that takes place even when power is applied

Presence or absence of charge is interpreted as a logical 0 or 1

We can appreciate the operation of DRAM by considering a 1-bit DRAM memory cell

Address line is activated when the bit value from the cell is to be read or written

A transistor acts as a switch that is closed (allowing current flow) if a voltage is applied to the address line and open if a voltage is not applied to the address line (no flow)



DRAM Operation

Write operation:

Voltage applied to the bit line (high for logical 1, low for logical 0)

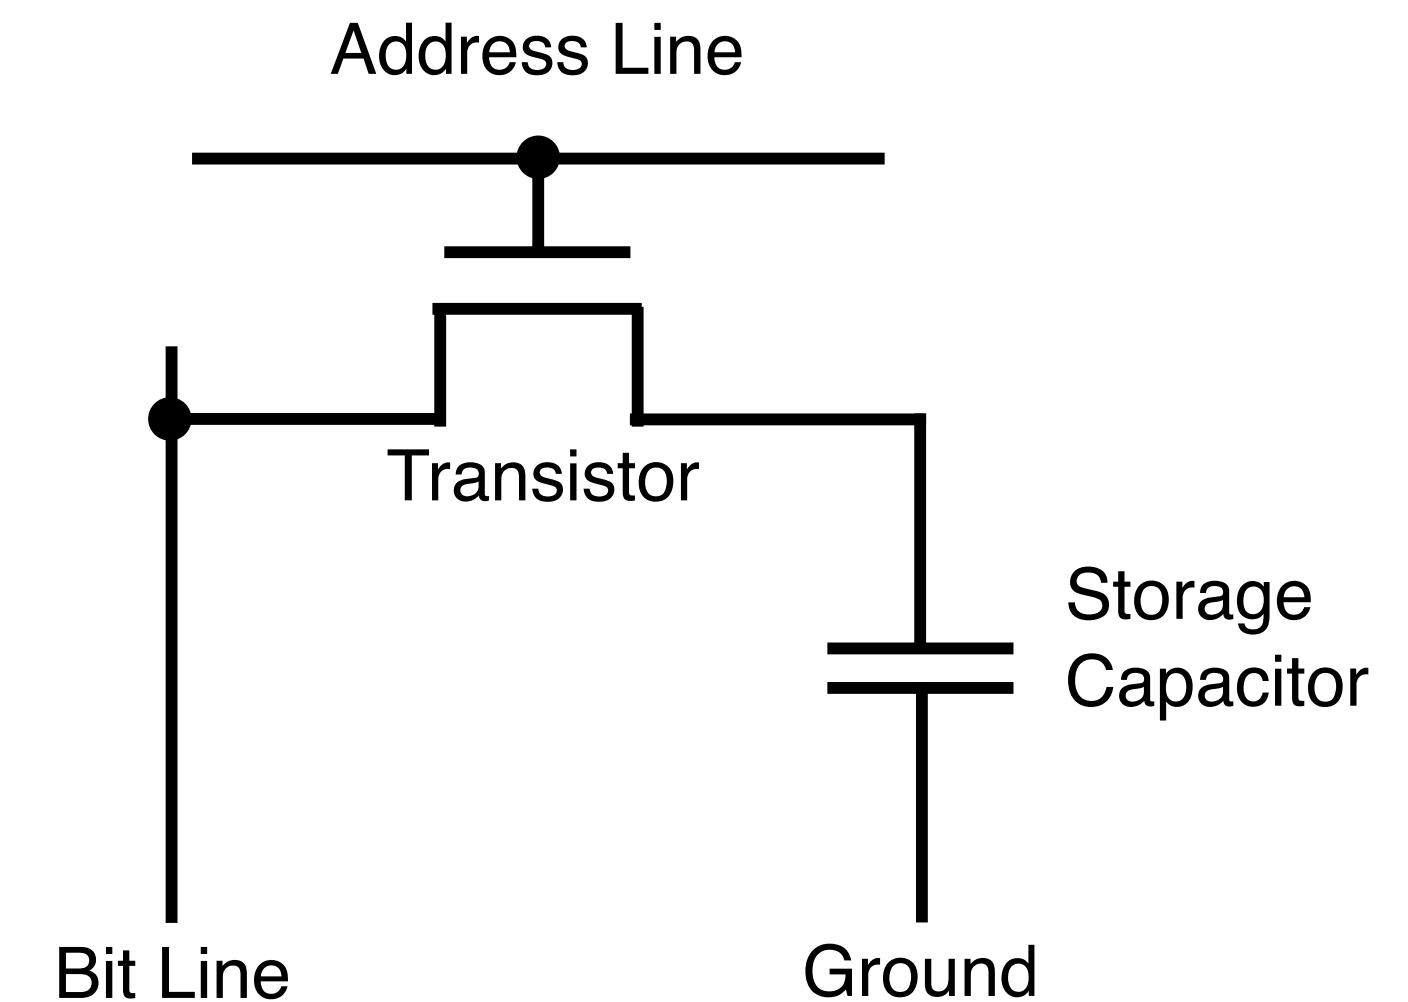
A signal is then applied to the address line, allow a charge to be transferred to the capacitor (data store)

Read operation:

Transistor turns on when address line is selected, allowing the stored charge to be fed out to the bit line and a sense amplifier

Sense amplifier compared the capacitor voltage to a predetermined threshold value to determine if the cell contains a logic 0 or 1

The read process discharges the capacitor, which means that it must be restored (as well as refreshed) following the read operation



Comparing SRAM and DRAM

Both are volatile - power must be continuously applied

Dynamic memory cells are generally simpler and more compact

Allows greater memory cell density (more cells per unit area)

Cheaper to produce than equivalent SRAM memory

Refresh circuitry incurs one-off cost that is only compensated for by larger memory capacities

SRAM are typically provide better read and write times than DRAM

Cache memory (on and off chip) is often implemented as SRAM, whilst DRAM is commonly used for main memory

Advanced DRAM Organisations

The interface to main memory is a critical performance bottleneck

Motivation for multiple high-speed SRAM caches between a processor and DRAM main memory

To achieve better system performance we can also consider advanced DRAM organisations

Synchronous DRAM (SDRAM)

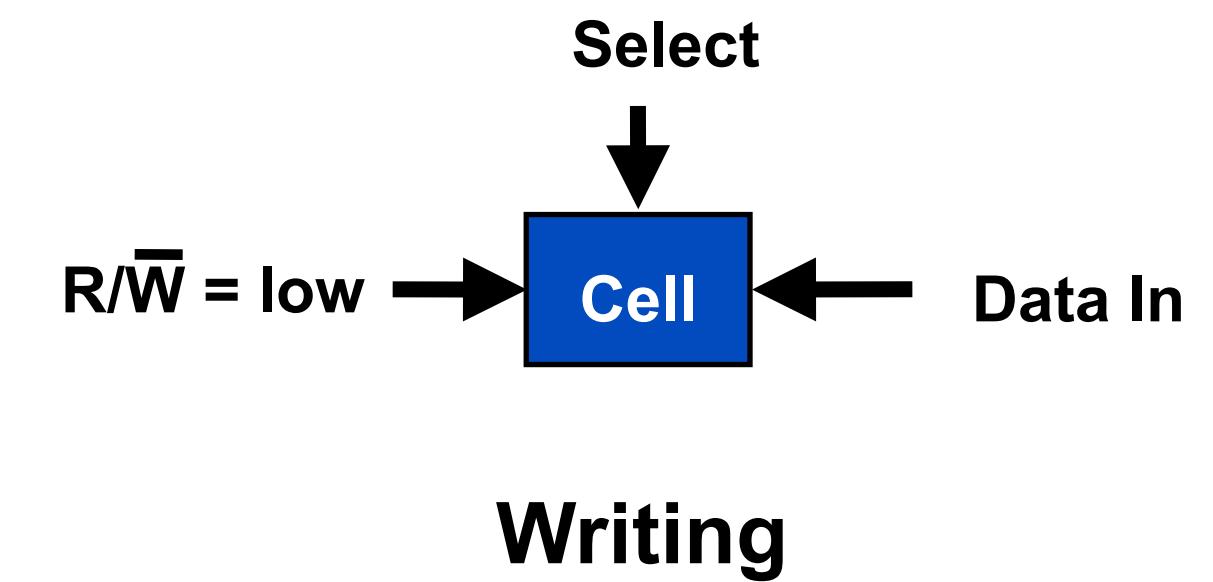
Rambus DRAM (RDRAM)

DDR SDRAM

Cache DRAM (CDRAM)

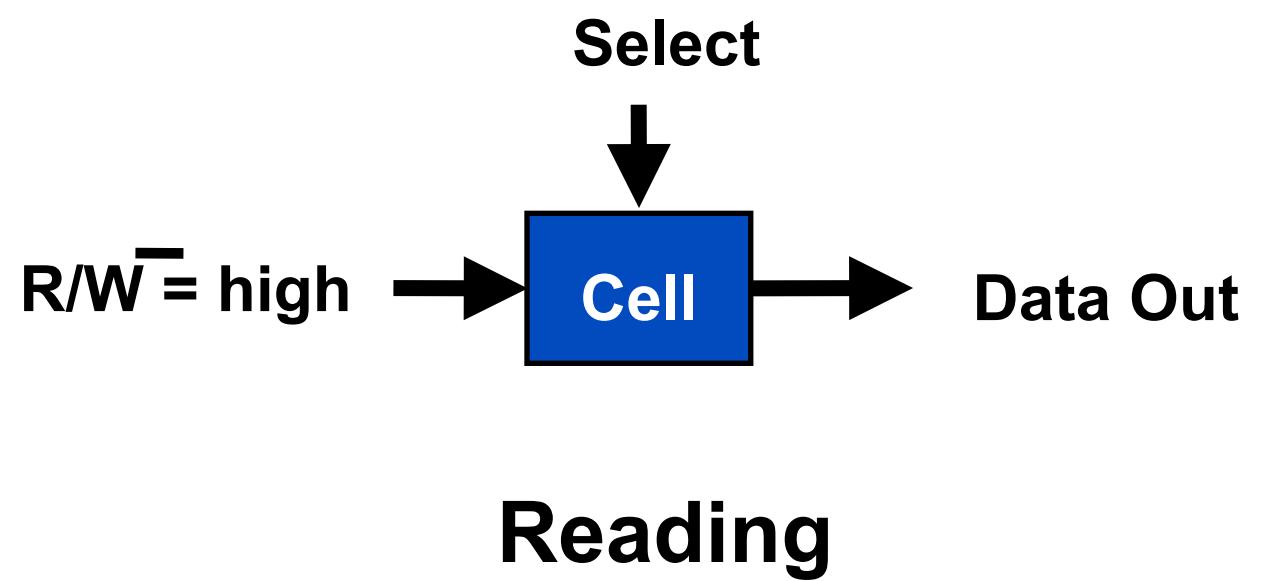
Memory Organisation

Basic element is a memory cell



Two states 1, 0

Capable of being written and read



Select - Enable a particular memory cell

R/W - Indicate read or write

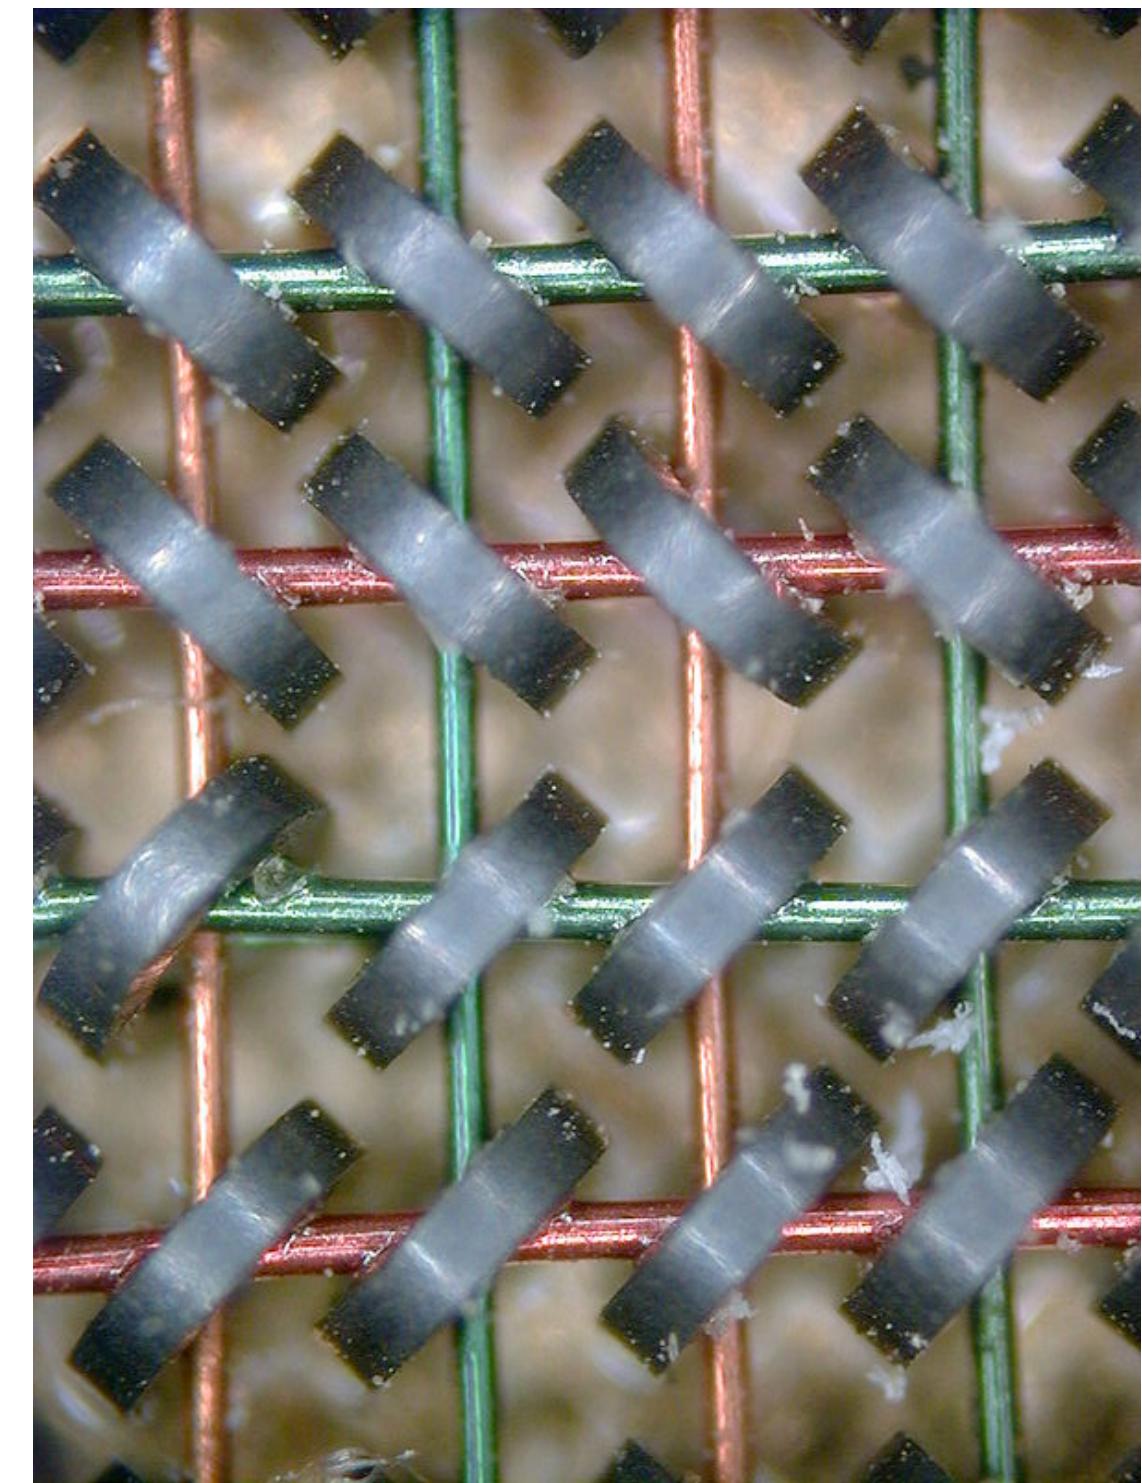
Data - Input when writing or output when reading

Core Memory

An early form of RAM where data is stored through the magnetisation of tiny metal rings – one bit per ring

Rings are organised into a grid array

Individual bits are accessed by energising row and column wires

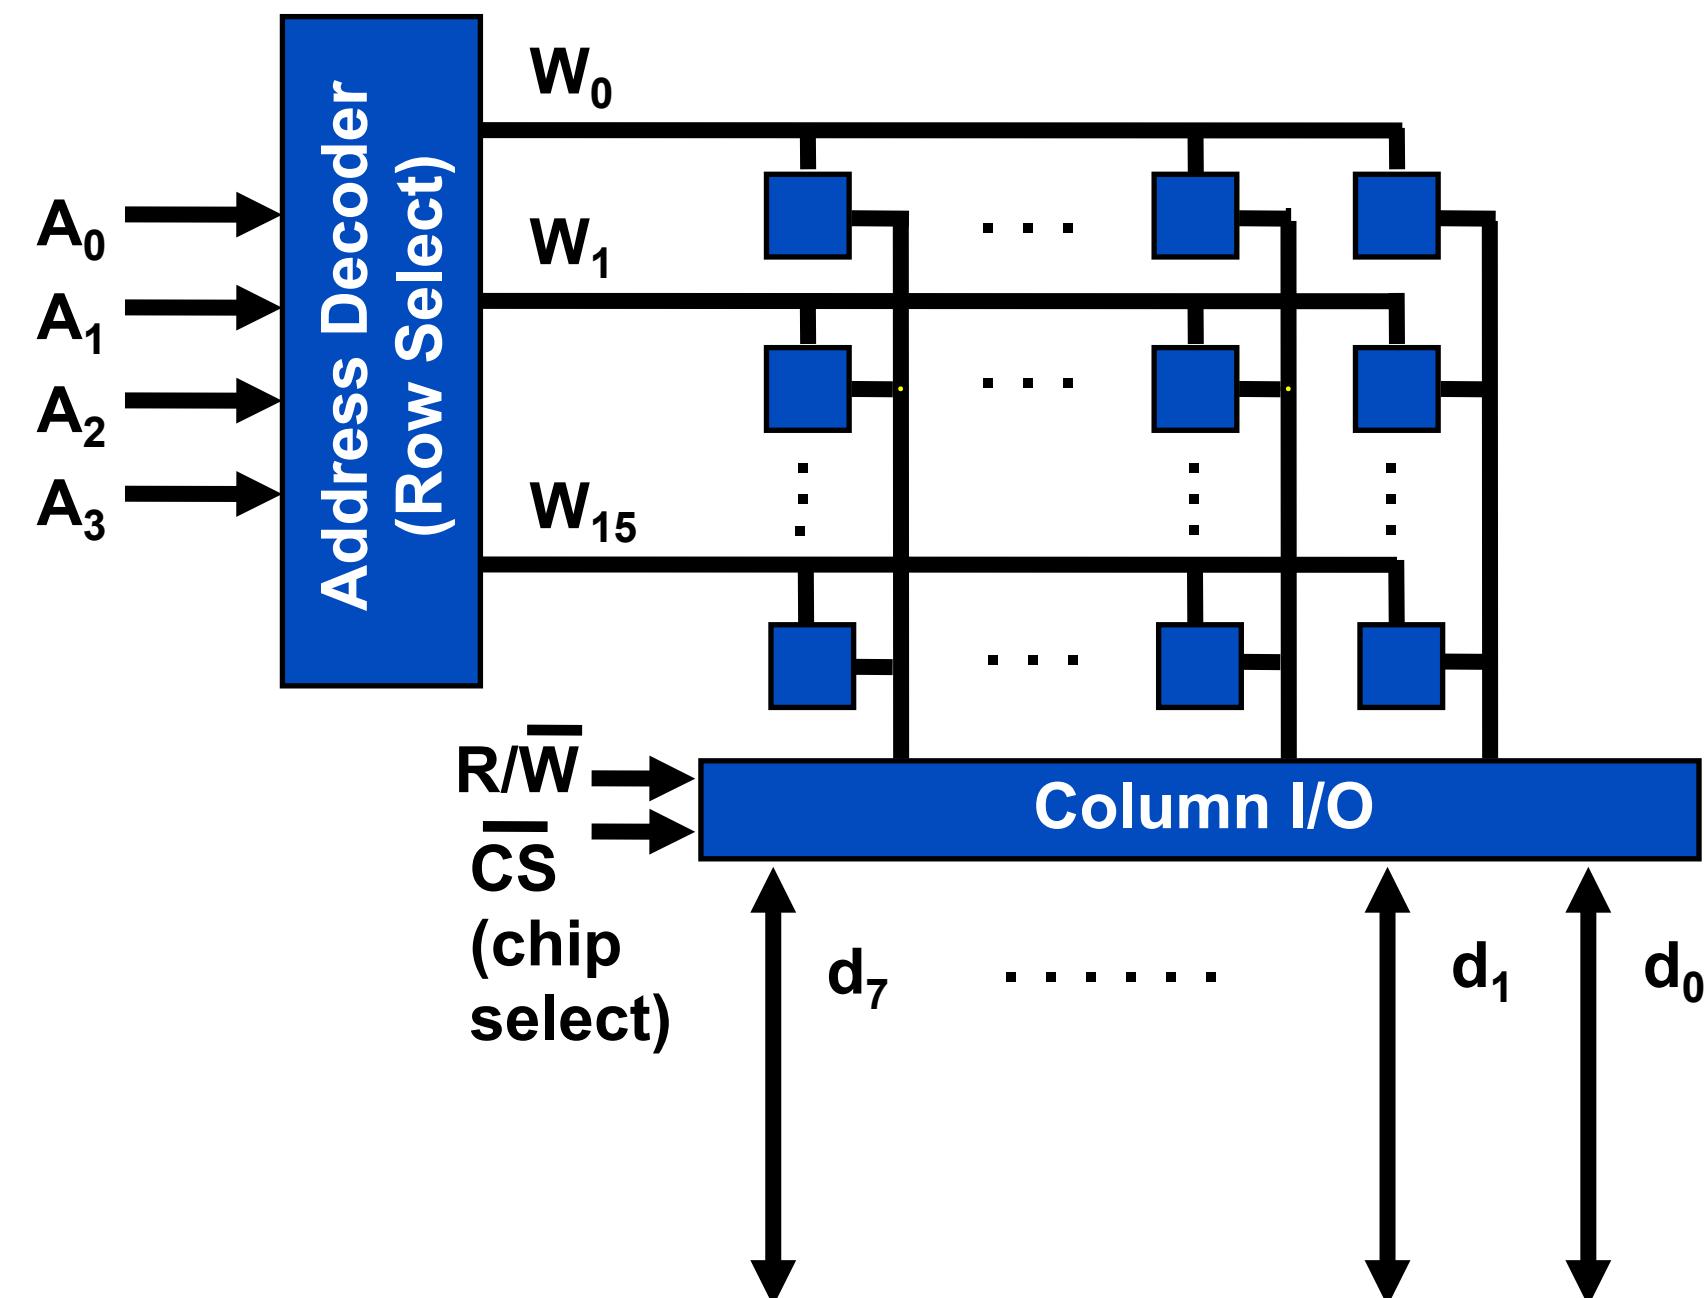


Organisation of a Memory Chip

As you might expect, it is critical to maintain a structured approach to the organisation of on-chip memory

Example of a 16x8
memory integrated
circuit

16x8 denotes
 $W \times B$, i.e., words x bits
16 words, each of 8 bits



Choices in Memory Chip Organisation

IC design for $16 \times 8 = 128$ cells required 4 address inputs ($= \log_2 W$, where W = no of words) and 8 data lines ($= B$, where B = word size)

Consider a 1Kbit = 1024 cells device

Possible to organise as a 128×8 memory cell array

Requires 7 address pins ($\log_2 128$) and 8 data pins

A total of 15 I/O pins, plus power, etc...

Poor Choices in Memory Chip Organisation

Alternative, it is possible to organise as a 1024×1 array

Requires 10 address pins ($\log_2 1024$) and 1 data pin

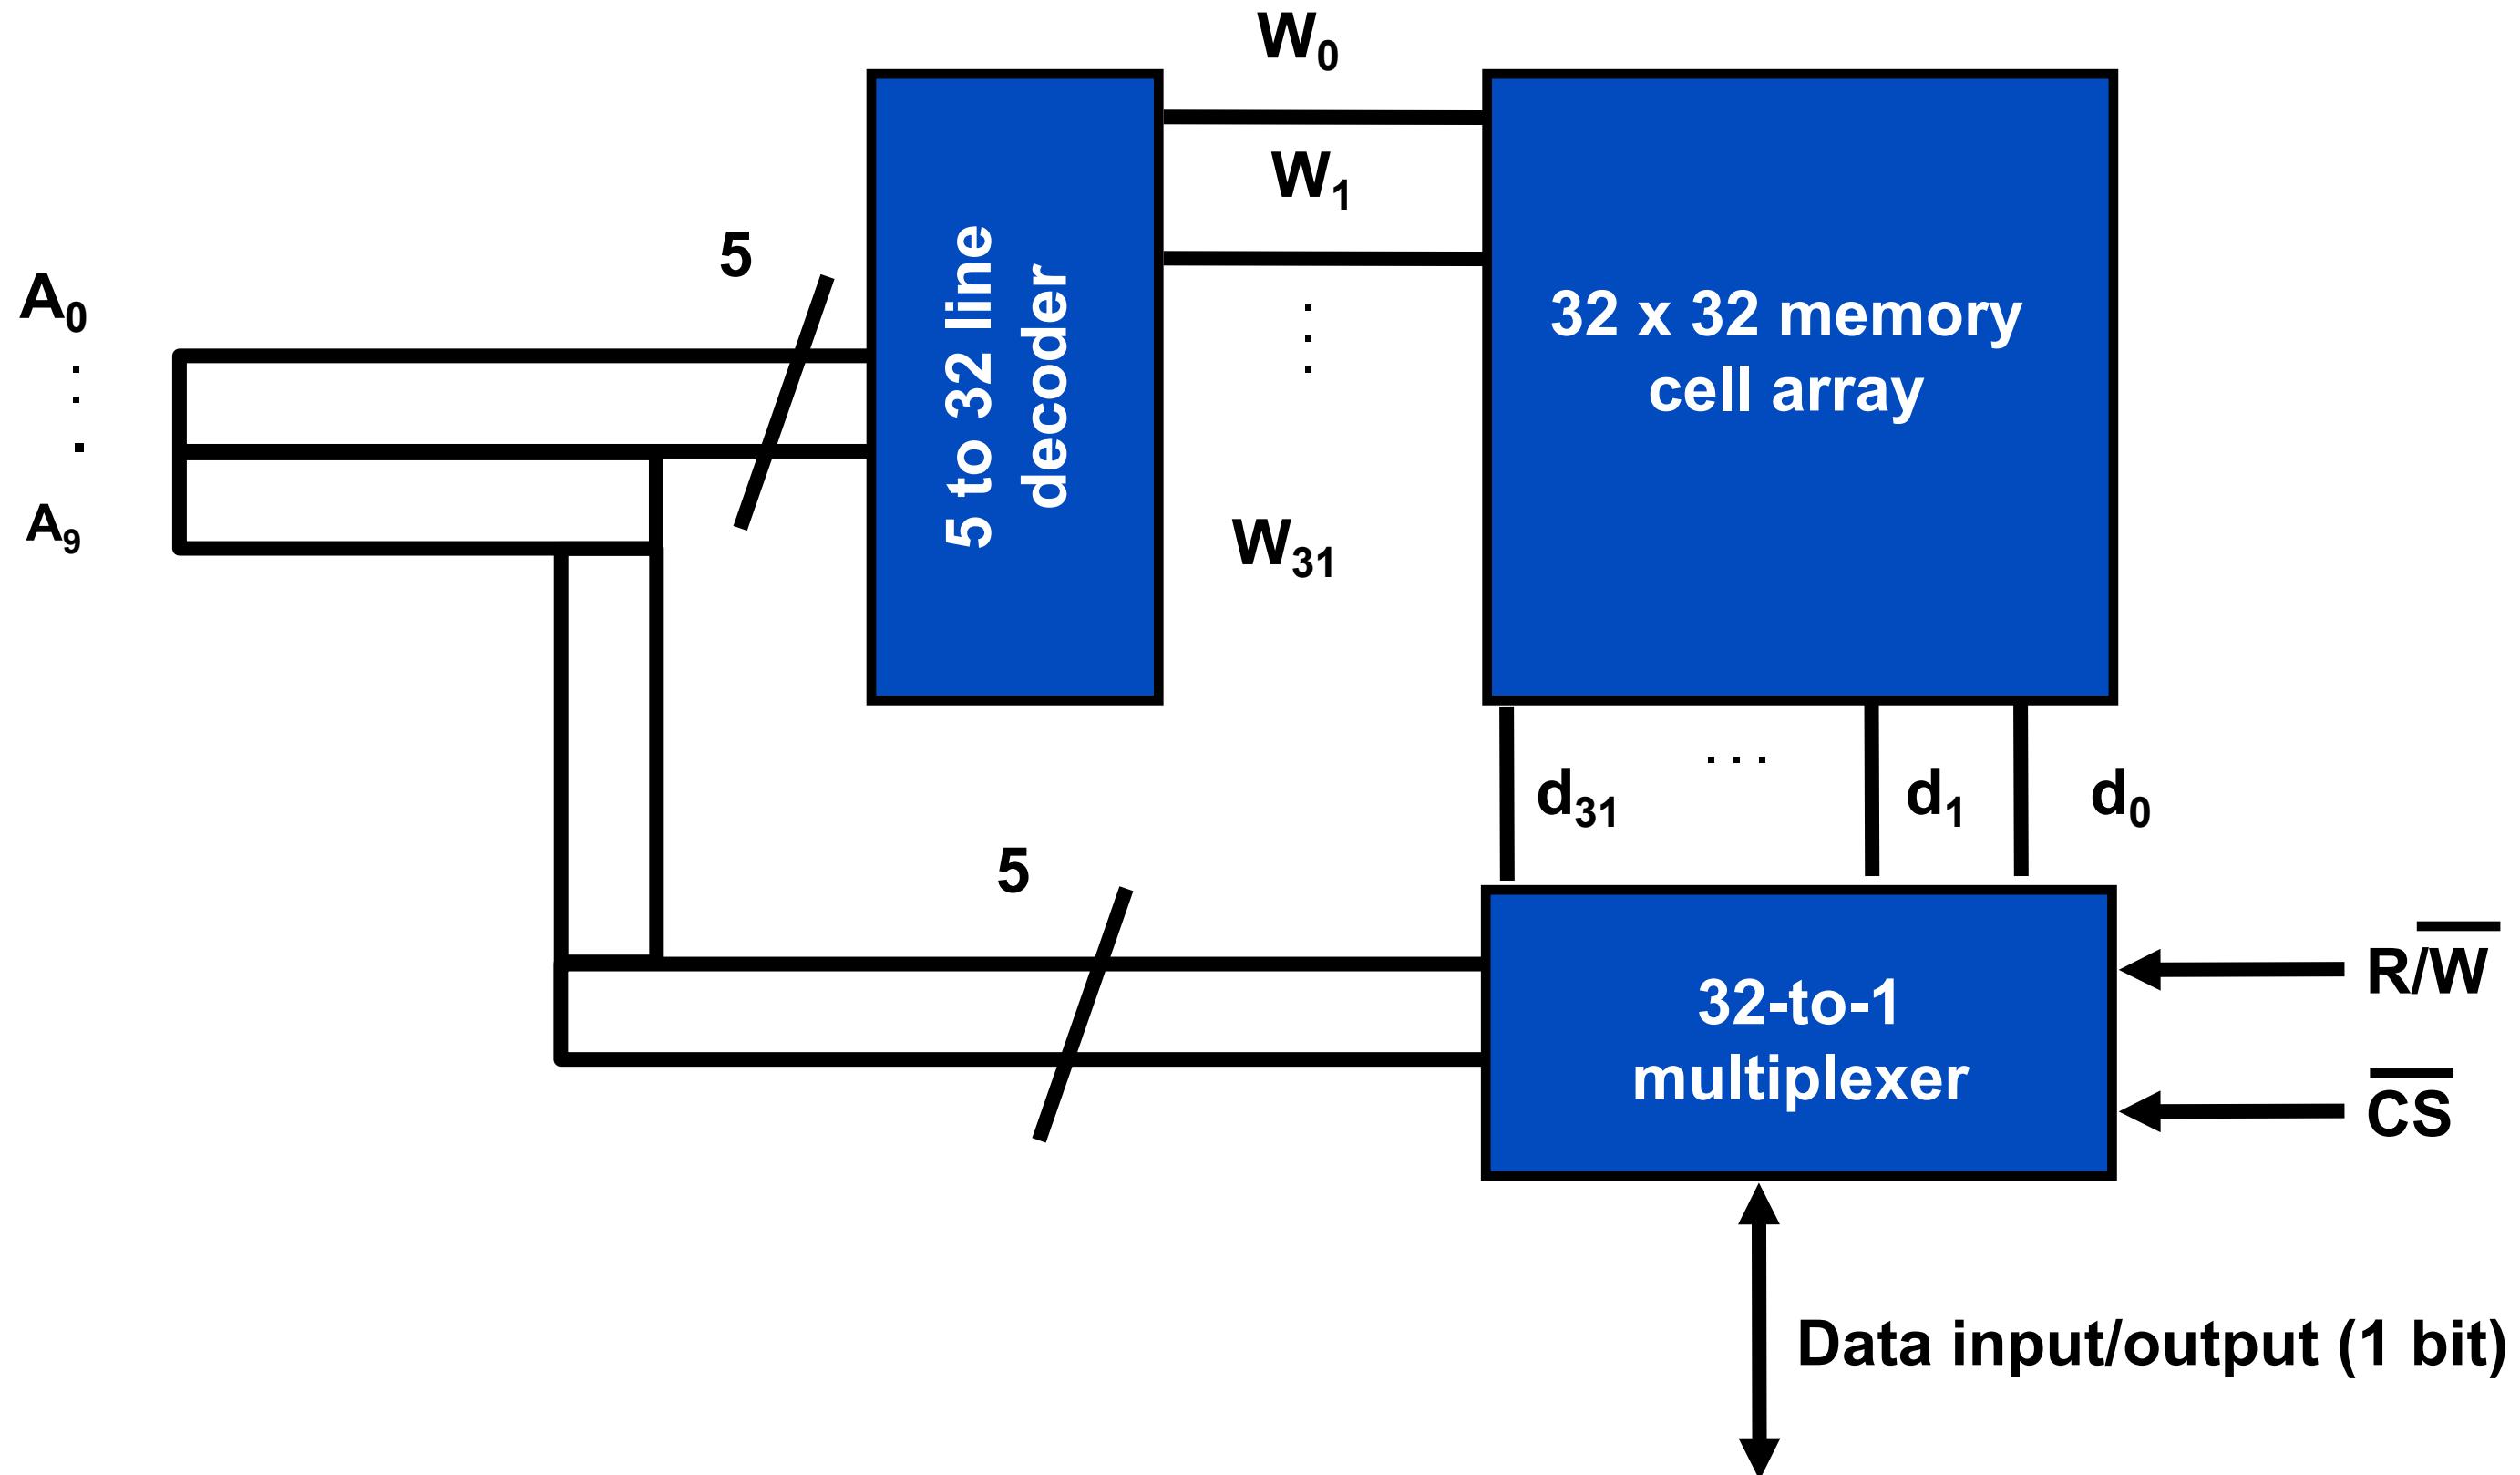
A total of only 11 I/O pins, plus power, etc...

Results in a long, narrow cell array with a large address decoder, which is an inefficient use of space

Minimising Address Decoding Space

Divide the address inputs into 2 parts - row address and column address

This minimises the space required for address decoding and hence maximises space used for memory cells



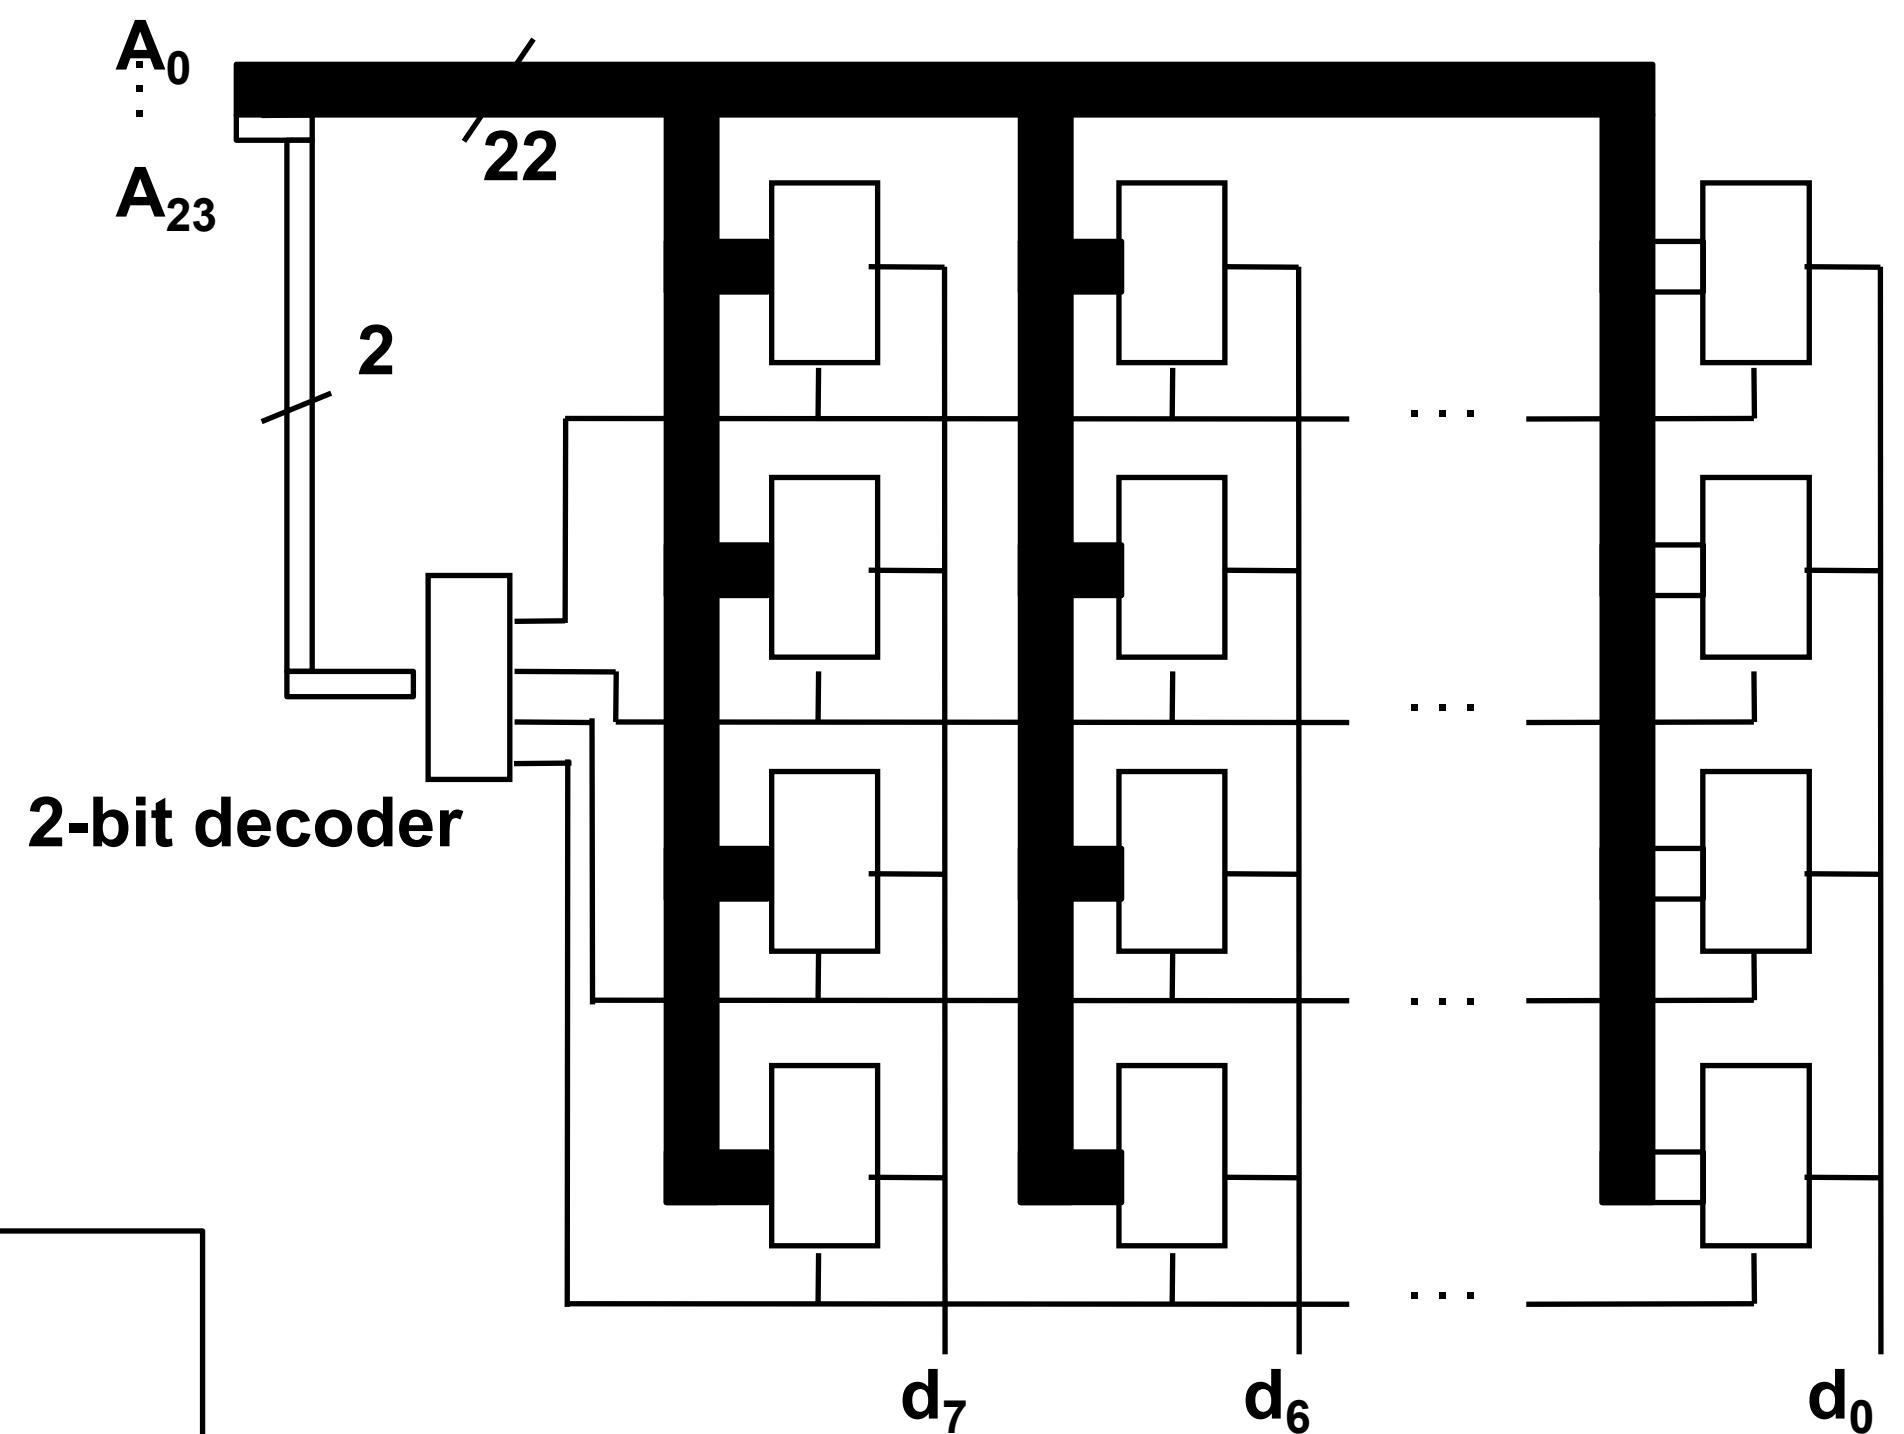
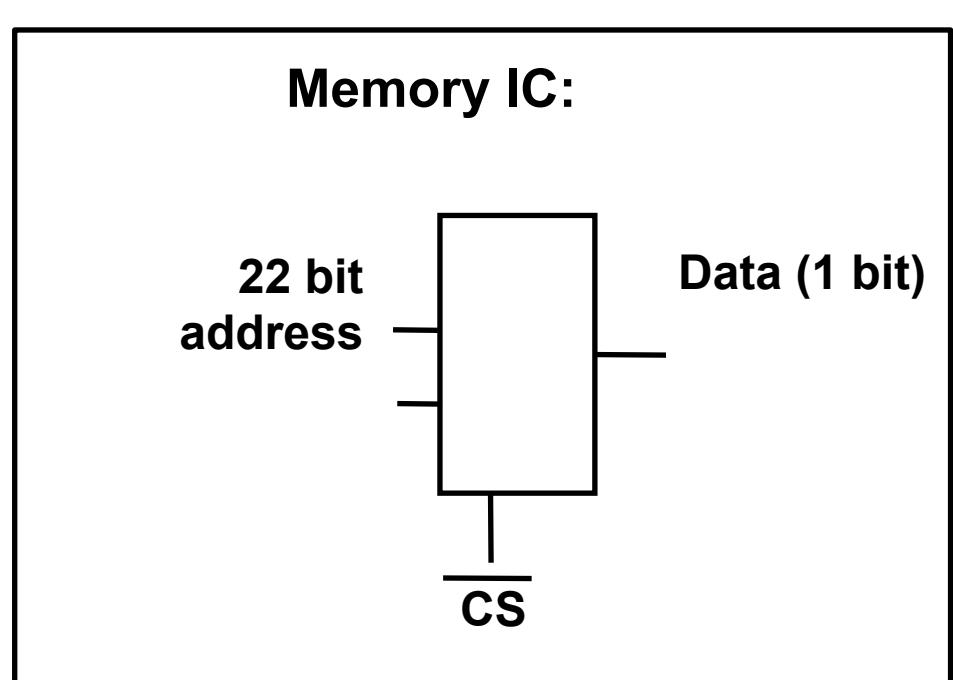
System Memory Organisation

We can use the same idea to organise multiple memory ICs

Here, A_{22} and A_{23} select 1 row of 8 chips, each providing one bit of the 8-bit word

Each chip is $4M \times 1$ (22 address pins, 1 data pin)

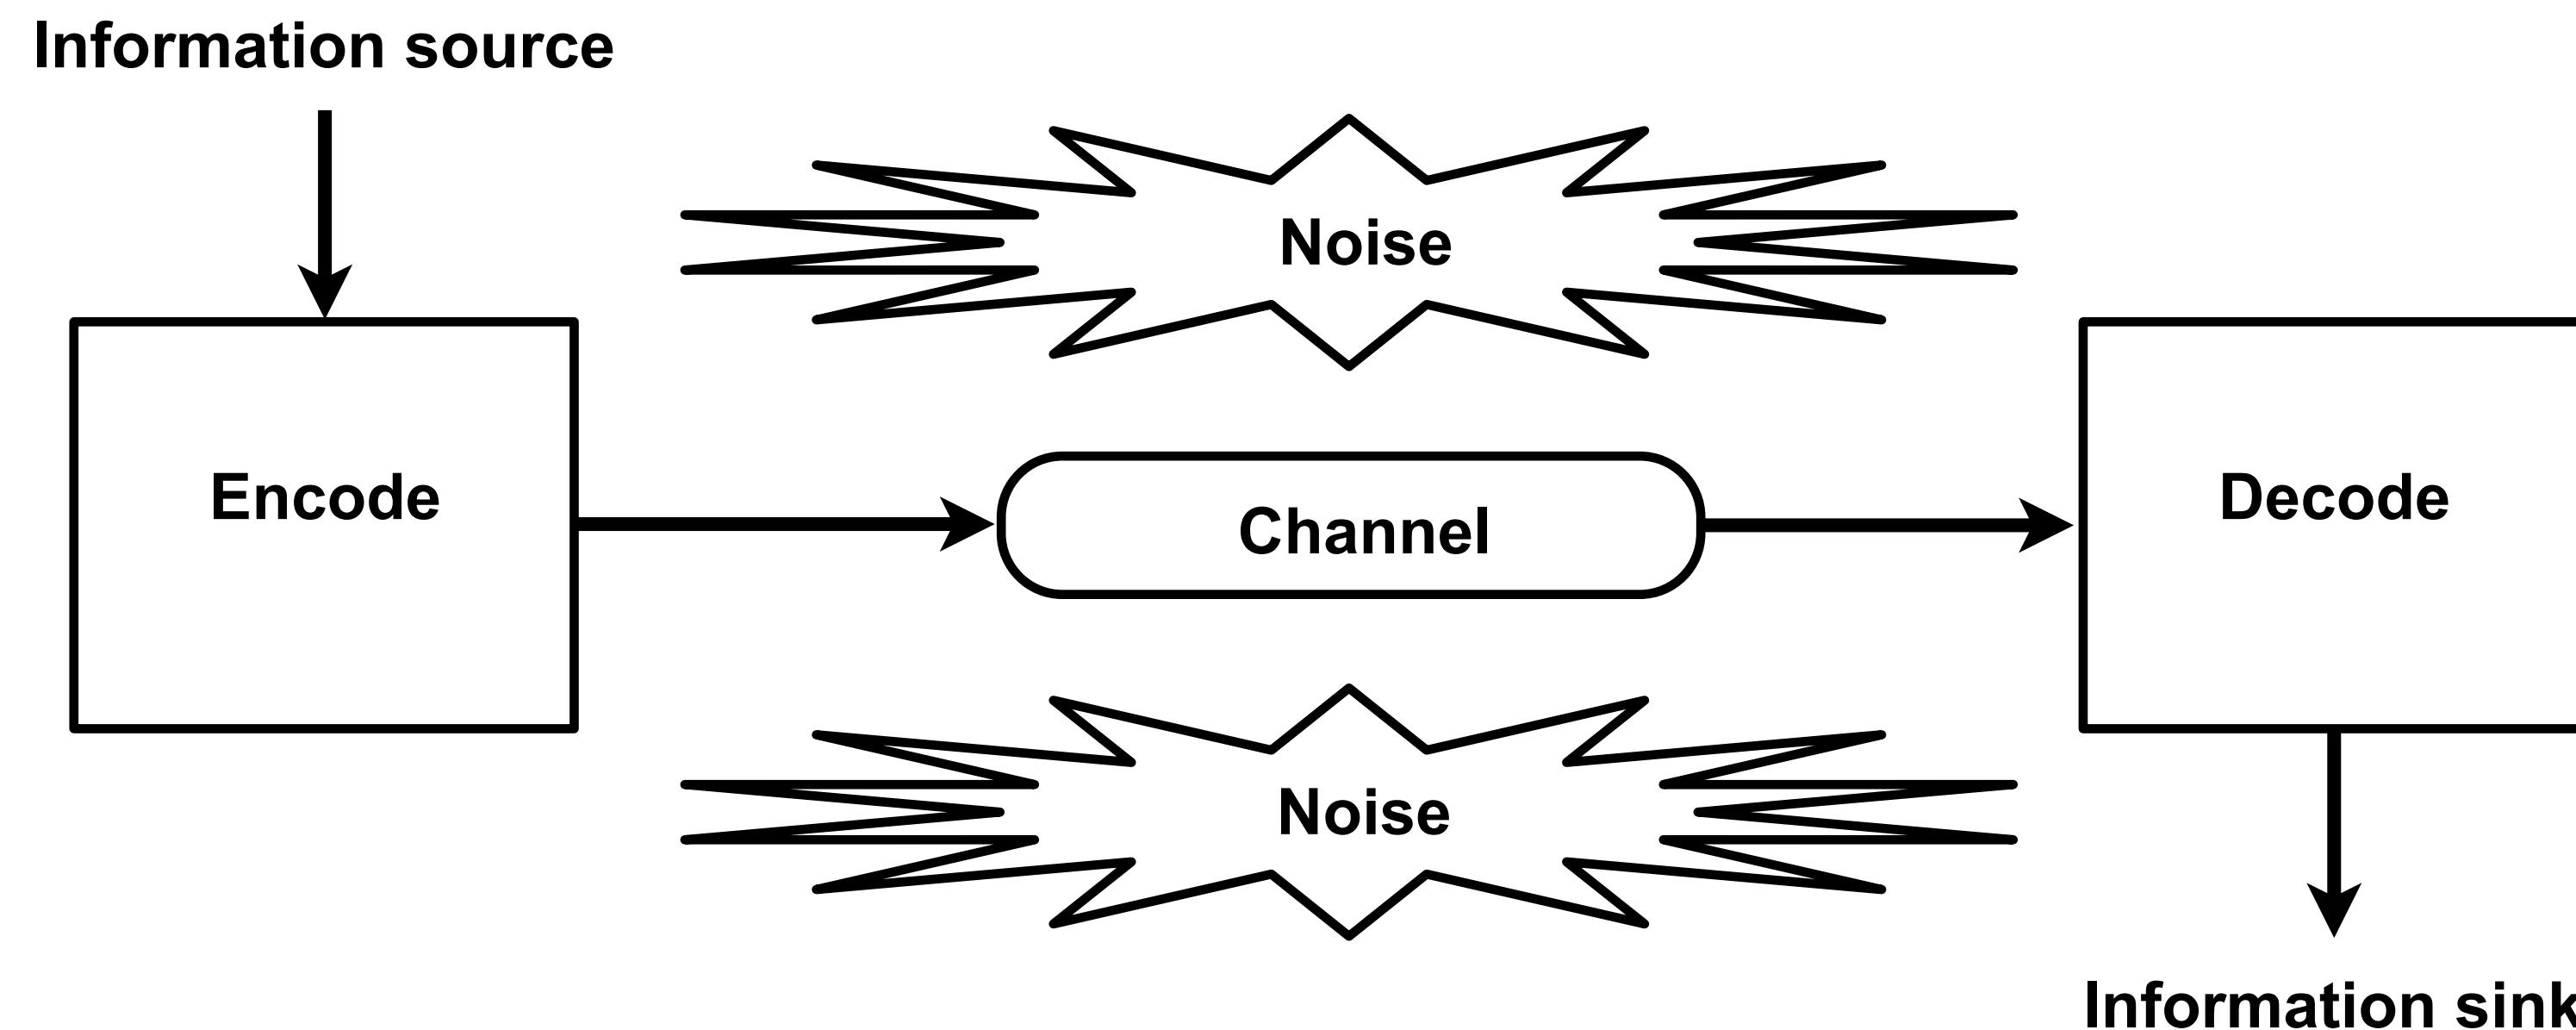
System has 4 rows, 8 chips per row, making the system total a $16M \times 8$ memory



Detecting and Correcting Errors

Error Detection and Correction

Errors occur within a computer system, e.g., in system memory, and in the communication between systems, e.g., in the transmission of messages



Noise

Noise is unwanted information

It comes in various forms, but is always present, and is one of the limiting factors in computer systems

Noise arises from the physical properties of devices

Thermal noise

Noise of electronic components

Noise of transmission circuits

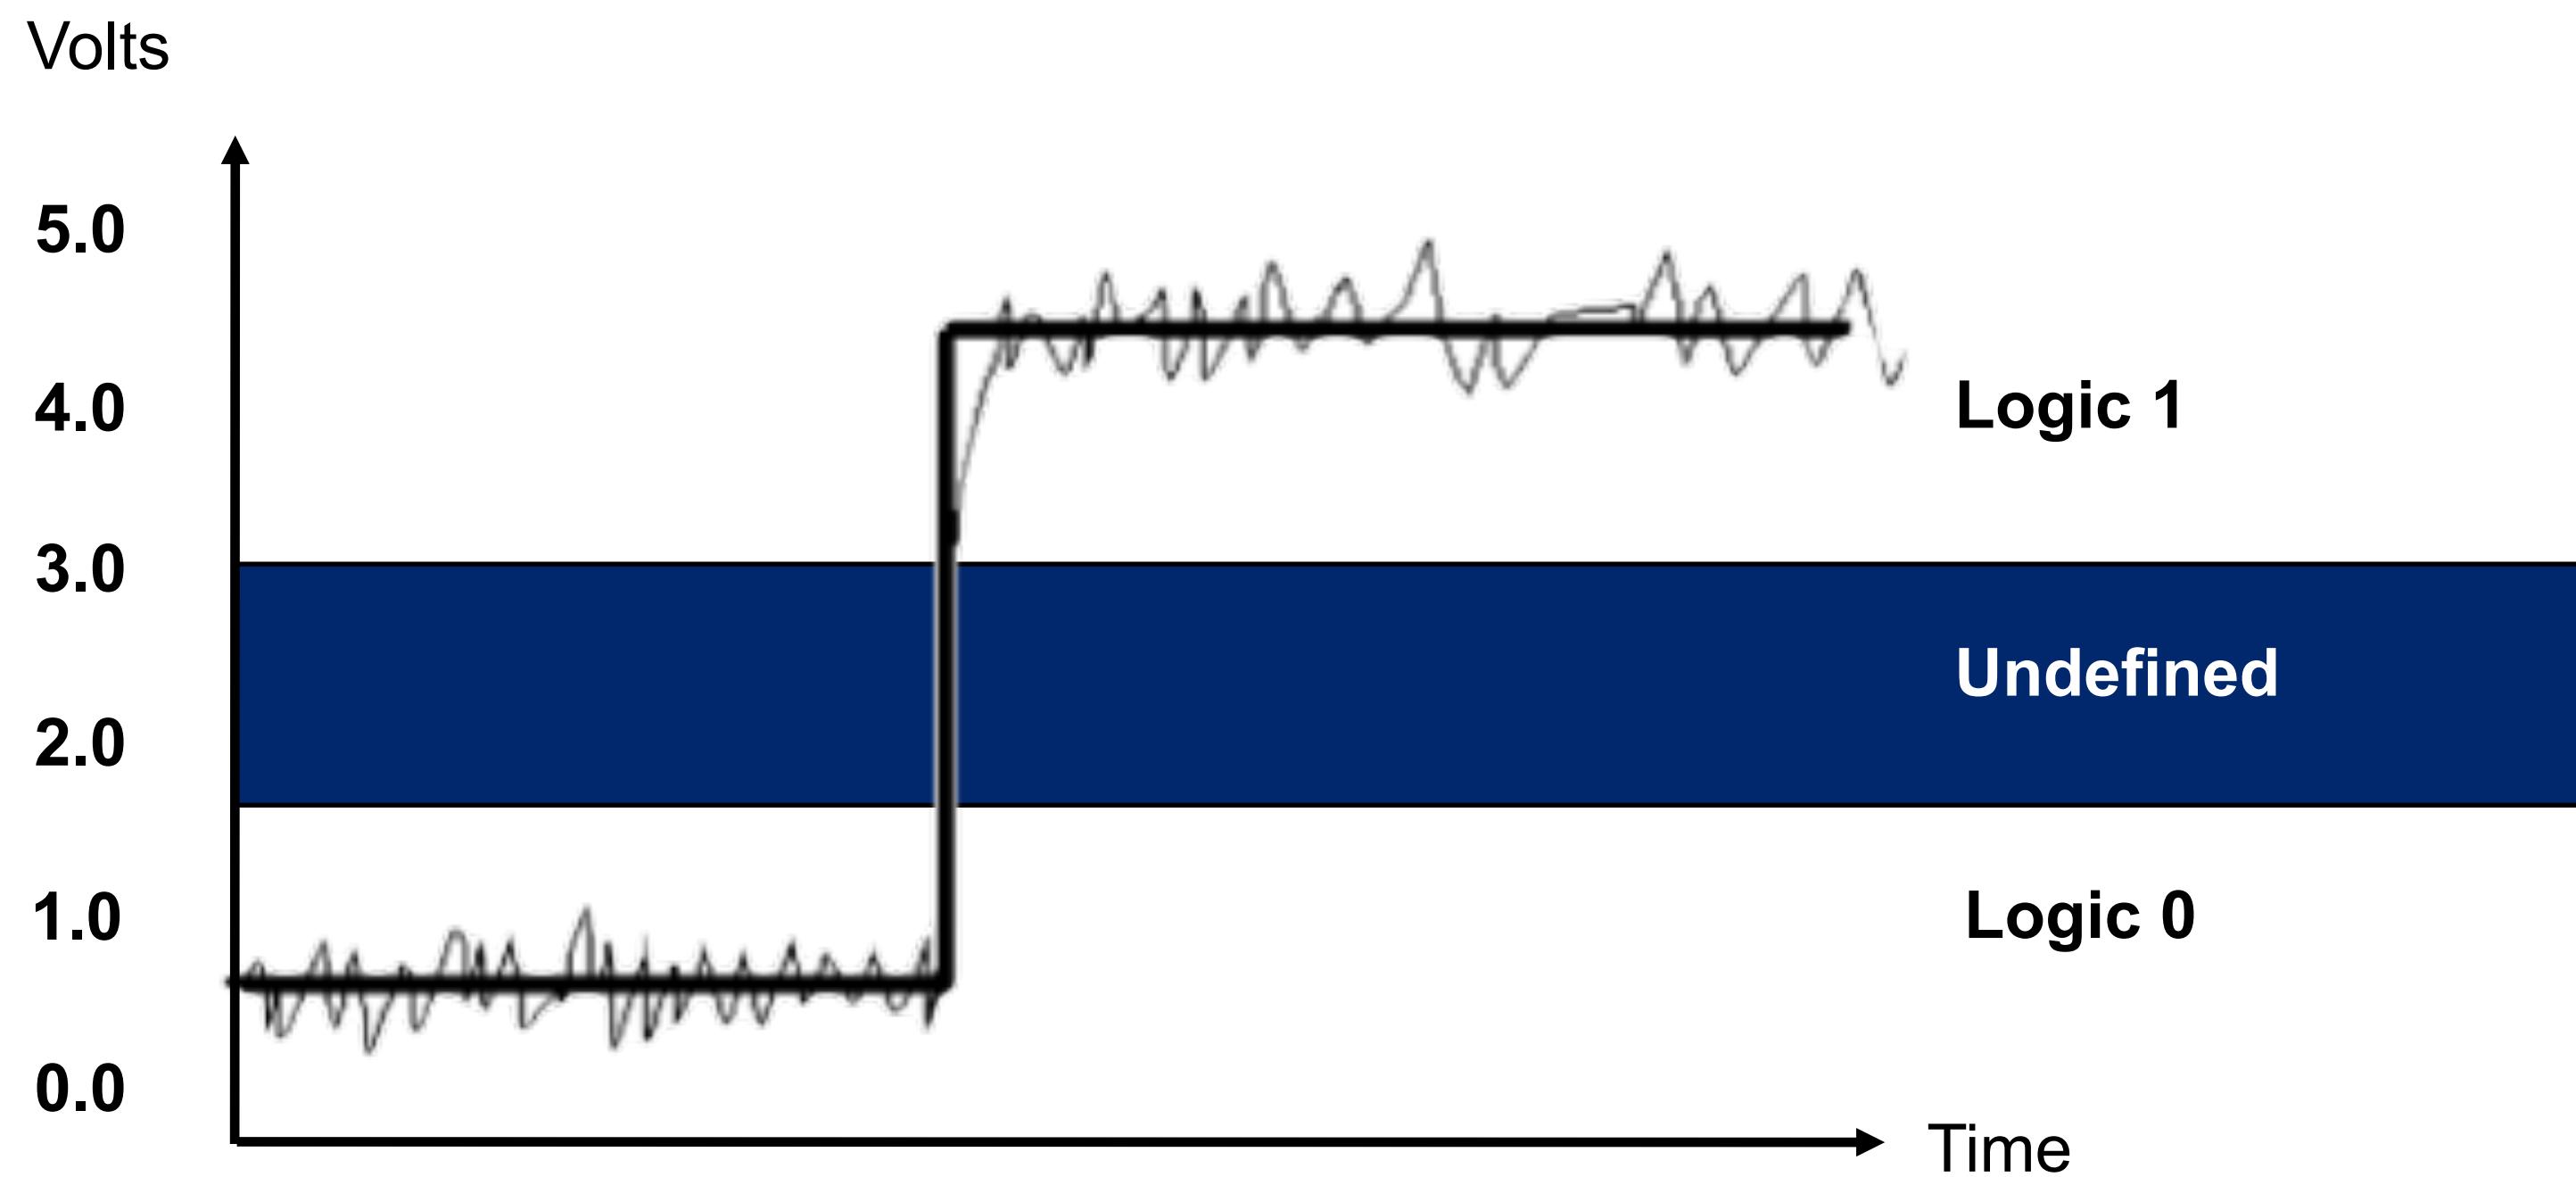
Magnetic media also have a classic form of noise due to random alignment of magnetic fields.

Decreased area to store a bit means noise gets worse, making errors likely

Noise and Digital Logic Devices

Digital logic gives us a high degree of noise immunity

Immunity collapses once the noise reaches a certain magnitude - our aim is to not let this happen



Detecting Single, Isolated Errors

These are considered to occur at random, usually due to noise

We could send the message three times and take a vote

Very expensive

Observation that, if the probability of an error in the channel is low, the probability of two errors close together is even lower

Thus, we can add a “parity” bit to the message every so often, which “summarises” a property of the message that we can check to determine whether the message has been altered

Much cheaper and adequate in many situations

Parity

Parity adds an extra bit

There are two types of parity system.

Even parity system - The value of the extra bit is chosen to make the total number of logic 1s an even number

Odd parity system - Make the total number of logic 1s odd

Parity bit value can be computed in software or hardware

7-bit ASCII for 'A':

A = 0100 0001

Using even parity system: A = 0100 0001

Using odd parity system: A = 1100 0001

Parity in Software

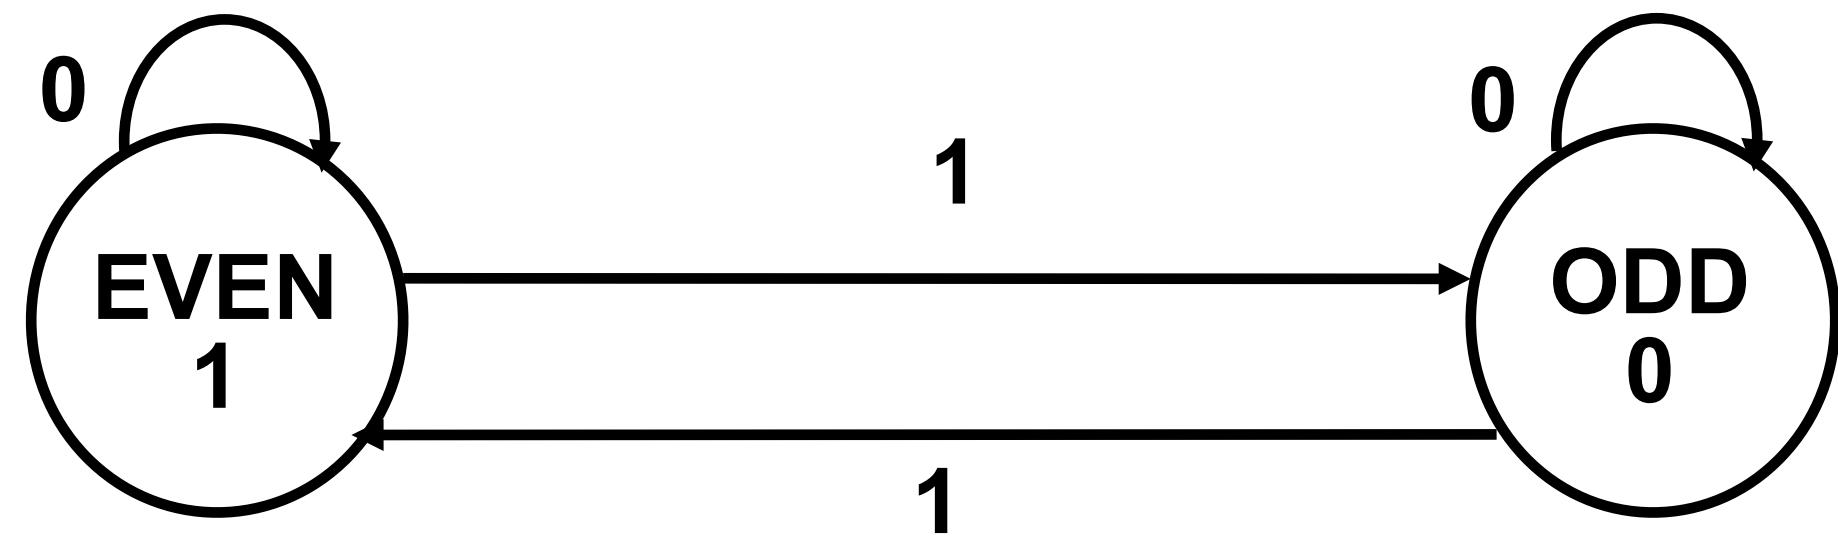
Odd parity system - Start at EVEN

Even parity system - Start at ODD

Each 1 in the message causes a transition of state

Final state at the end of the message is the resulting parity bit

Parity may be computed using a two state Finite Automaton



Parity in Software - Example

For example, try sending 110 using an even parity system

Start with the ODD state, then

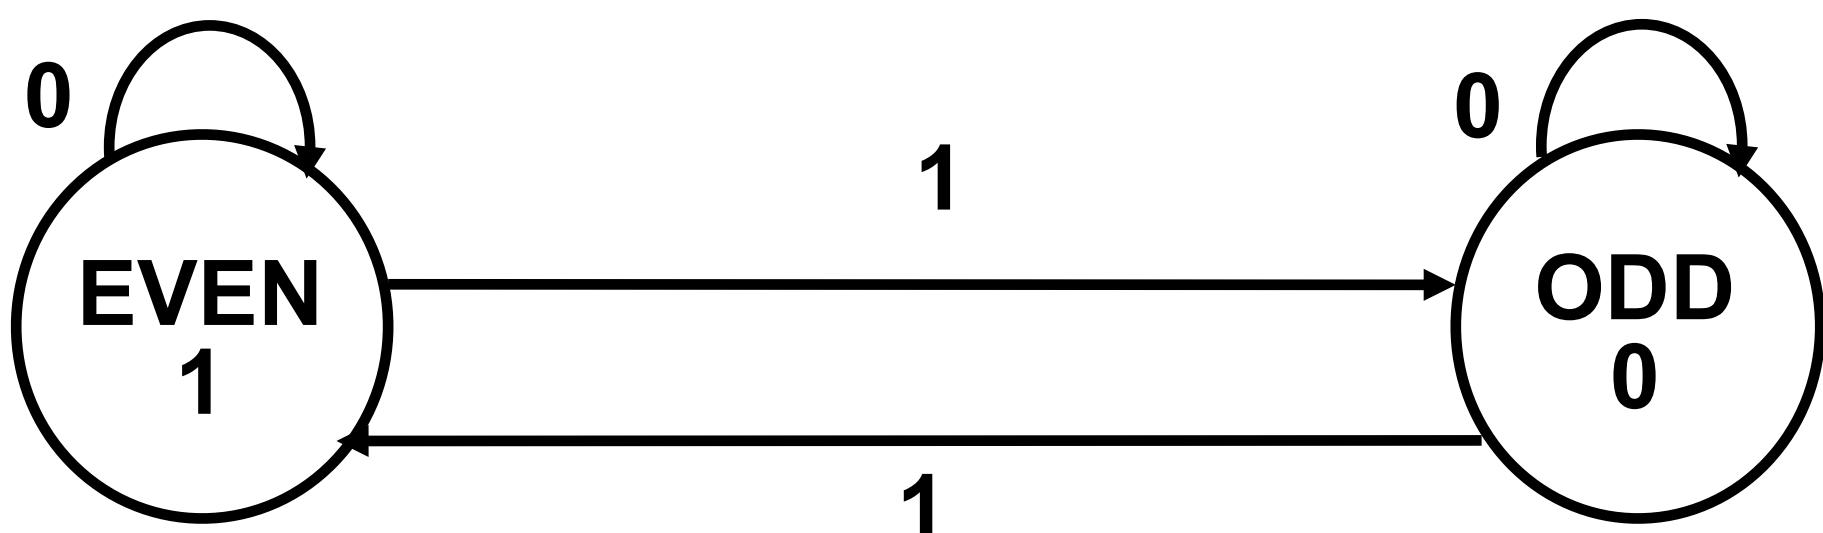
Bit 0 = 0, so go from ODD to ODD

Bit 1 = 1, so go from ODD to EVEN

Bit 2 = 1, so go from EVEN to ODD

Therefore the parity bit is 0 and the message to be sent is 0110

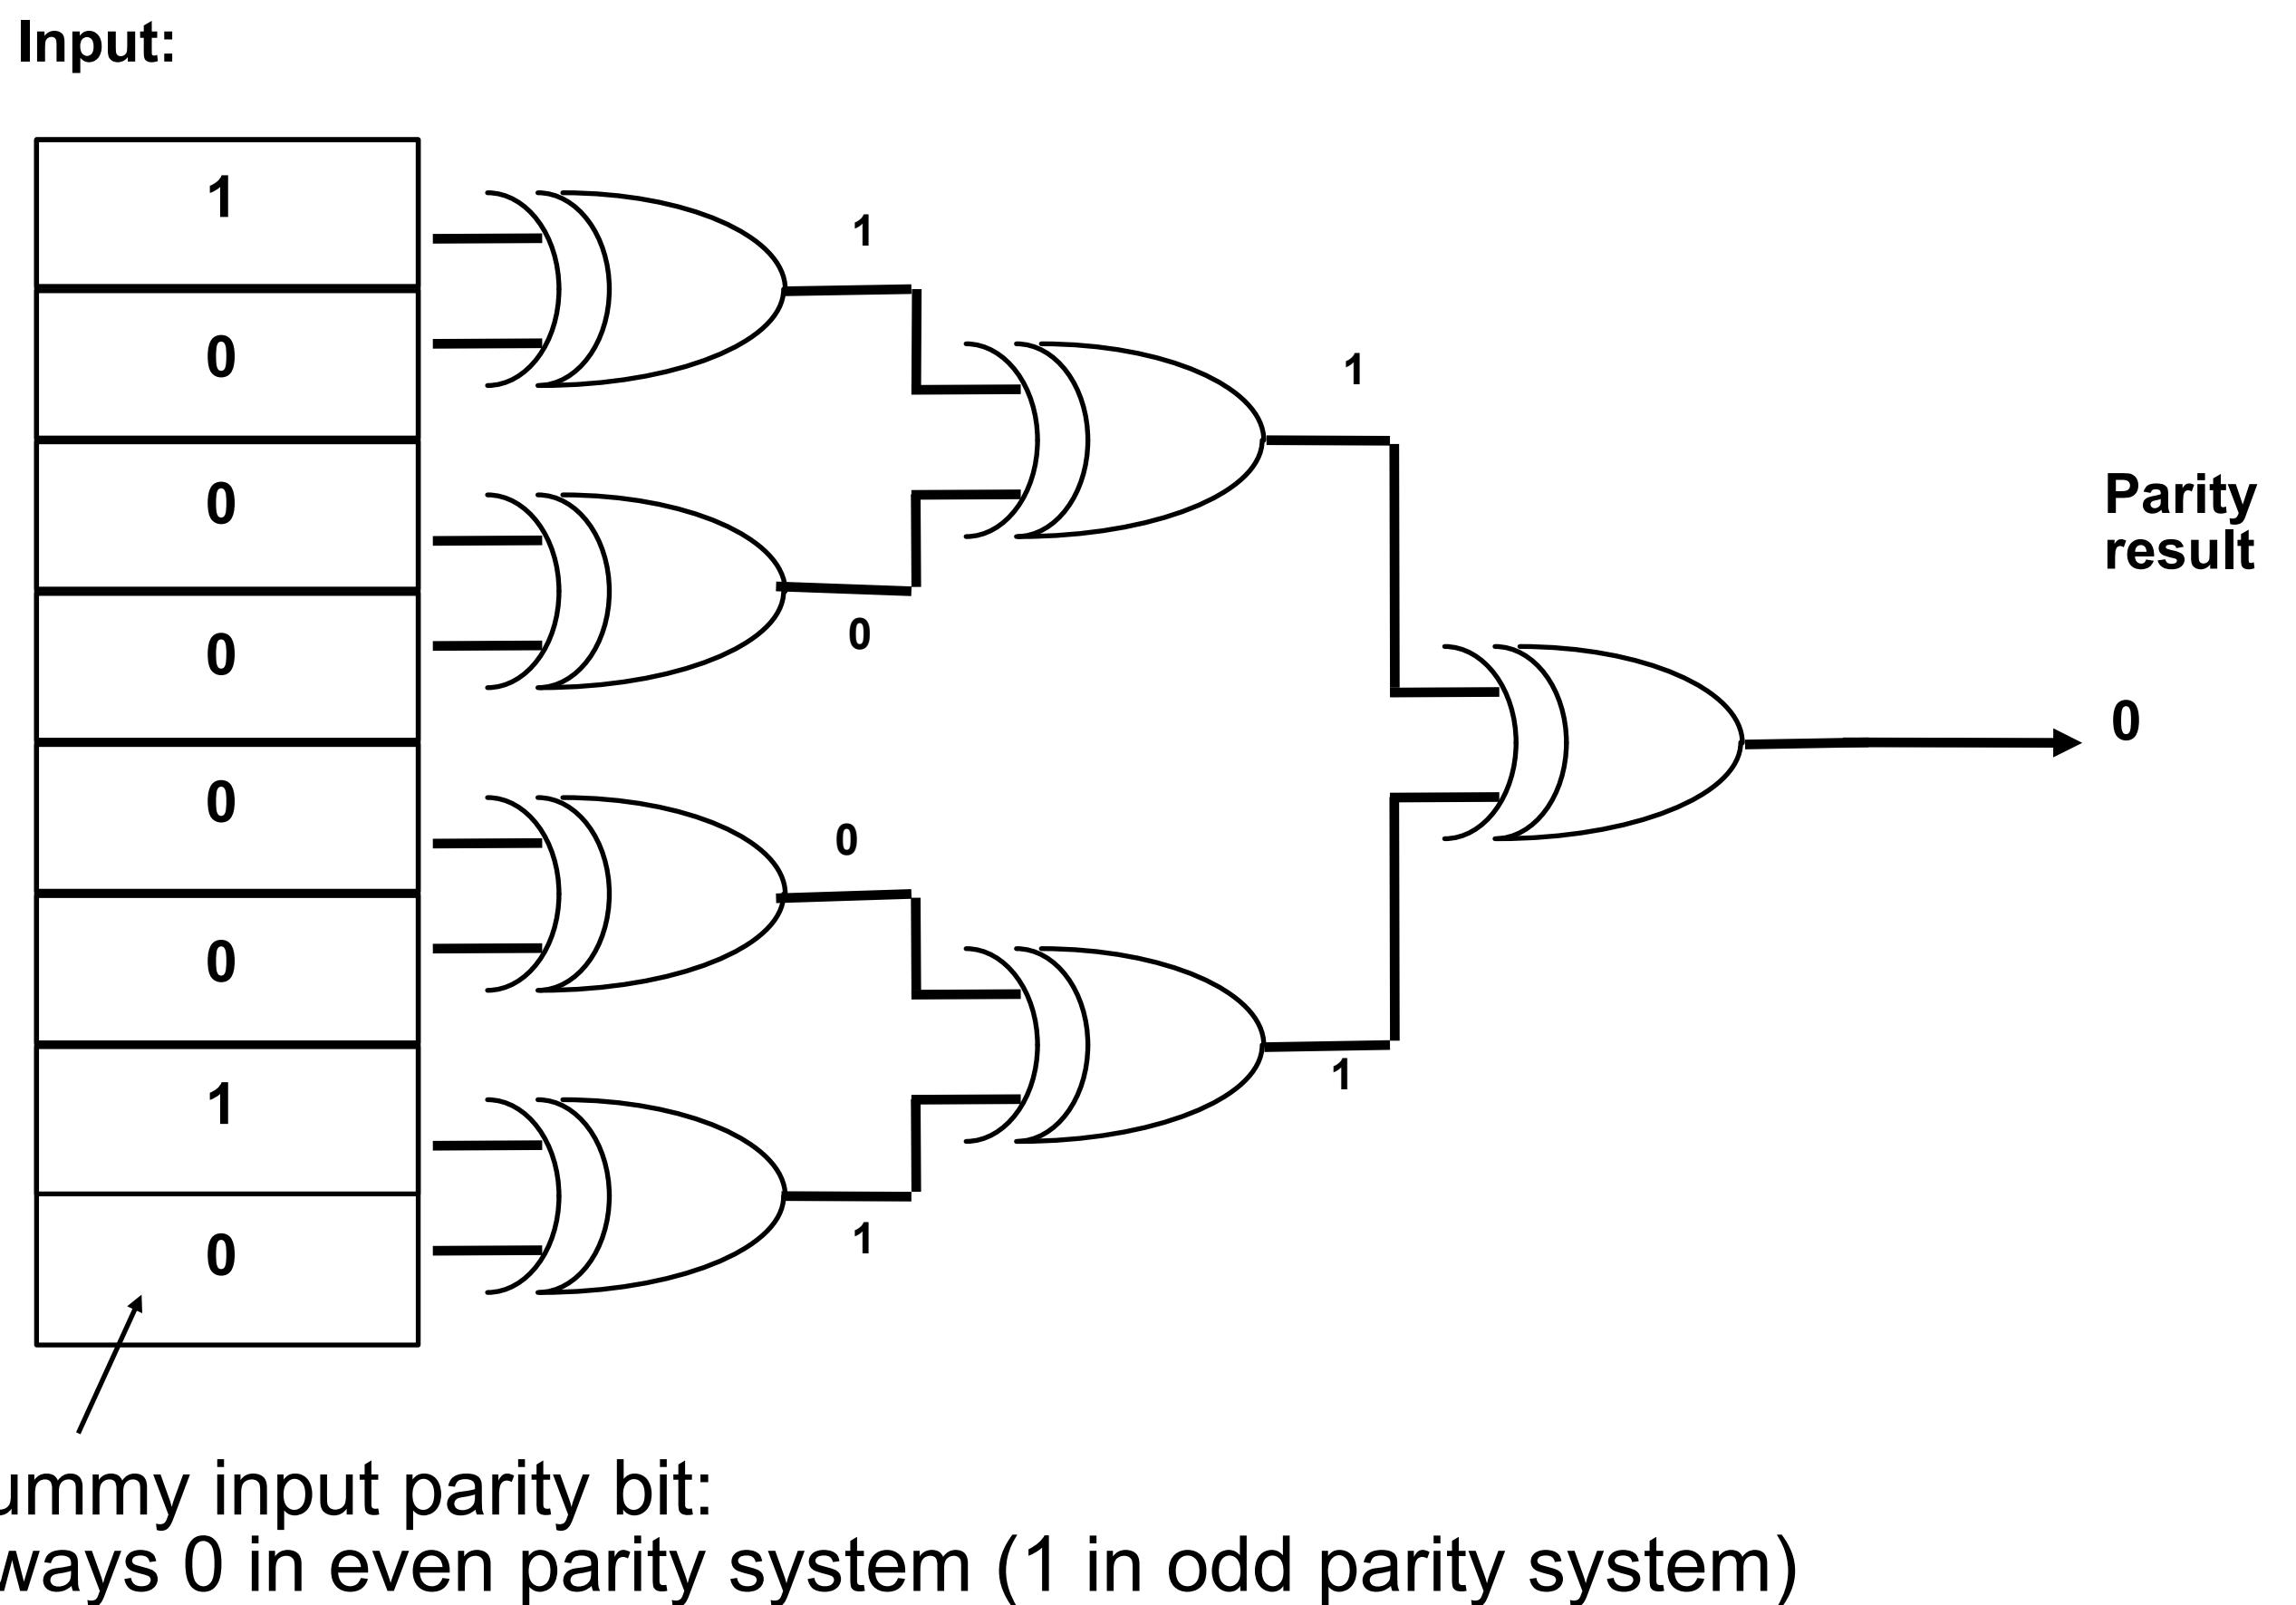
Parity may be computed using a two state Finite Automaton



Parity in Hardware

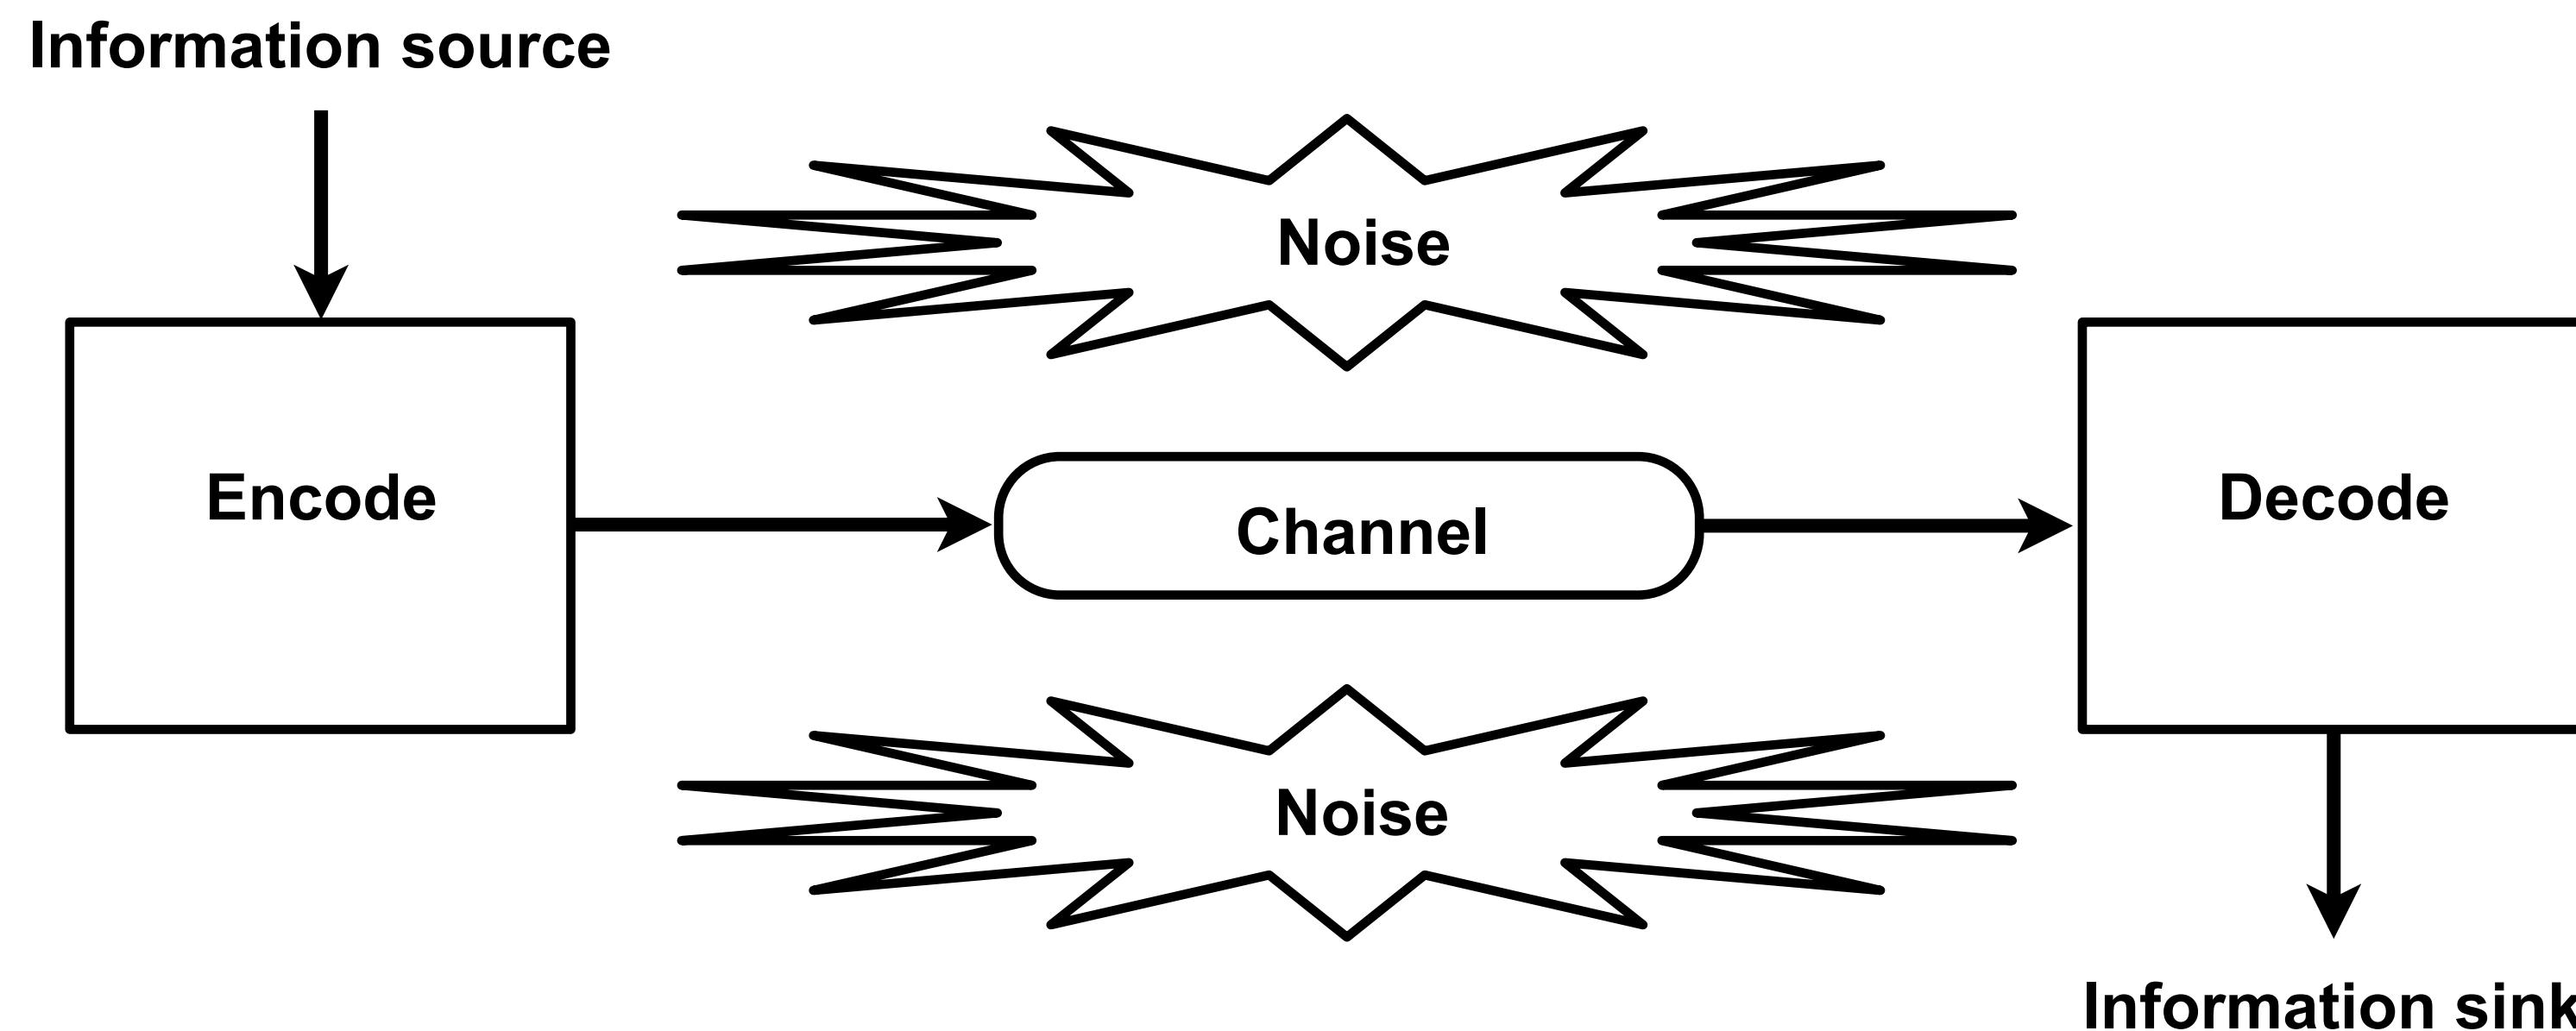
Computing parity in hardware
is popular

Hardware implementations
generally perform better than
software implementations,
despite being a relatively
simple digital logic circuit



Overview of Parity

Parity provides error detection without the need to significantly change our model for communication



Error Detection Using Parity

At our transmitter a message, e.g., 7 bits, is encoded

Compute parity (P) and append parity to message

Send modified message across channel

At our receiver the message, now 8-bits including parity, is decoded

Compute parity (Q) from data bits of received message

Compare computed parity Q with received parity P in message and flag ERROR if these are not consistent

Error Detection Using Parity - Example

Transmitter and receiver have been designed to use even parity

Message: 100 0101

Transmitter computes P=1

Transmits: 1100 0101

An error in the channel changes one bit

Received: 1101 0101. Receiver notes received P=1

Receiver uses the data bits (101 0101) to compute Q=0

Receiver compares received P and computed Q.

An error has been detected, thus receiver will ask for retransmission

Burst Errors

Remember that, if the probability of an error in the channel is low, then the probability of two errors close together is even lower

In practice, multiple errors often occur in bursts, e.g., a 1ms dropout on a 56k baud modem link can error up to ~58 bits

Worse on broadband, e.g, ADSL

Parity can detect single isolated errors but an even number of errors close together leaves parity unchanged, e.g., 0100 0001 => 0100 1101 => parity unchanged!

The form of parity we have seen is inadequate for burst errors

A possible solution is to compute “checksum” values, i.e., the parity of each bit-column over multiple bytes

Detecting Burst Errors - Example

Checksums vastly improve our error detection capabilities and allow for the detection of burst errors

1. To send the following message over a channel:

(NB 7-bit ASCII without parity)

2. Calculate bit-column parity "checksum" values,
e.g., using even parity

100 1011 = ASCII 'K'

3. Send "Message**K**" instead of "Message"

M	100 1101
e	110 0101
s	111 0011
s	111 0011
a	110 0001
g	110 0111
e	110 0101

This system detects all burst errors < 14 bits
Fails if an even number of errors occurs in a bit-column

Error Correcting Codes

The checksum codes we have seen (those with bit-column parity) could detect some burst errors

We can extend to an Error Correcting Code (ECC) by calculating both character and bit-column parity

M	0	100	1101	1. ← character parity
e	0	110	0101	
s	1	111	0011	
s	1	111	0011	
a	1	110	0001	
g	1	110	0111	
e	0	110	0101	
K	0	100	1011	

2.	bit- column parity
----	--------------------------

ECCs - Example

At transmitter we computer and transmit

Compute character parity bits, then bit-column parity value, then transmit message with these added bits

At receiver we re-compute, compare and correct if needed

Compute character parity bits, ignoring received parity bits

For each character, compare re-computed character parity with received parity

If not equal, ERROR somewhere in row

Re-compute bit-column parity value

Compare re-computed value with received value

If not equal, ERROR somewhere in column

If a single error detected, correct by inverting bit in message at (row, column)

ECCs - Graphically

A graphical representation allows us to think clearly about the potential issues and capabilities of a particular ECC

✓	0	0	100	1101	M
✓	0	0	110	0101	e
✓	1	1	111	0011	s
x	0	1	11 <u>0</u>	0011	c
✓	1	1	110	0001	a
✓	1	1	110	0111	g
✓	0	0	110	0101	e
	0	100	1011		
	0	101	1011		
	✓✓x	✓✓✓✓			

Key:

- 1101 – received message bits
- 0 – bit changed by noise
- 1011 – received parity bits
- 1011 – re-computed parity bits
- ✓✓x – comparisons

This ECC can detect and correct single errors

It can also detect, but not correct, certain combinations of multiple error

Common Memory Components

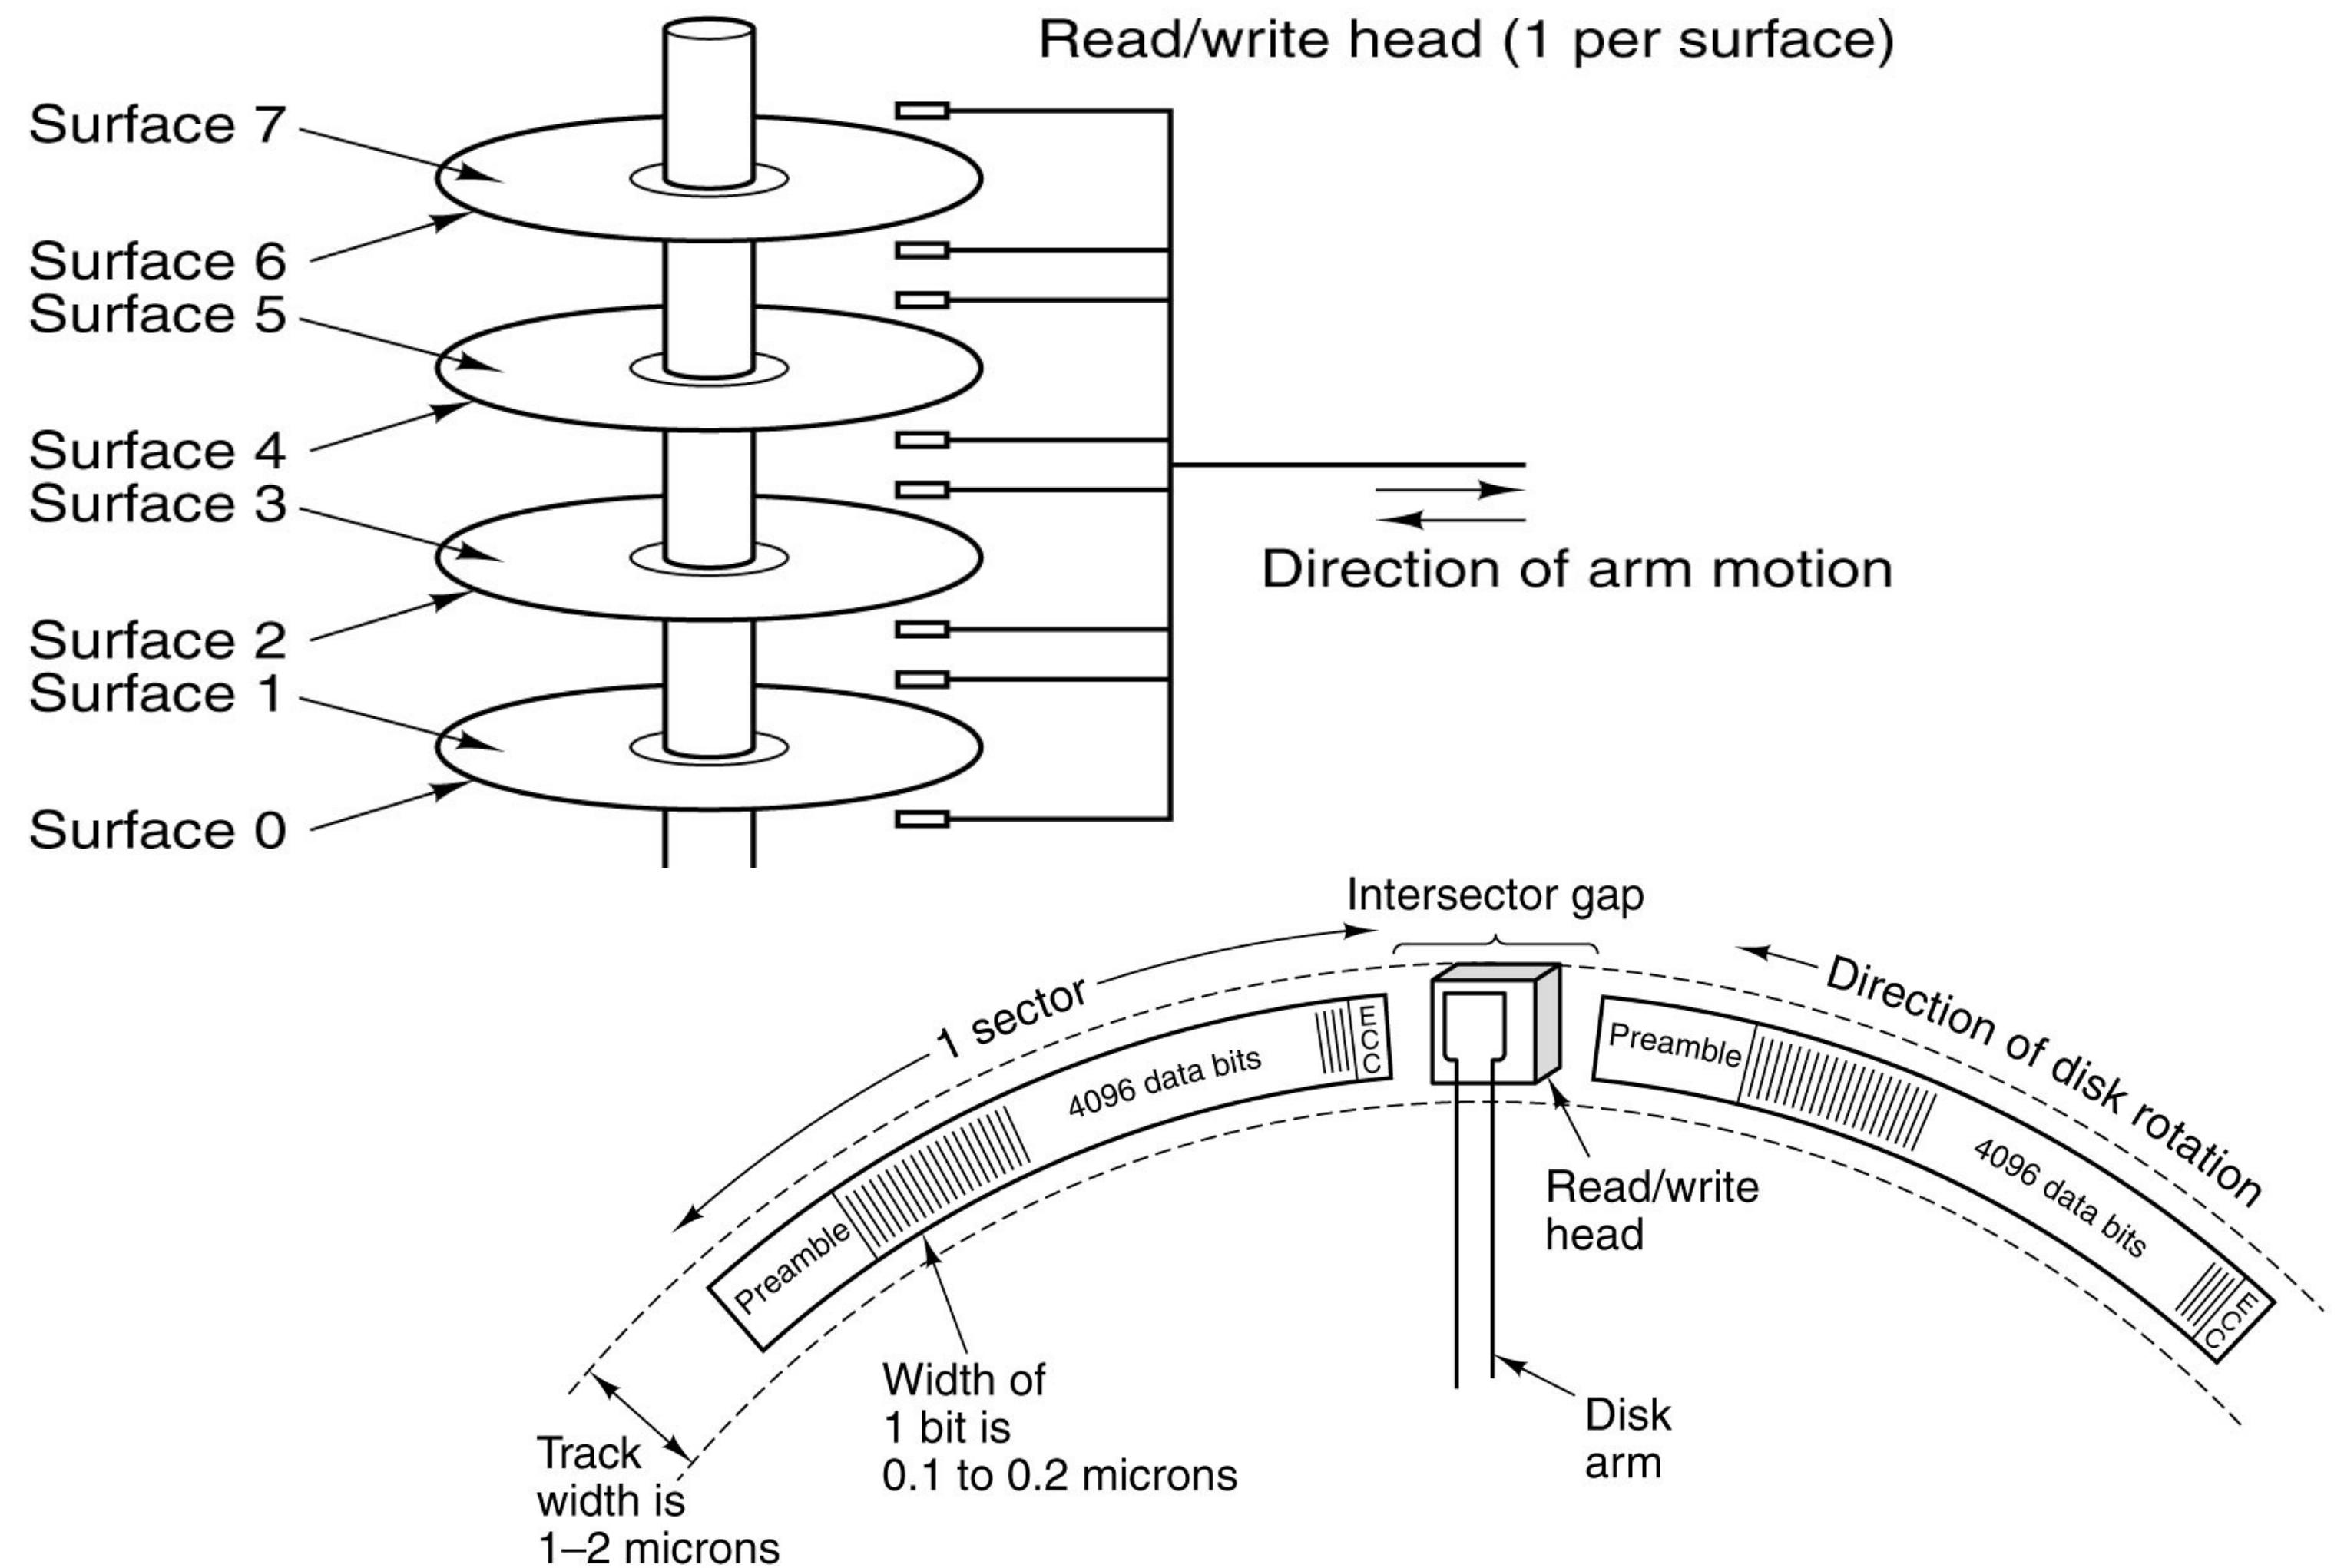
Magnetic Hard Disks - Platters, Surfaces and Heads

Record data by magnetising a thin film of ferromagnetic material on a disk

One track contains many sectors

Each sector is separated by an inter-sector gap.

Sector contains preamble to allow head to be synchronised before read/write, data and ECC



Magnetic Hard Disks - Performance - What Can We Learn?

Seek times - time to move arm to correct track

Typically about 5ms between different tracks but can be below 1ms for consecutive tracks

Rotational latency

Constant angular velocity, i.e., motor spins at constant speed

Drives available with rotational speeds of 5400 RPM, 7200 RPM or 10800 RPM

Average delay (half a rotation) = 3 to 6ms

Sector read time is about 0.013ms, which means that seek time and rotational latency dominate access time.

Formatting (preambles, ECC etc...) reduces capacity by about 15%

Hard Disks - Increasing Capacity with Zones - More To Learn?

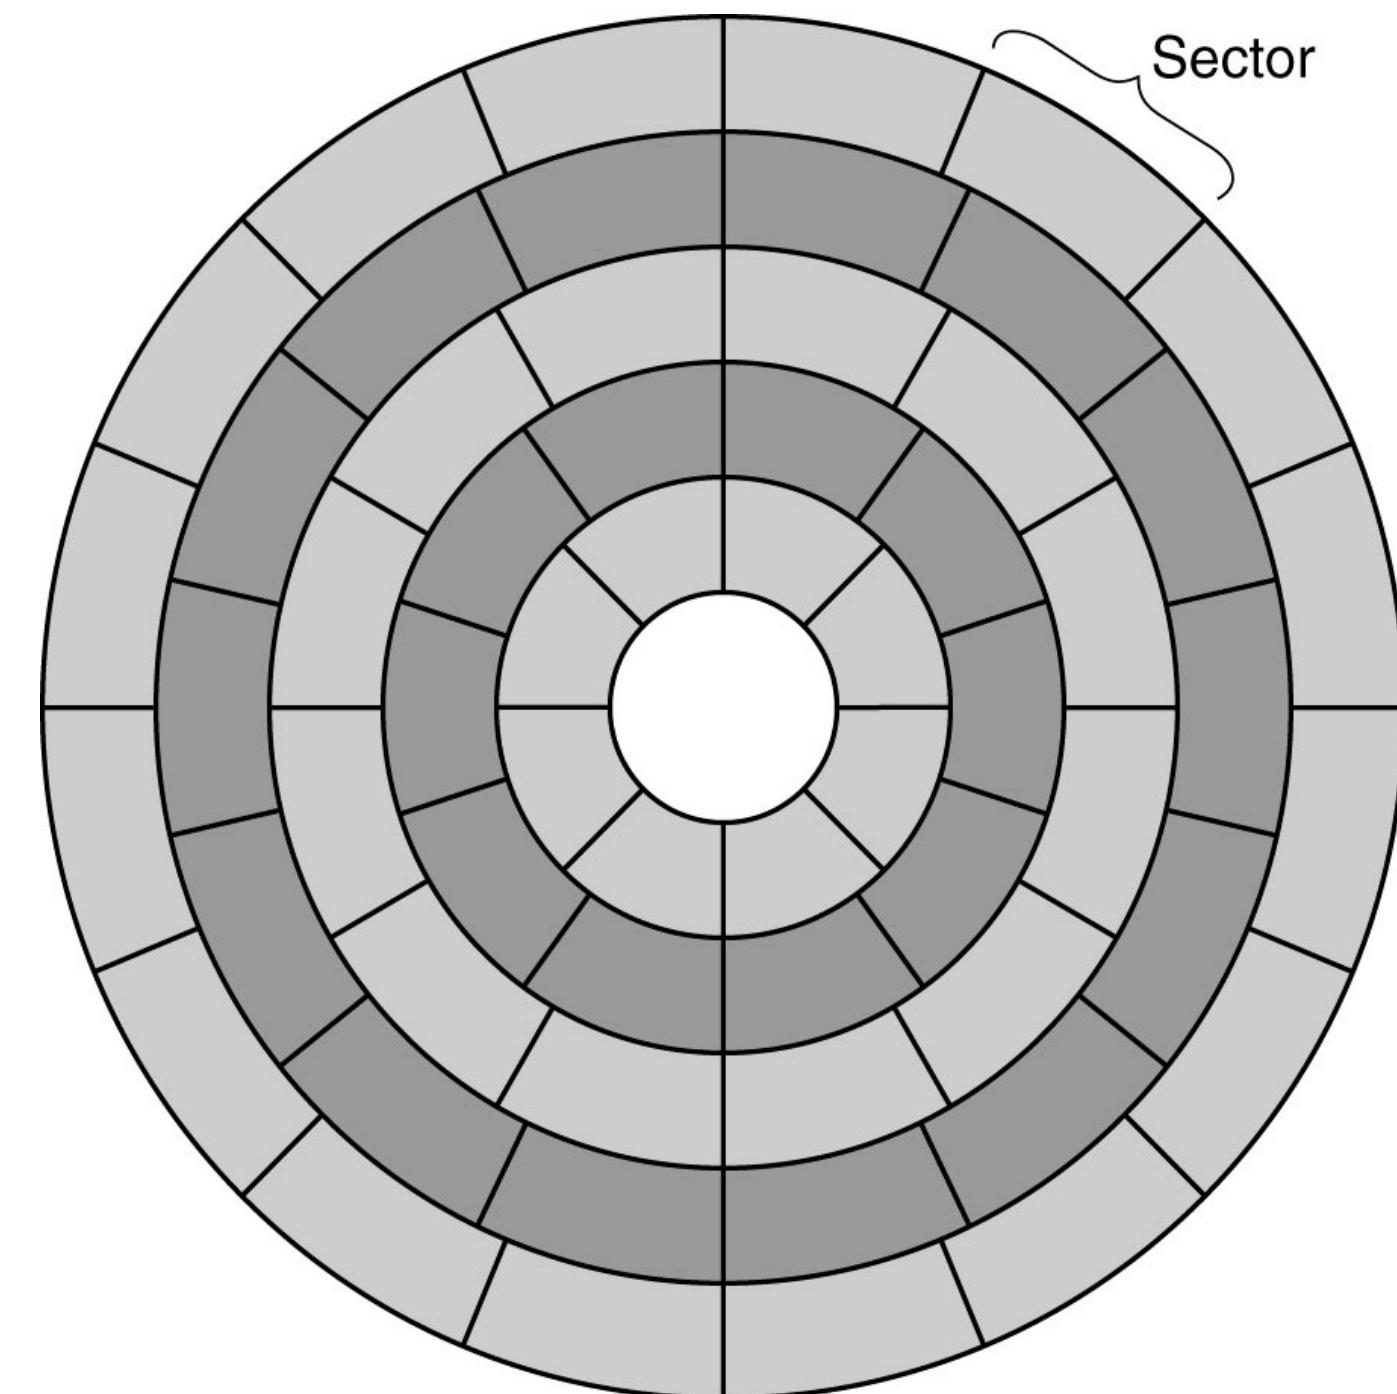
Older drives have a constant number of sectors per track, which meant that, due to constant angular velocity, bit density is lower on outer tracks

Most Recent disks are zoned (10-30 zones), with more sectors in the longer (outer) tracks

Many tracks per zone

More sectors gives higher capacity

Graphic shows a disk with 5 zones



Optical Disc - Spiral Track

Originally designed to hold music

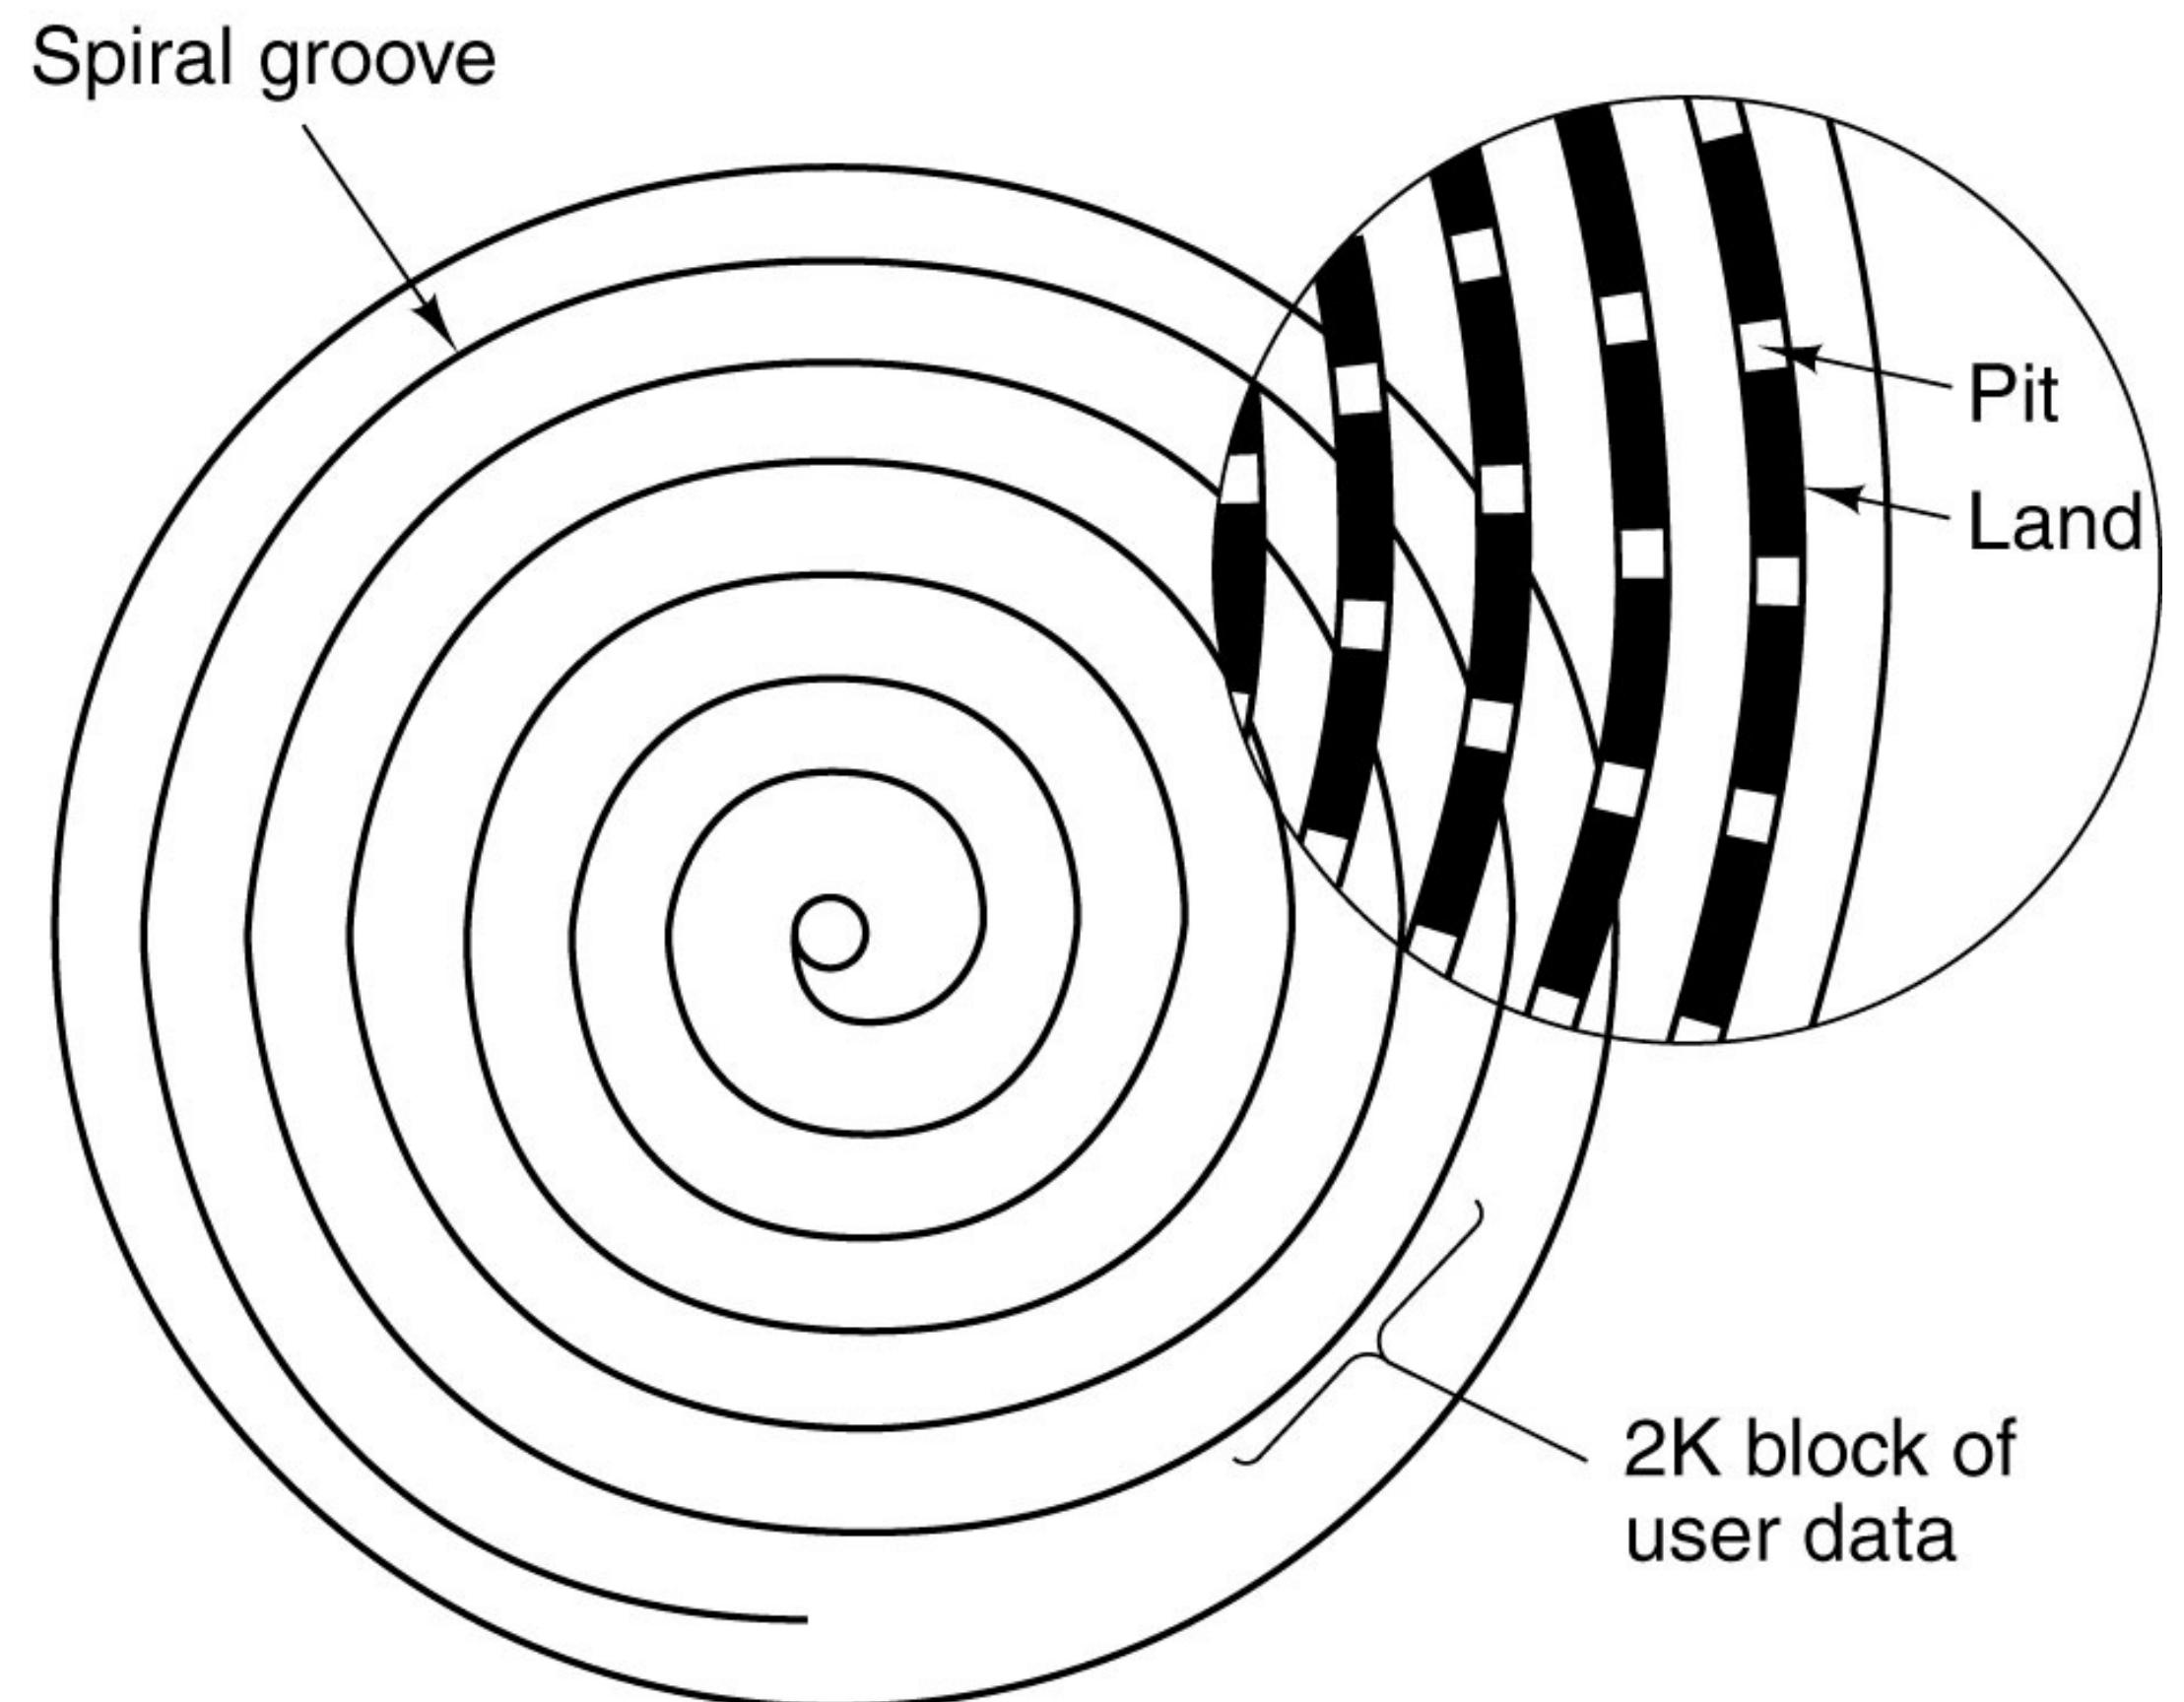
Max 74 mins - Beethoven's Ninth Symphony!

Spiral is 3.5 miles long if unwound

6,000 tracks/cm

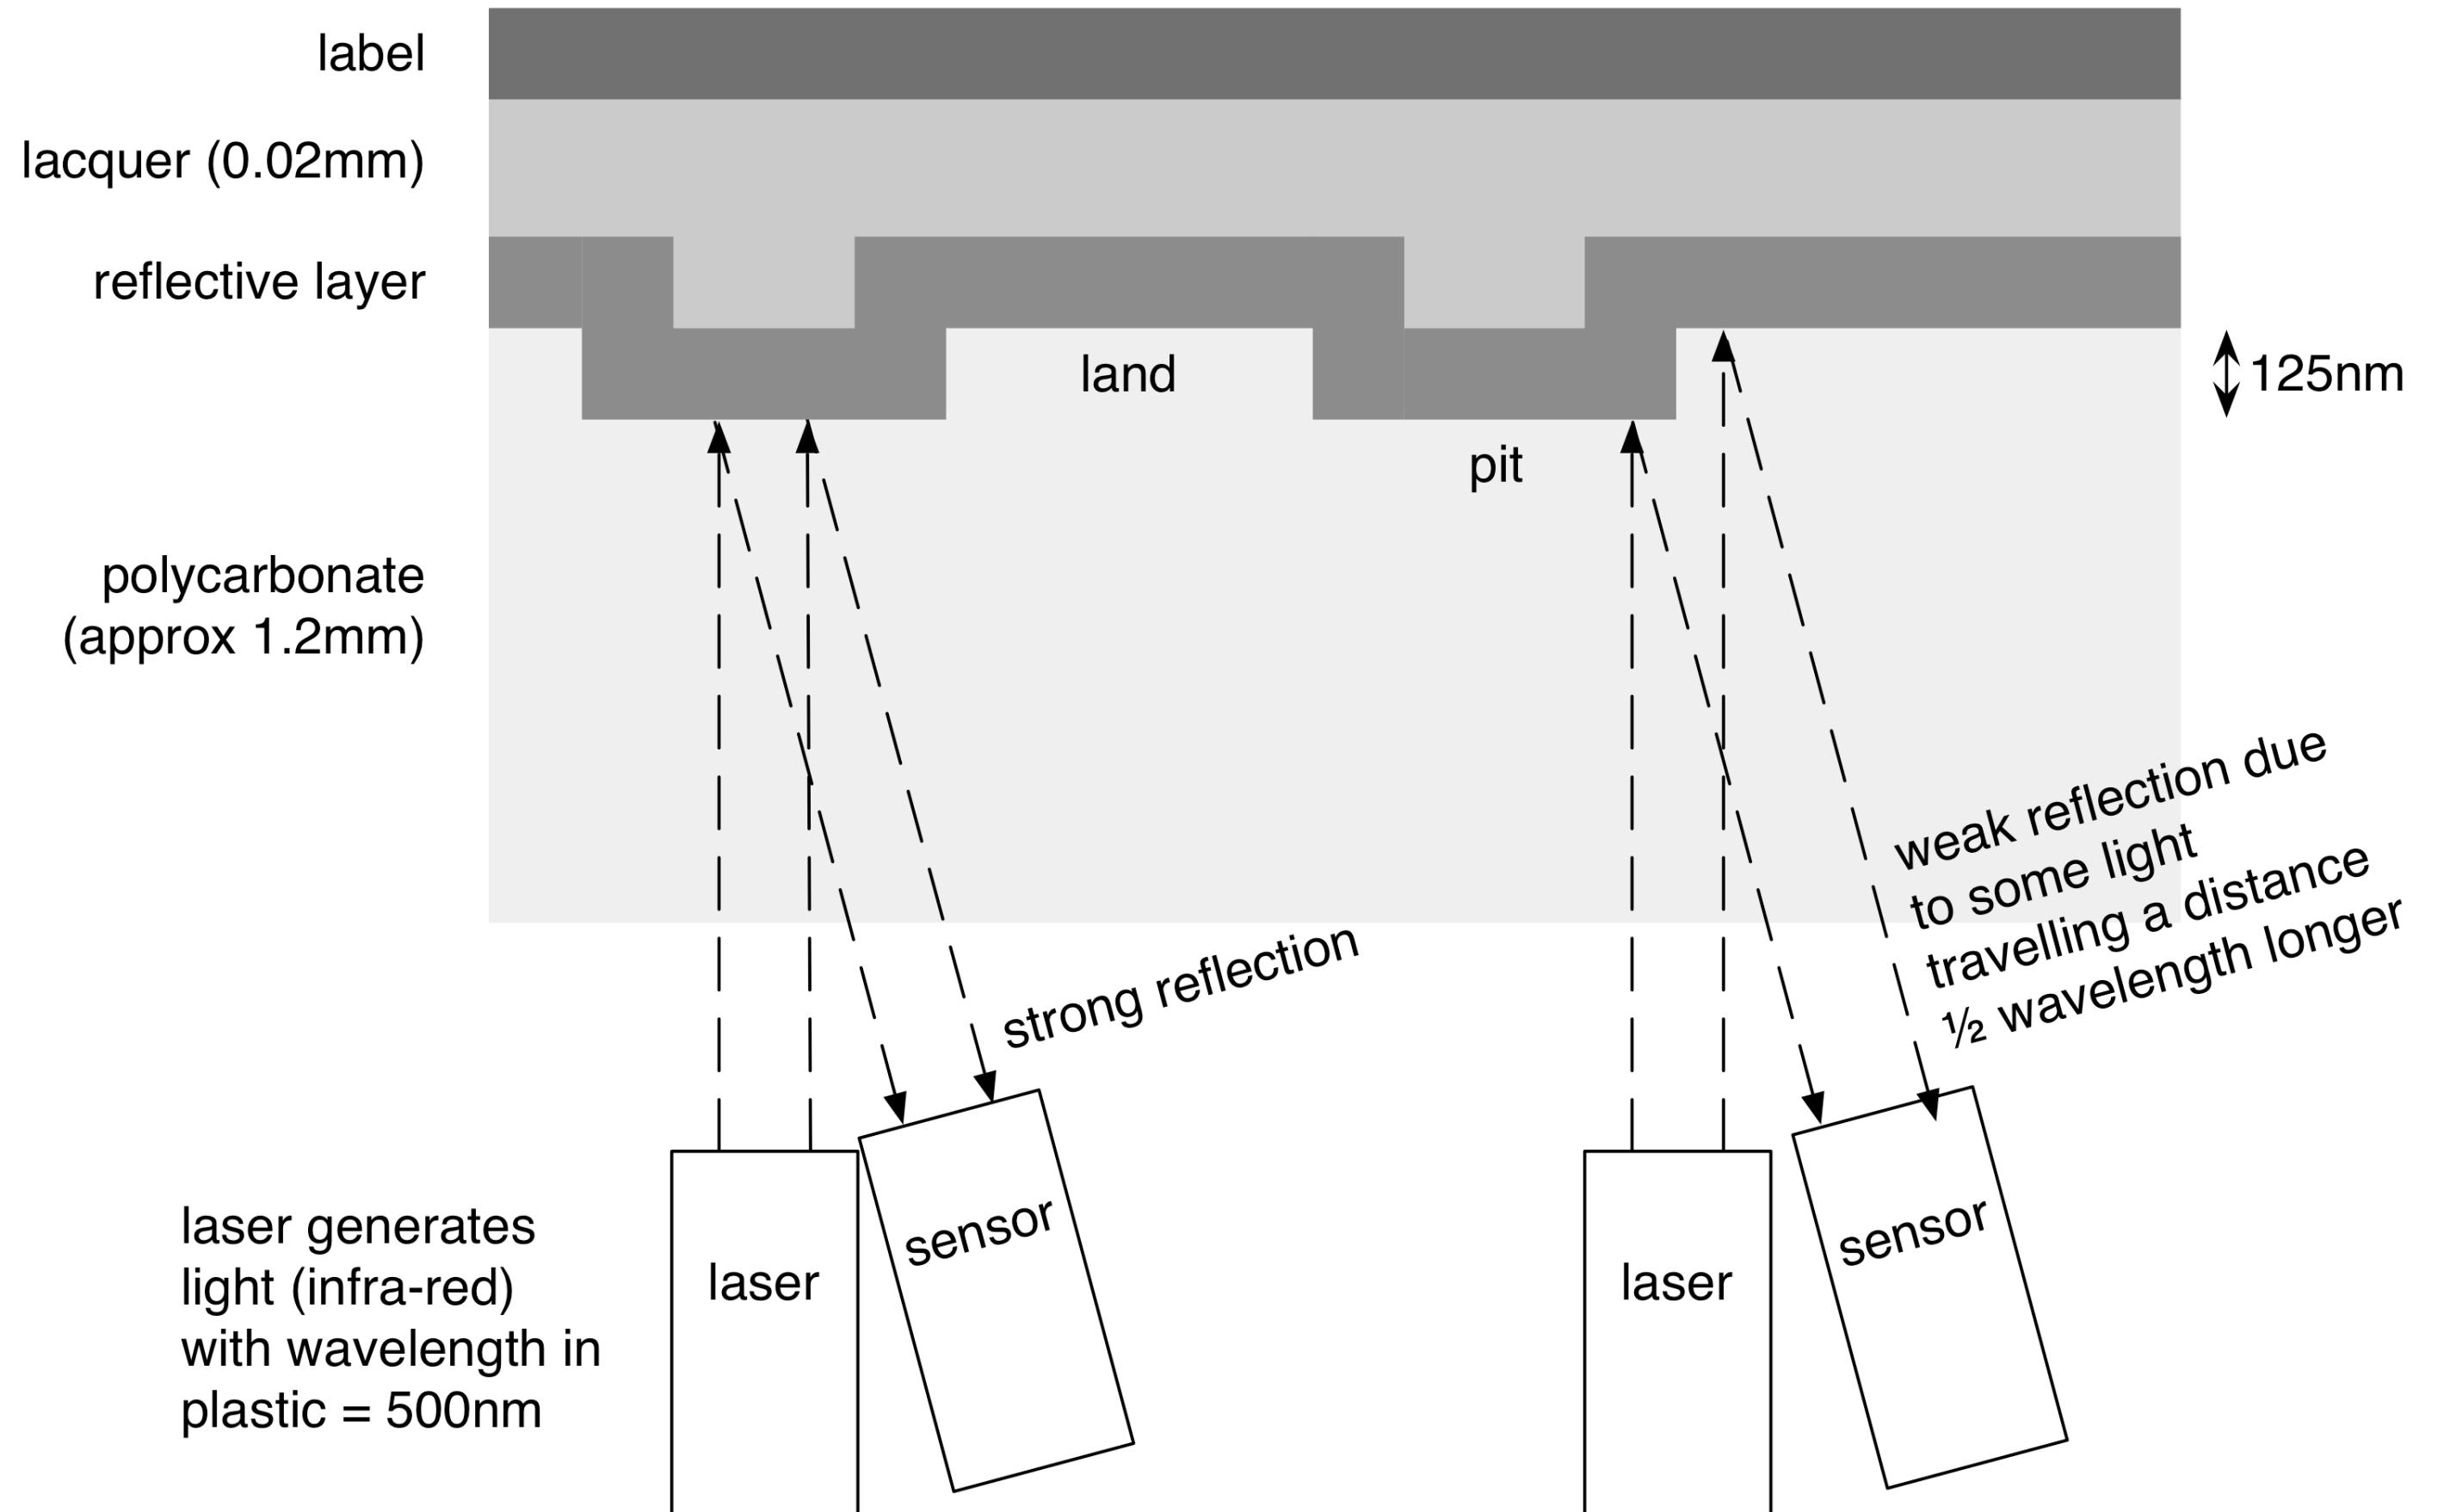
Data encoded as "pits" and "lands"

15,000 pits and lands/cm



Optical Disc - Spiral Tack Detail

Consider a much closer look at how optical discs store data.



Optical Disc - Data Encoding

A change in state (pit/land, land/pit) signifies data is a 1

Data represented as bytes, but each 8 bit byte is mapped to a 14 bit symbol (first level of CD-ROM ECC)

Symbols are chosen such that there are always two 0s between 1s and a max of ten 0s in succession

Of the 214 possibilities, only 267 values fit the rules

This is more than sufficient for the 256 values (for the 8 bits) that are to be represented

42 symbols make a frame (containing 24 bytes of data)

98 frames make a sector (containing 2048 bytes)

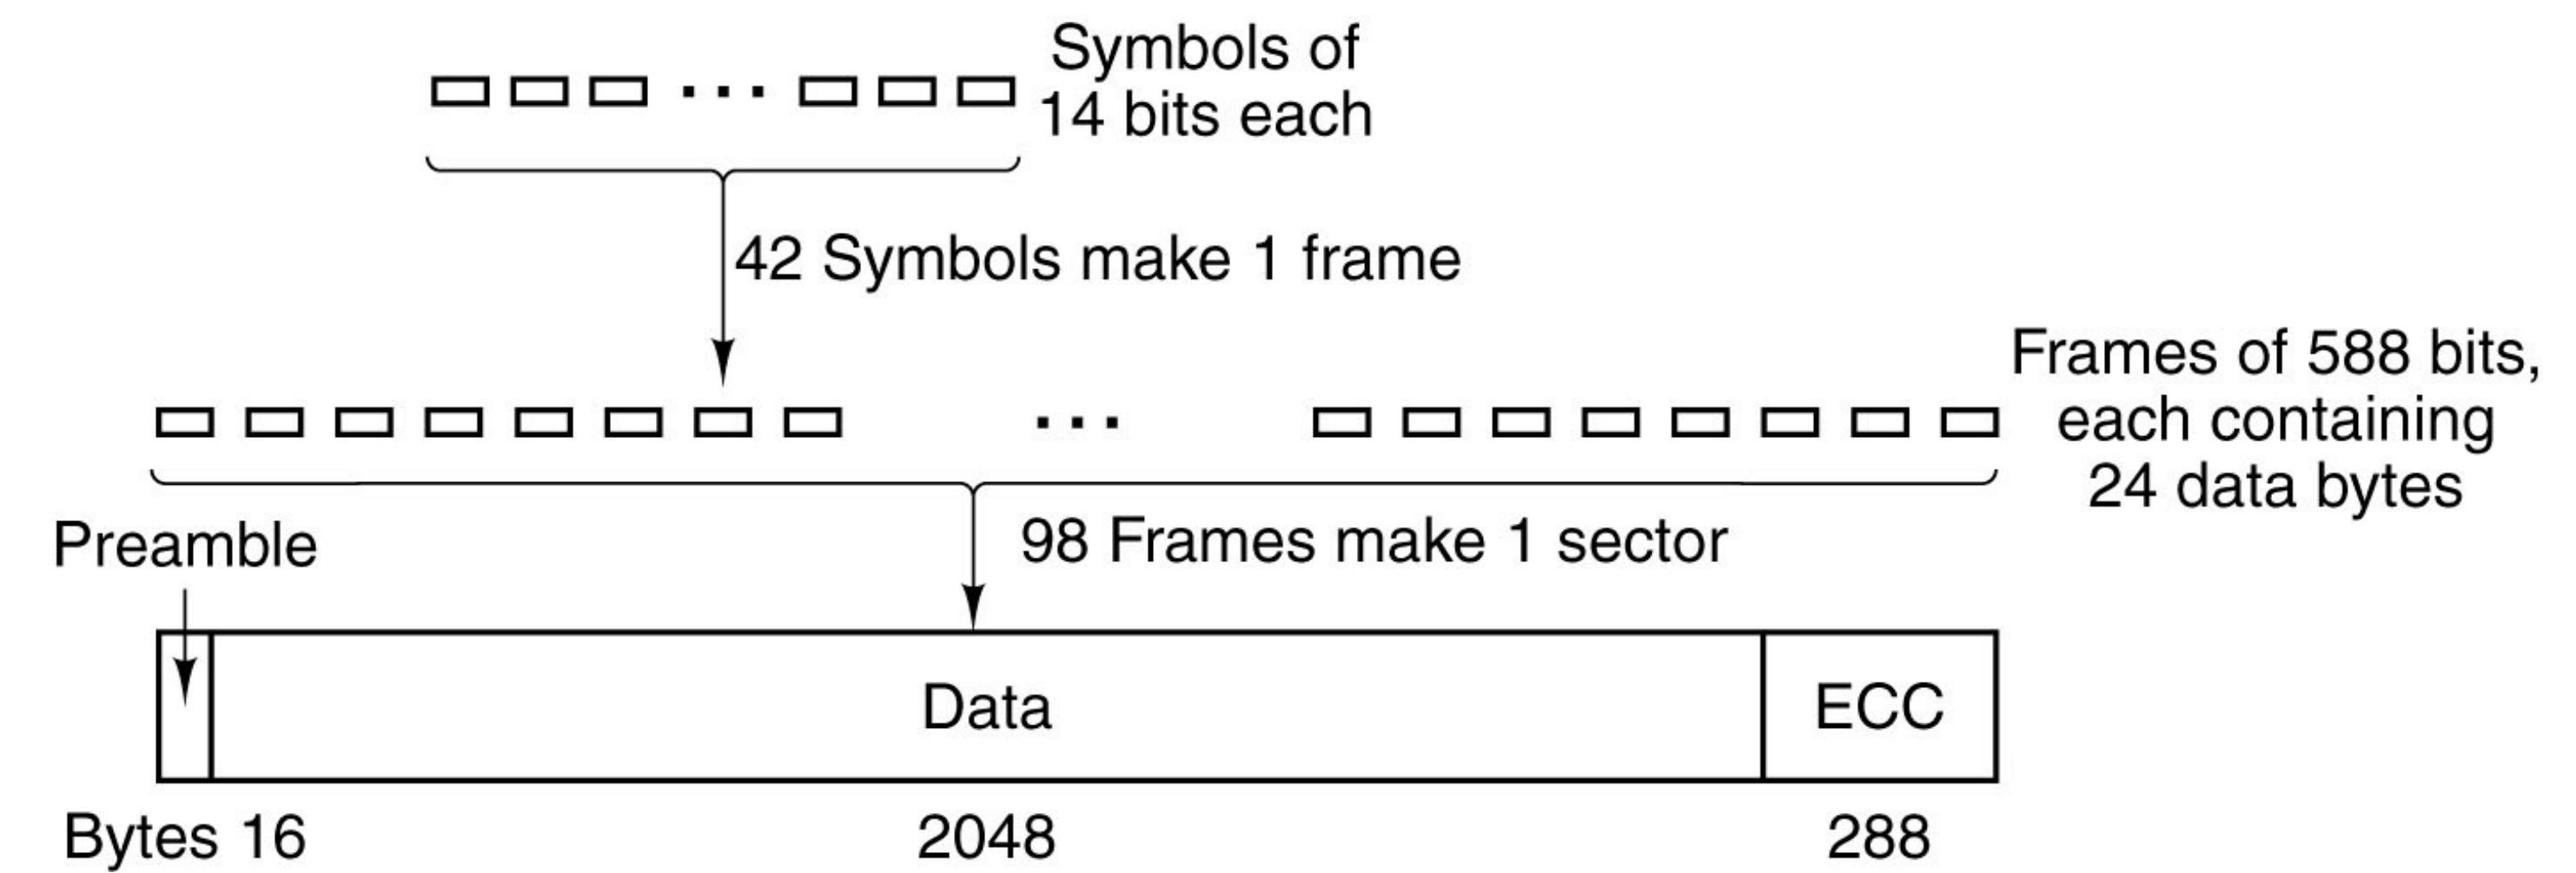
Optical Disc - Data Encoding Graphically

Multiple separate error
correcting schemes

Single bit errors
corrected at symbol level

Short burst errors
corrected at frame level

Any residual errors
corrected at sector level



Optical Disc - Error Prevention

Most errors are burst and local, e.g., those caused by scratches

Cross Interleaved Reed-Solomon (CIRC) code is used to spread data out over the surface of the disk.

Burst errors of up to 2mm track can be corrected

Data accounts for only 28% of bits stored, thus we pay a heavy price for great reliability

Optical Disc - Performance

Constant linear velocity, i.e., motor spins at variable speed, to keep the pits and lands moving over the laser at a constant speed

For a “1x” drive, i.e., audio CD speed we have 120 cm/sec (data rate of 150kbytes / sec) and 200 RPM (outer) – 530 RPM (inner part of spiral)

An “8x” drive spins 8 times faster (1600-4000 RPM)

Above 12x speed, CAV is usually used, and the “x” number denotes data transfer rate at the outer edge of the disc, i.e., the max rate

Seek times averages are typically 80ms, which is at least an order of magnitude worse than hard disks

